

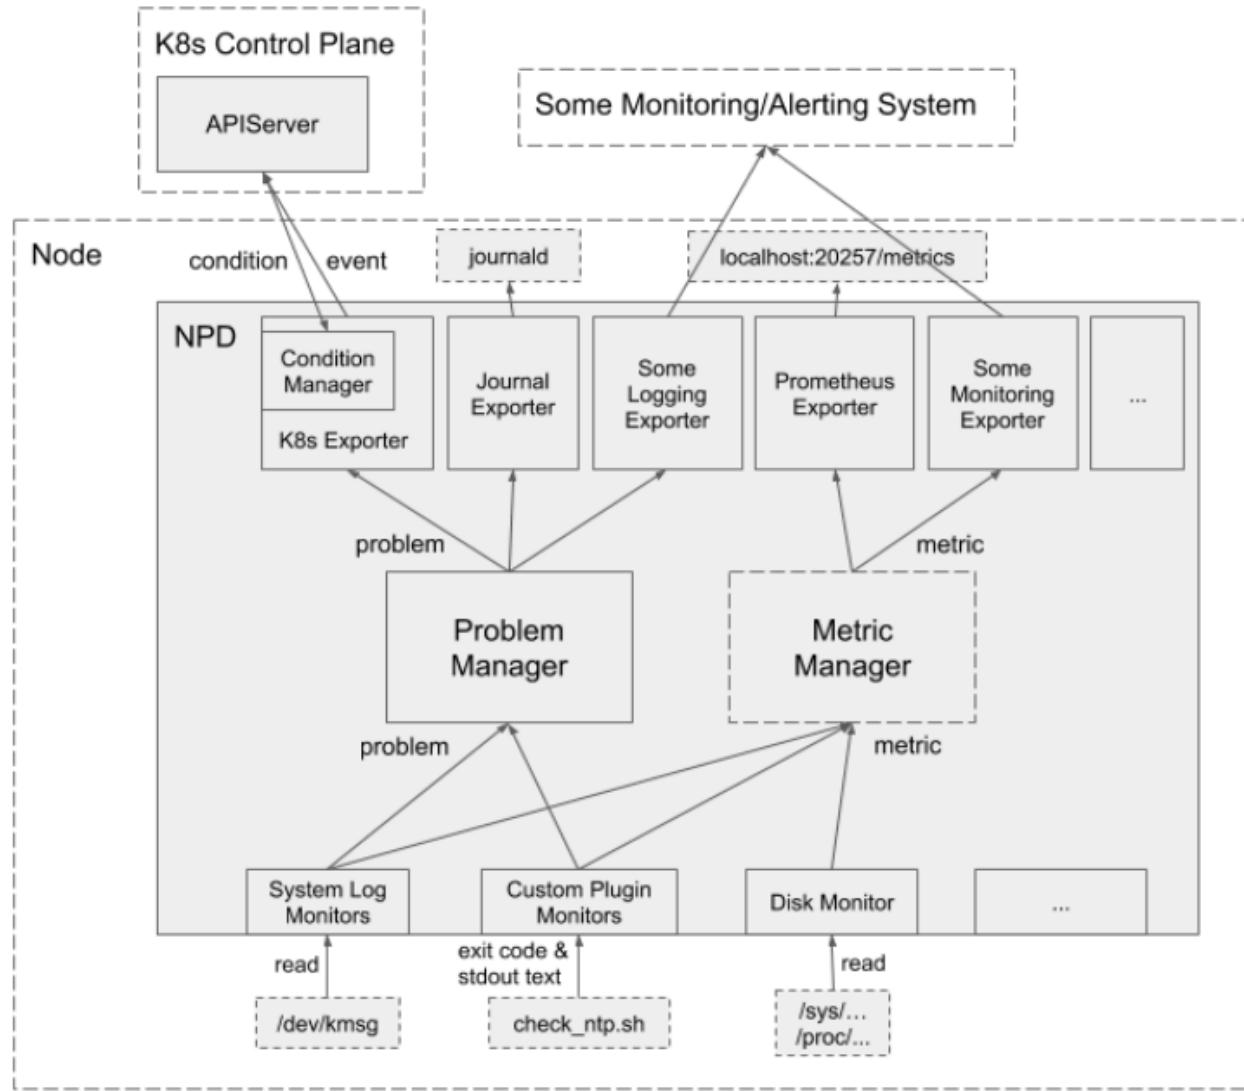
Predictive Maintenance

Topics

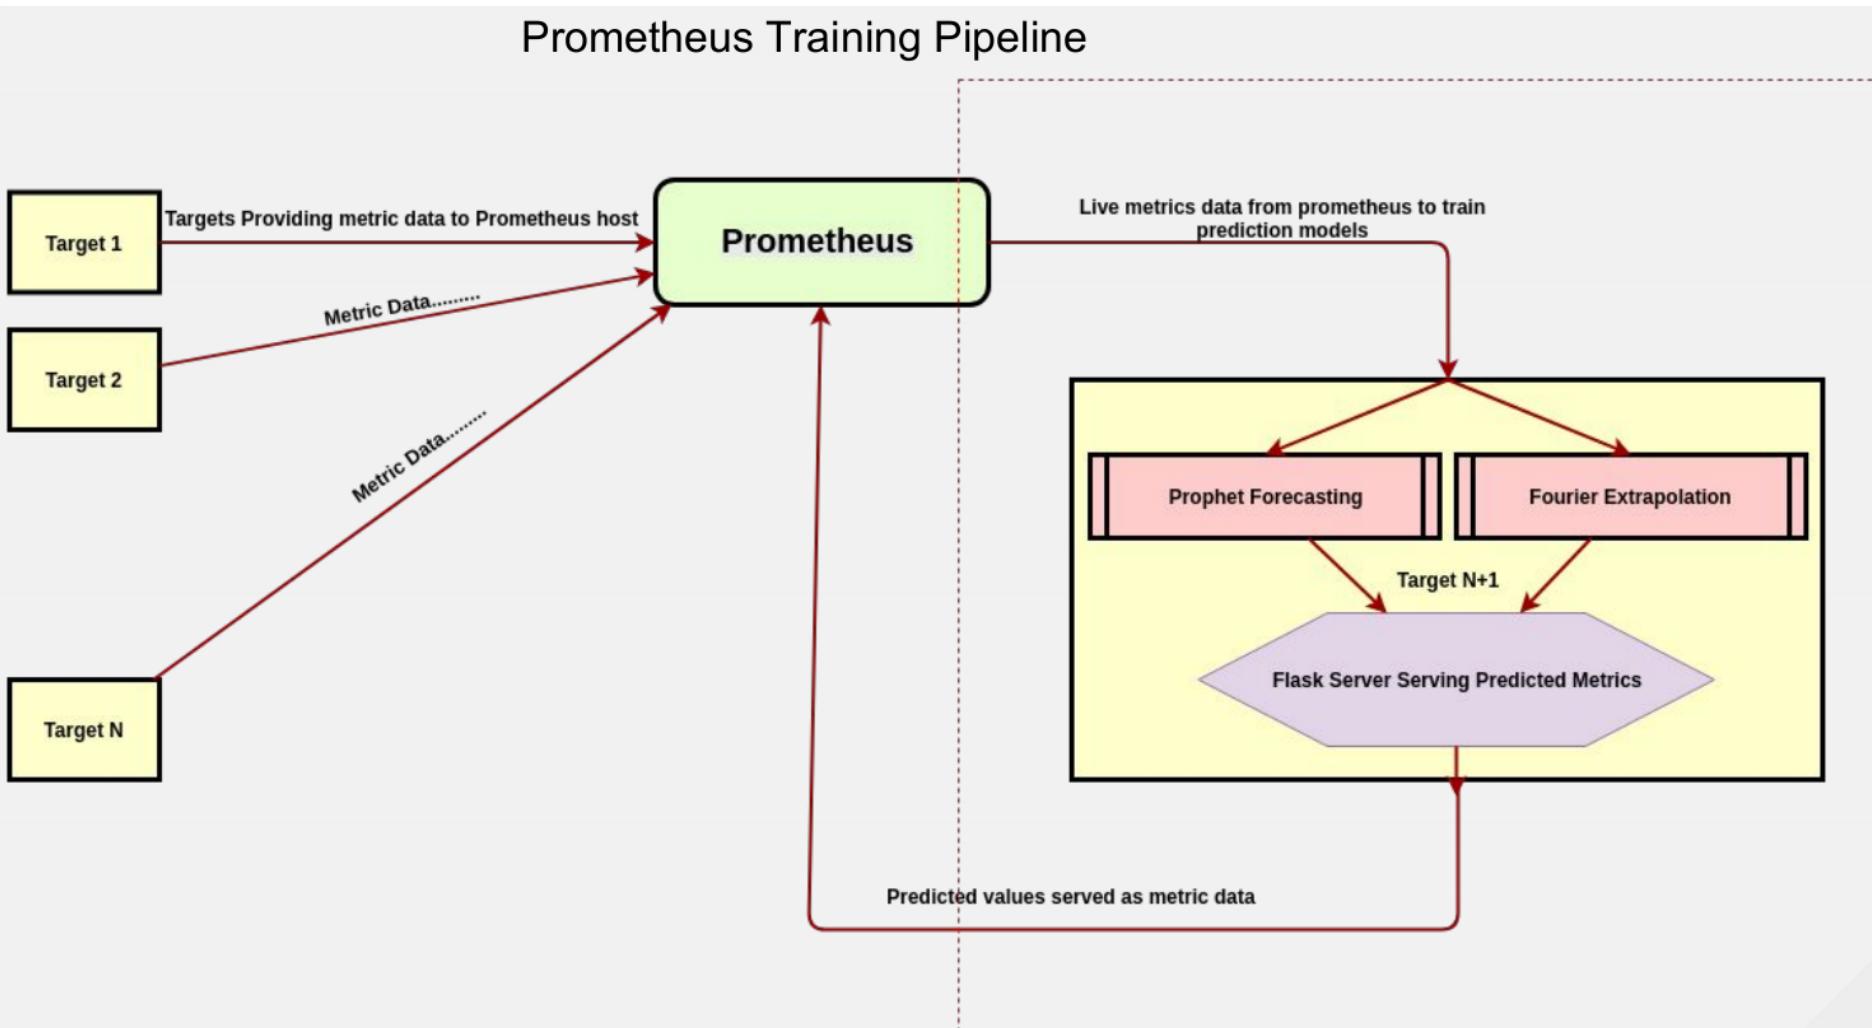
- Maintenance Scheduling
- Anomaly Detection
- Failure Prediction
- Improving Utilization
- Utilization Prediction
- AIOps platform
- CAP保证

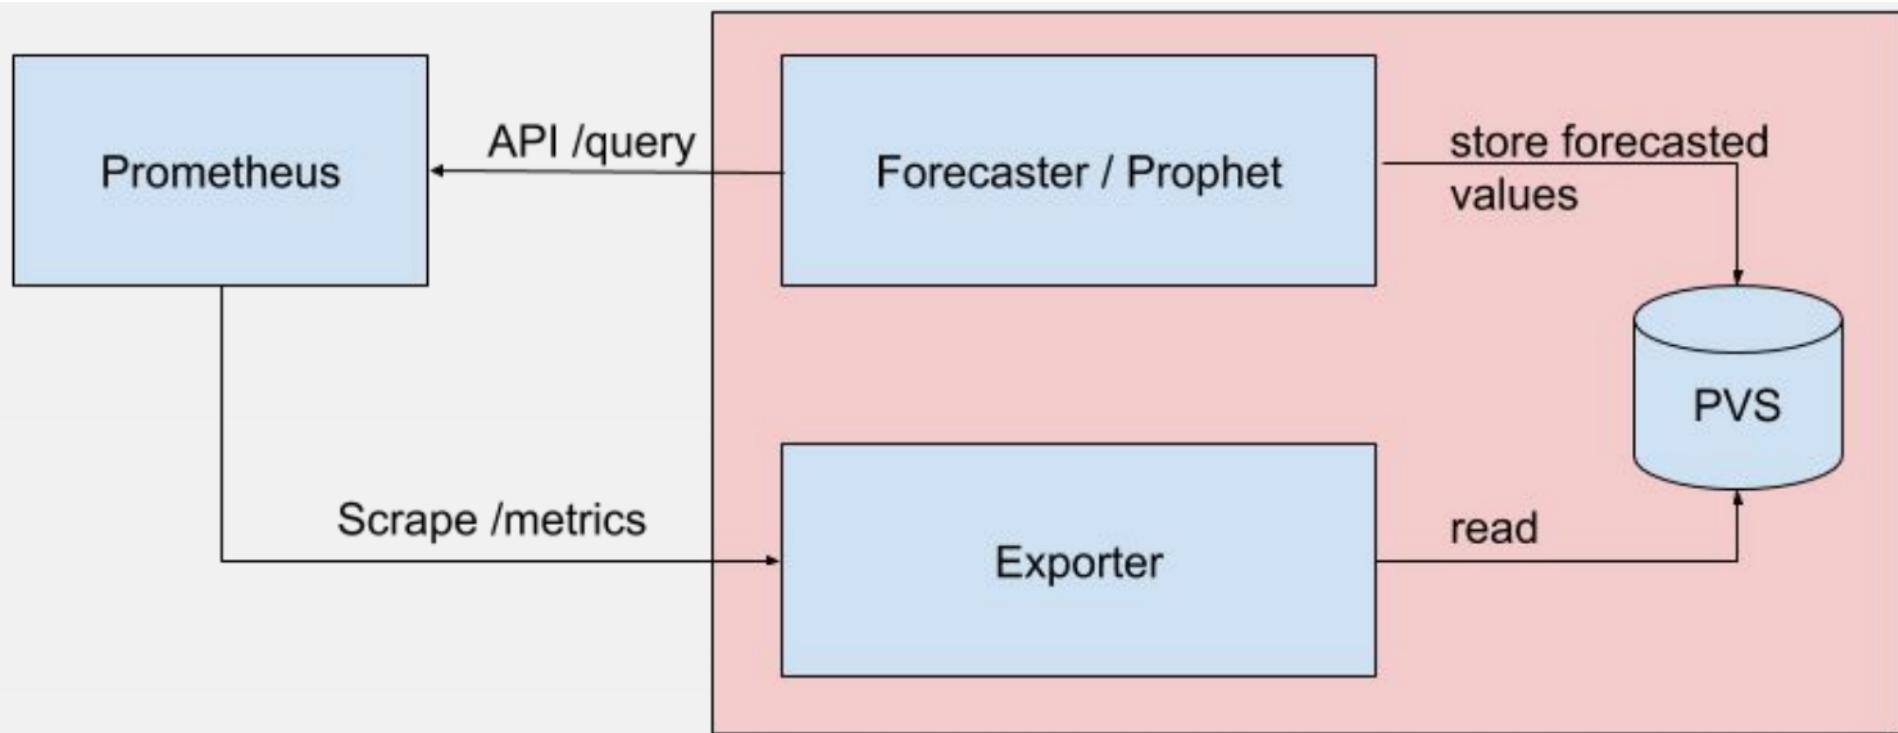
Python TensorFlow models serving with Go

- <https://tech.travelaudience.com/training-tensorflow-models-in-python-and-serving-with-go-1b2a9386b0ff>



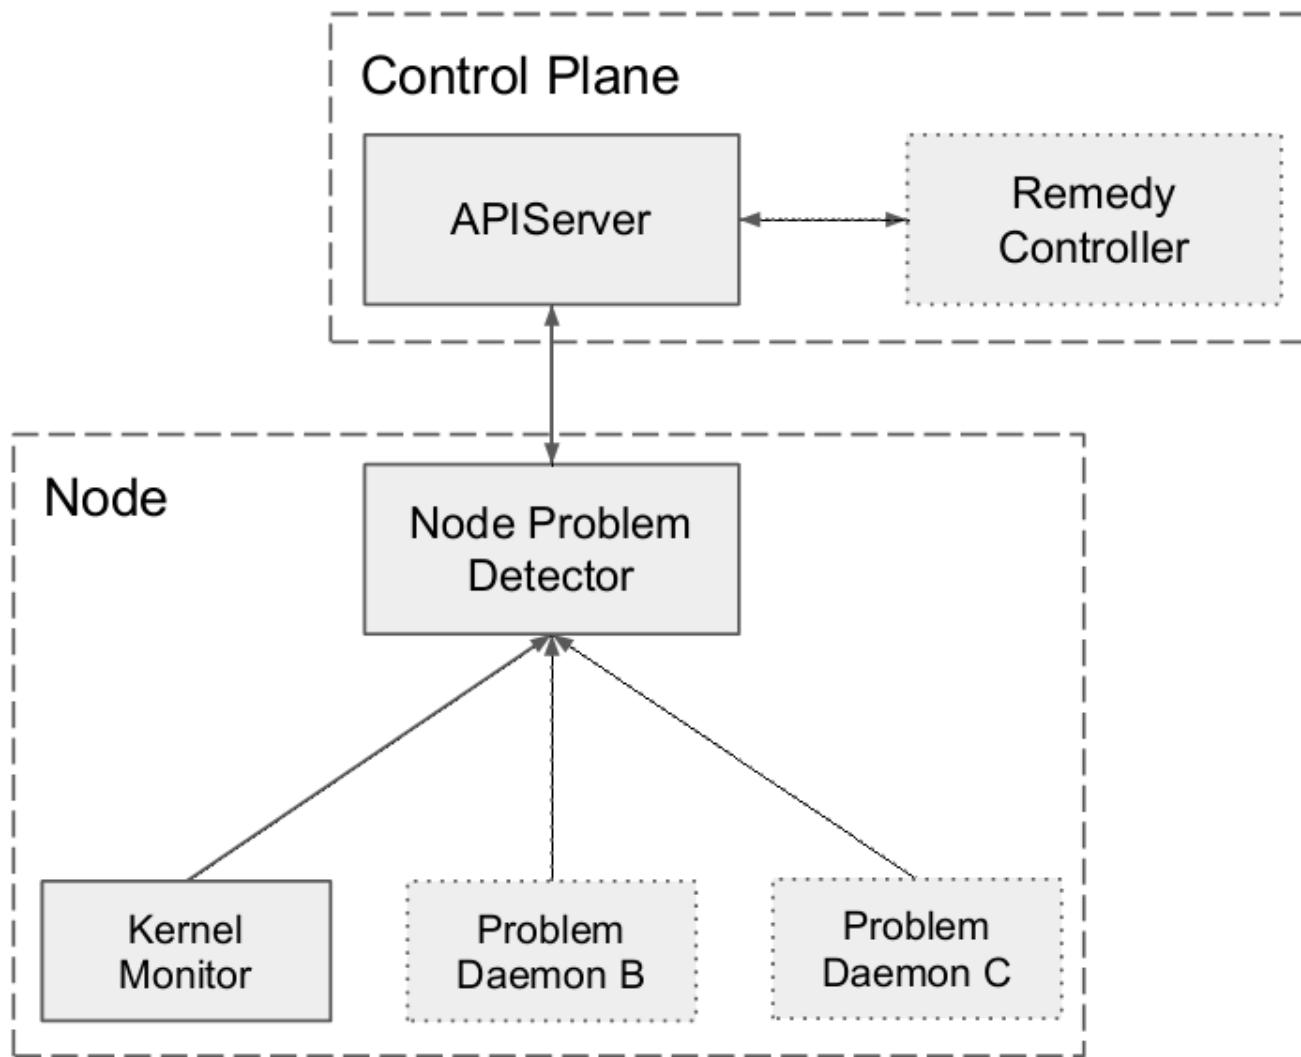
Training Pipeline



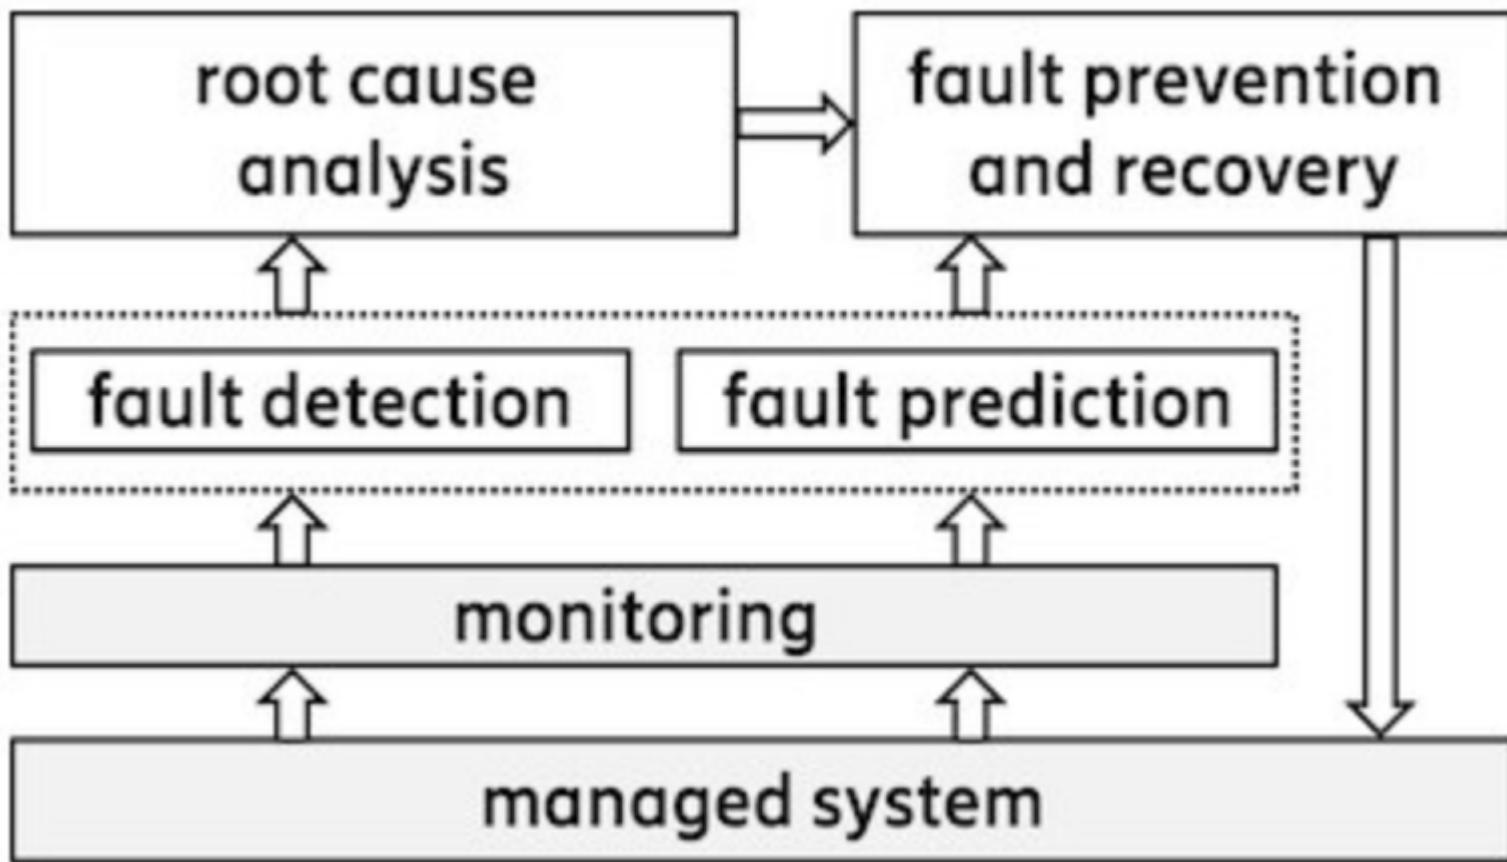


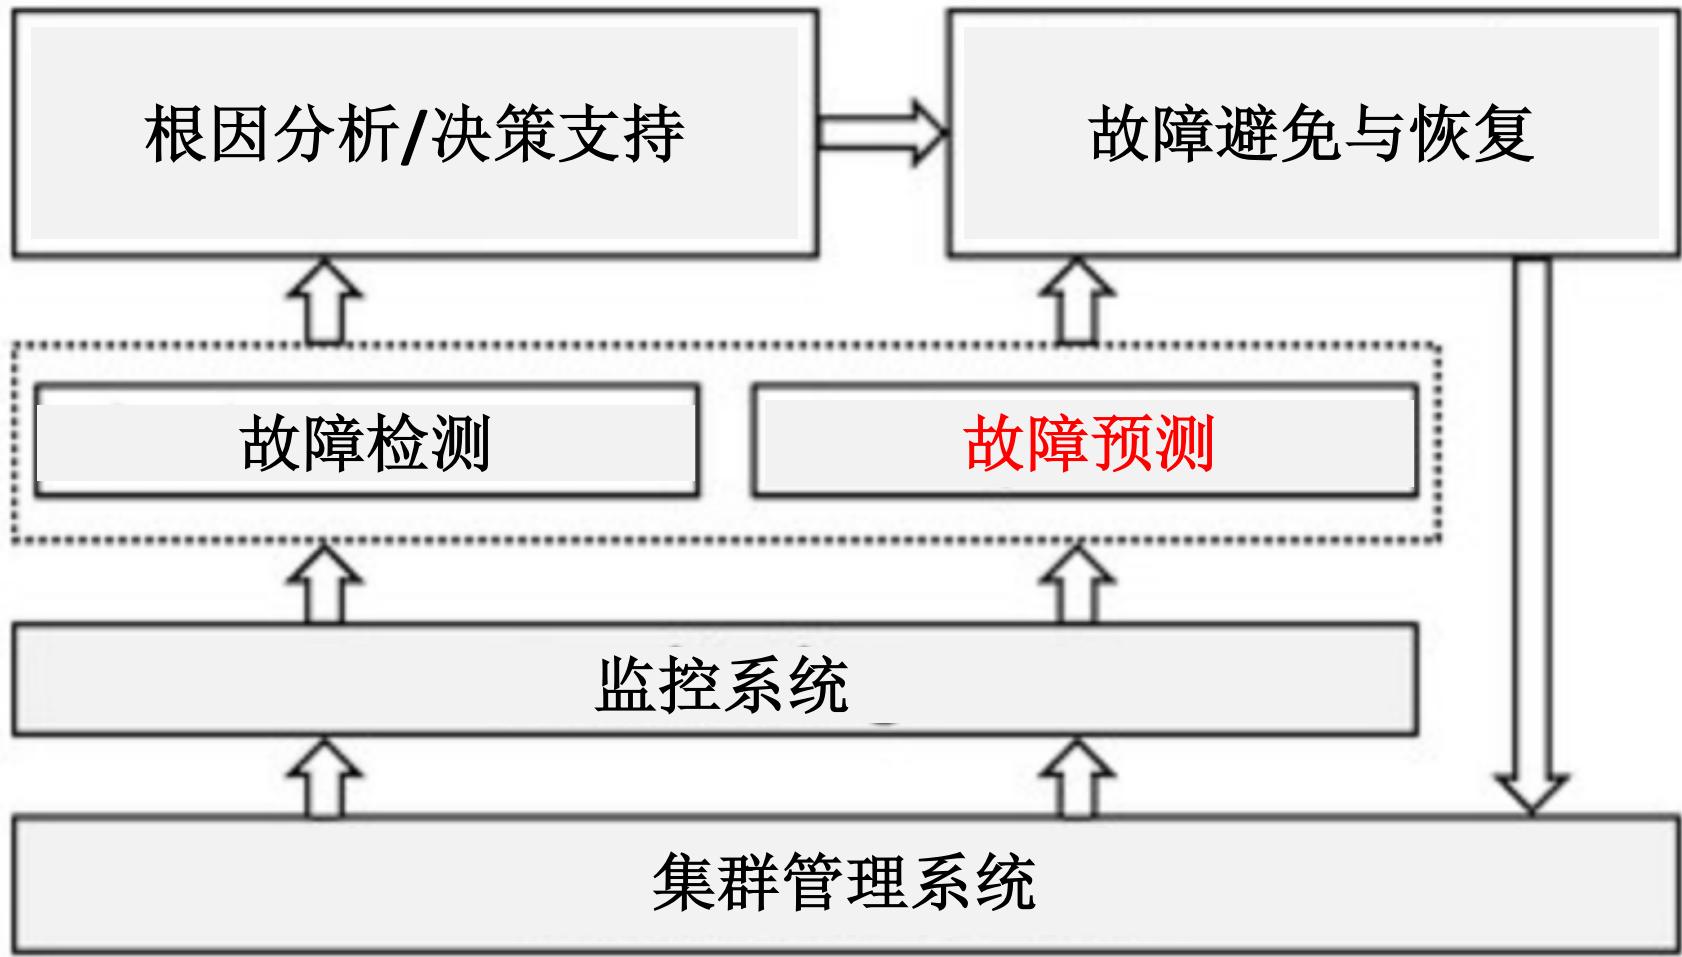
Remedy System

- Problem Manager
 - Problem Daemon
- Metric Manager



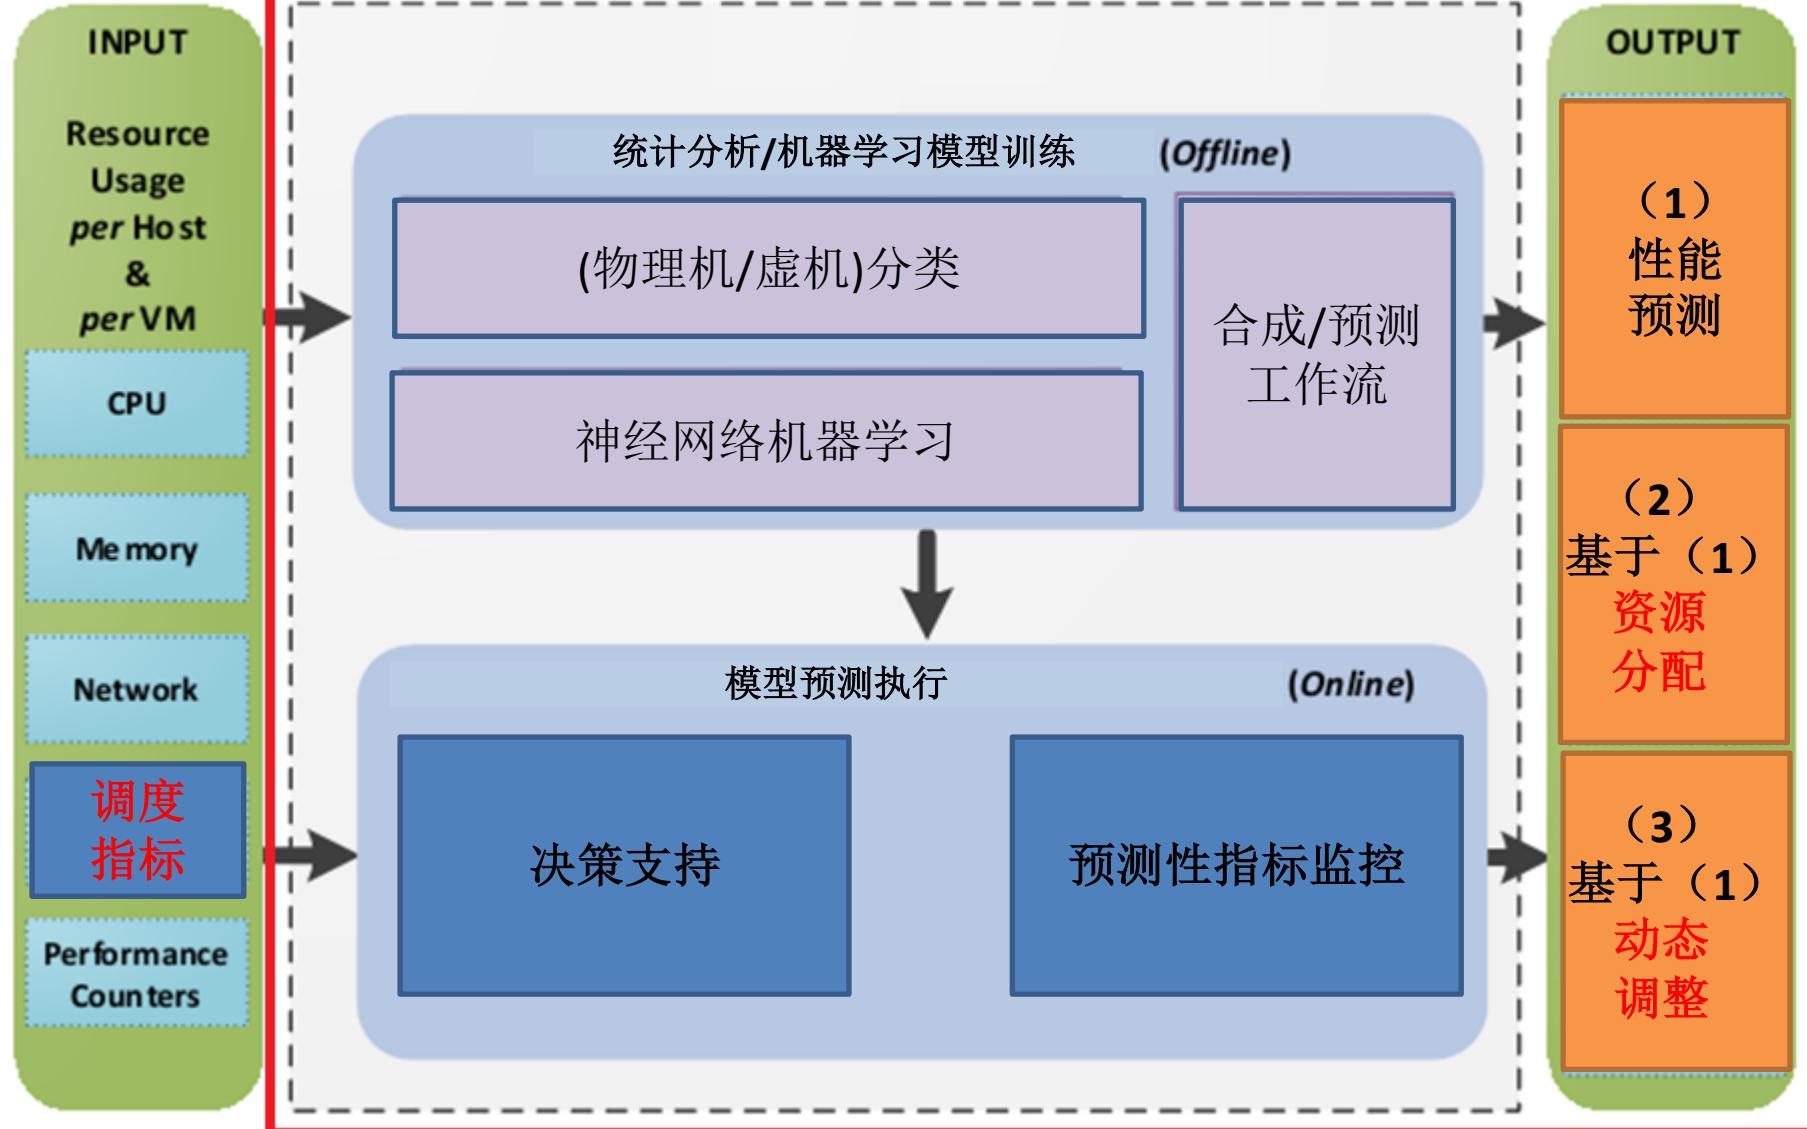
Fault Management System

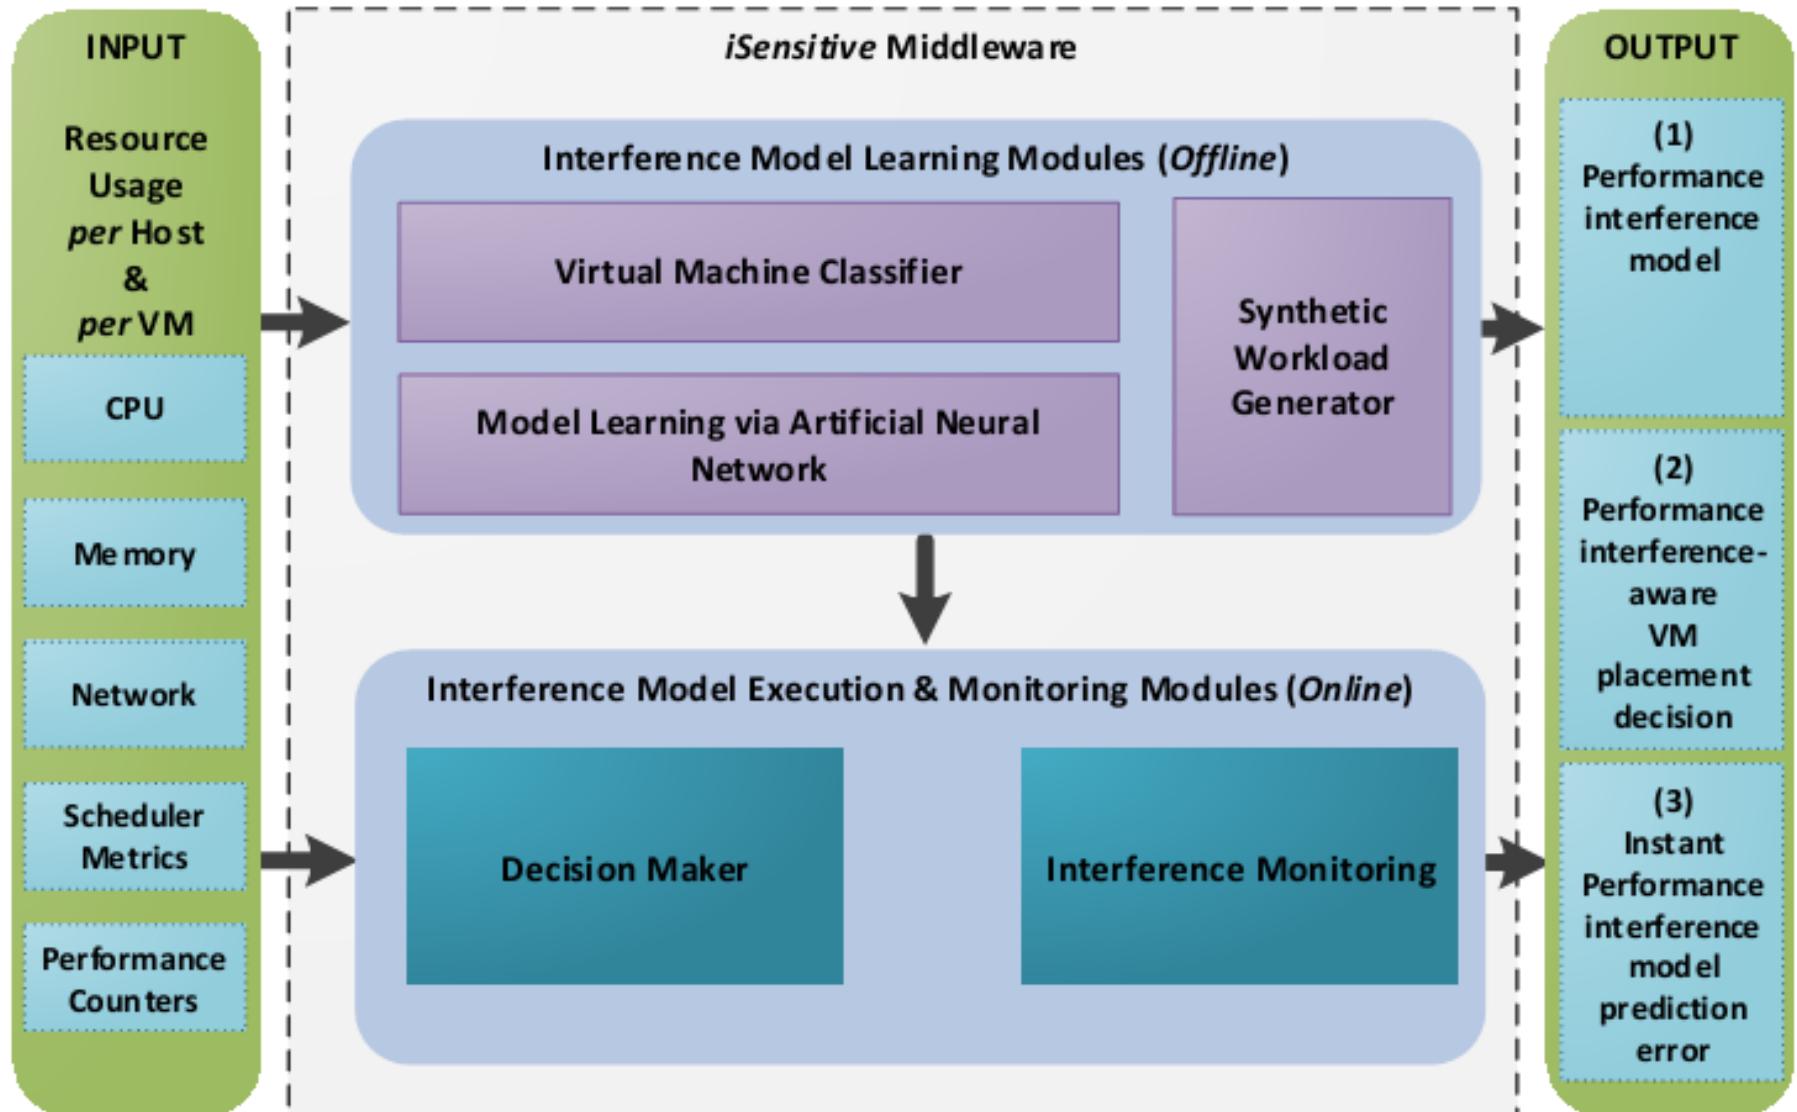


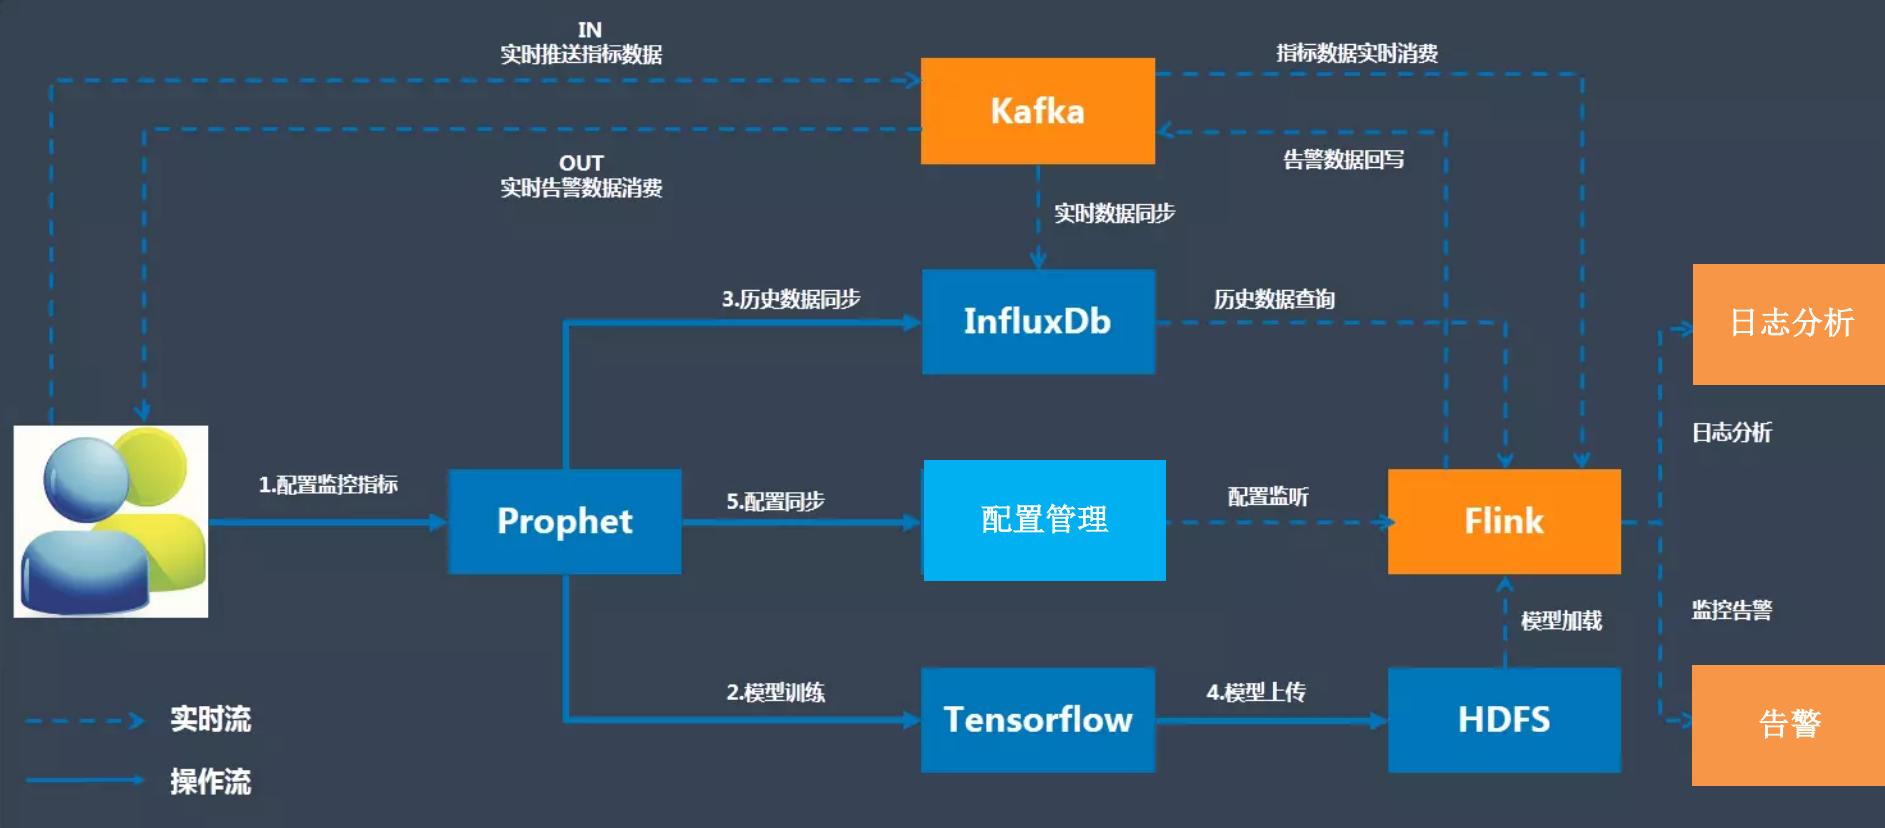


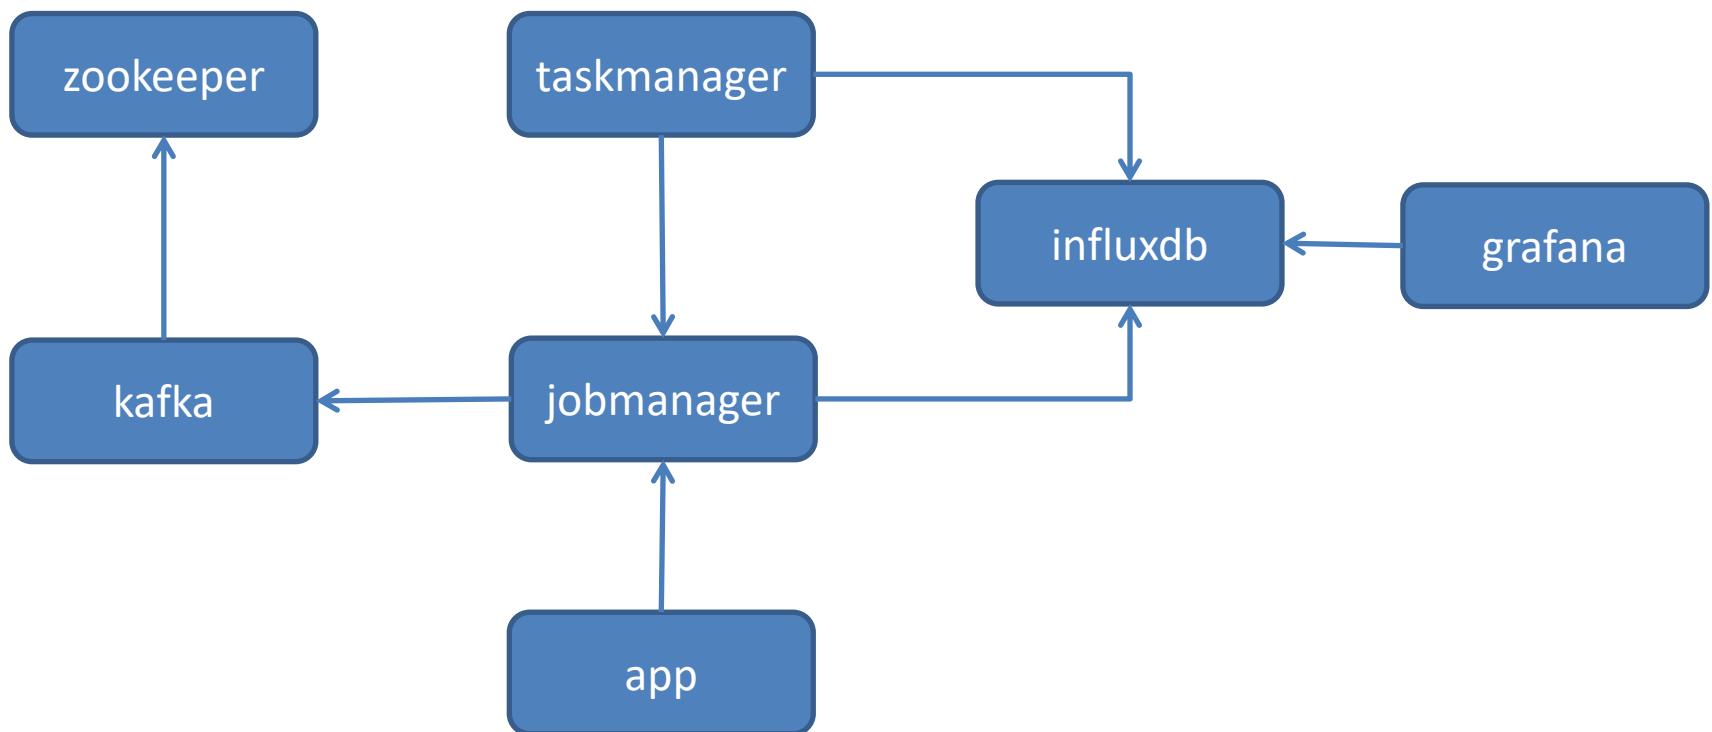
Resource Usage Prediction

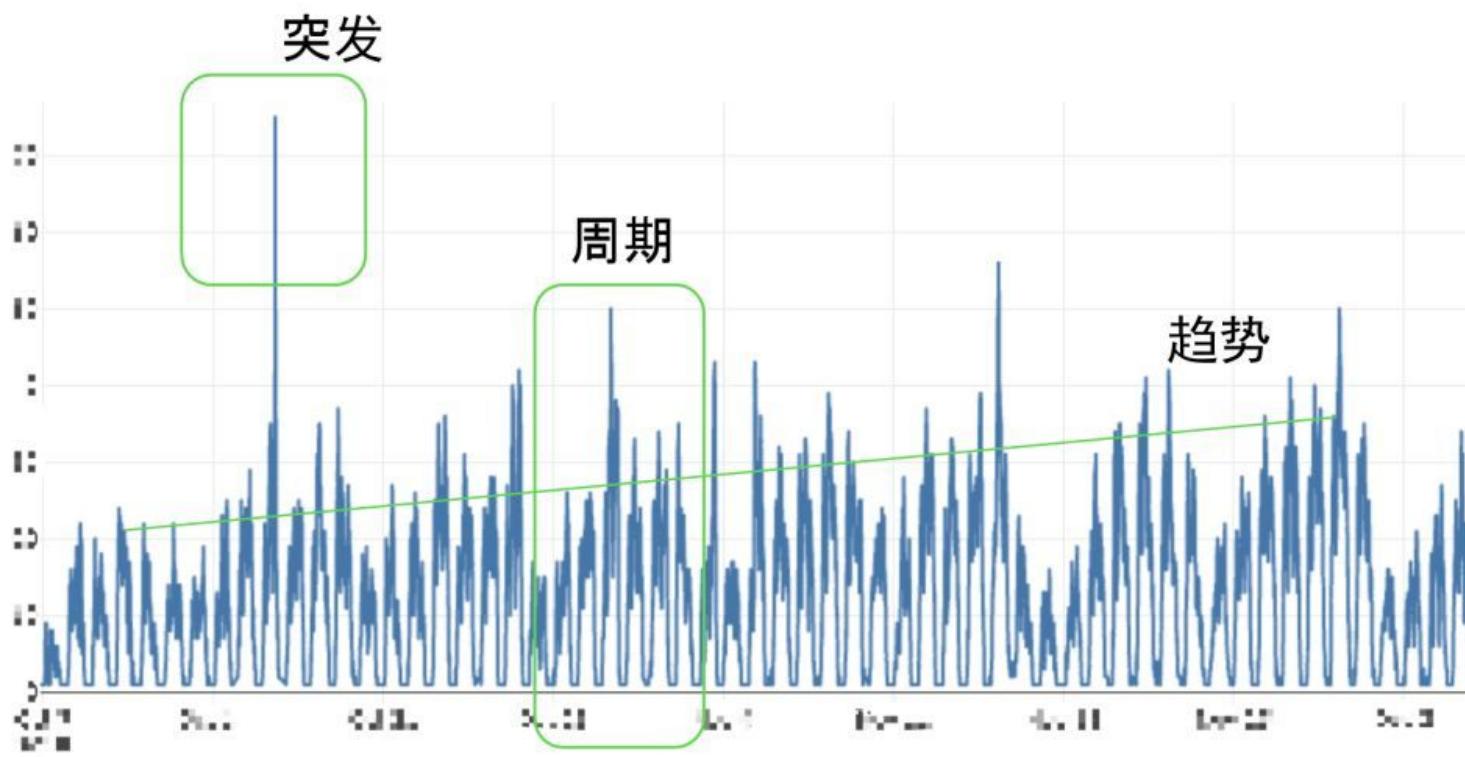
- Two layers, feed forward ANN
 - Cpu usage ratio, Memory usage ratio
 - Run vms
 - Ultility model

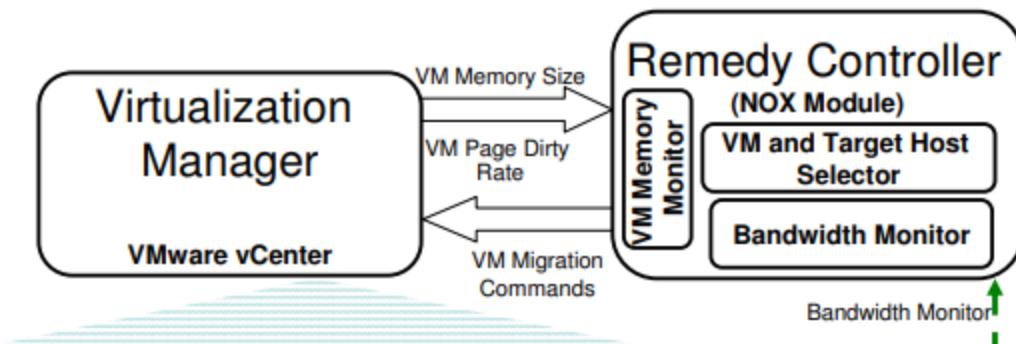


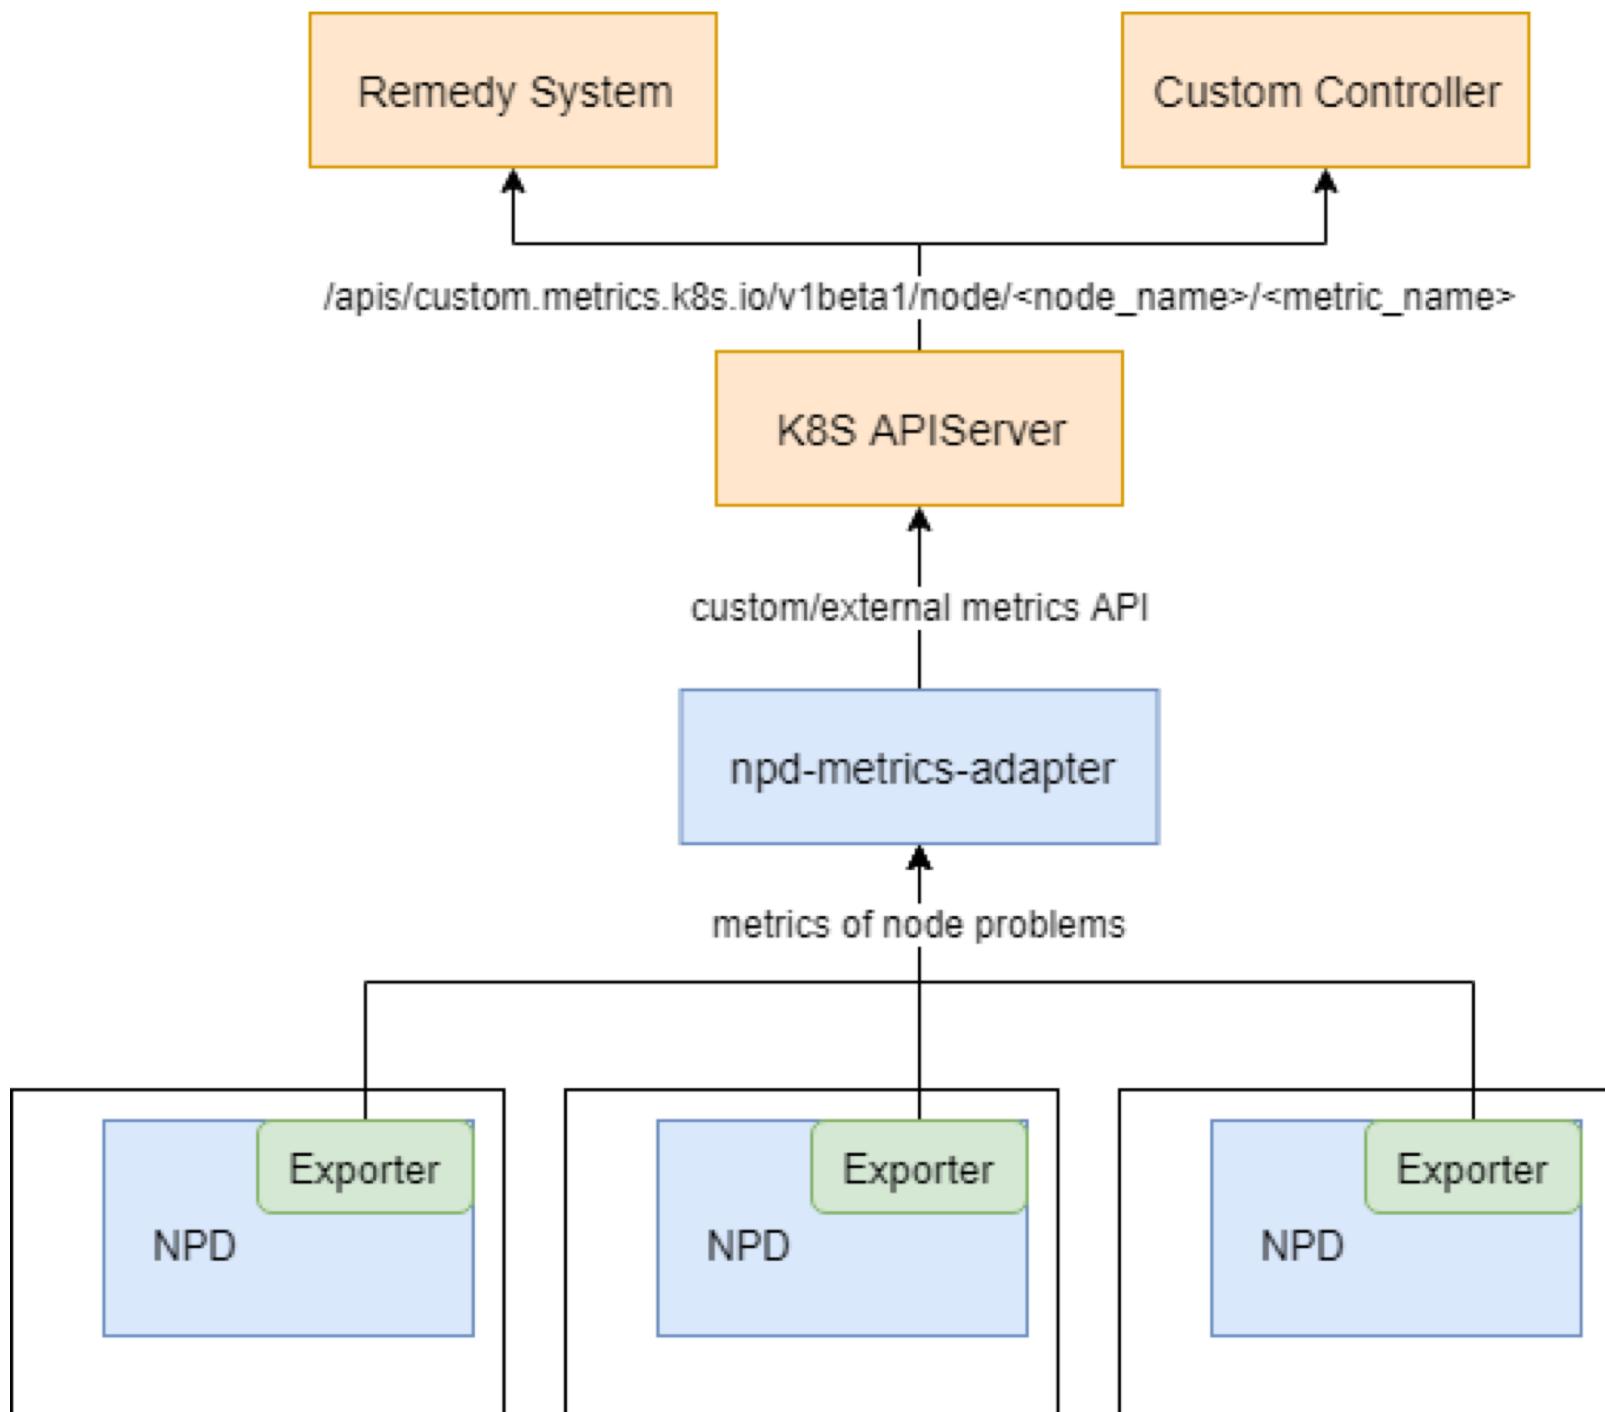












Task分析

Case	Py	Input	Output
task_uasge_500_preprocess	data_preprocess.py	task_usage_500_title.csv	test.csv
	data_inverse.py	test_data.csv	test_data_inversed_bycode.csv
	data_GRU_run.py	dataset_inversed.csv	plot
	google_cluster_dataset_500_predict/ google_cluster_dataset_500_predict.py	test_data_inversed.csv	plot
task_usage_total_preprocess	preprocess/with_maid/ data_predict.py	task_usage_preprocessed.csv	plot
	preprocess/with_maid/ data_preprocessing.py	task_usage.csv	task_usage_preprocessed_ud.csv
	preprocess/with_maid/ dataset_drop_mean.py	task_usage_preprocessed_ud.csv	task_usage_preprocessed.csv
	preprocess/Without_maid(ignore)/ data_total_preprocessing.py	task_usage.csv	task_usage_total_inversed_bycode.csv
	optimizer/Bidirectional_Swish/ bidirectional_swish.py	task_usage_preprocessed.csv	plot
	optimizer/Swish_activate/ swish_activate.py	task_usage_preprocessed.csv	plot

Trace分析

Case	Py	Input	Output
cluster	Alibaba_cluster/ Alibaba_cluster_LSTM_predict.ipynb	Machine_usage_groupby.csv	Plot(failed)
	Alibaba_cluster/ Alibaba_cluster_GRU_predict.ipynb	Machine_usage_groupby.csv	plot
	Alibaba_cluster/ Alibaba_cluster_drop_predict.ipynb	Machine_usage_groupby.csv	plot
cluster_predict	Alibaba_cluster_predict/code/ preprocess.py	machine_usage_50000.csv	Machine_usage_fill.csv
	Alibaba_cluster_predict/code/ Averange_time_process.py	Machine_usage_fill.csv	Machine_usage_groupby.csv
	Alibaba_cluster_predict/code/ data_run.py	Machine_usage_groupby.csv	plot & RMSE
	Alibaba_cluster_predict/code/ Plot_cluster_data.py	Machine_usage_groupby.csv	plot
cluster_predict_Total	cluster_predict_Total/ Alibaba_cluster_predict_runafterpr ocession.py	Machine_usage_groupby.csv	
	cluster_predict_Total/ data_give_head.py	machine_usage.csv	machine_usage_headed.csv
	cluster_predict_Total/ data_test_head.py	Machine_usage_groupby_100.csv	nil
	cluster_predict_Total/ data_preprocess.py	machine_usage_headed.csv	Machine_usage_groupby.csv

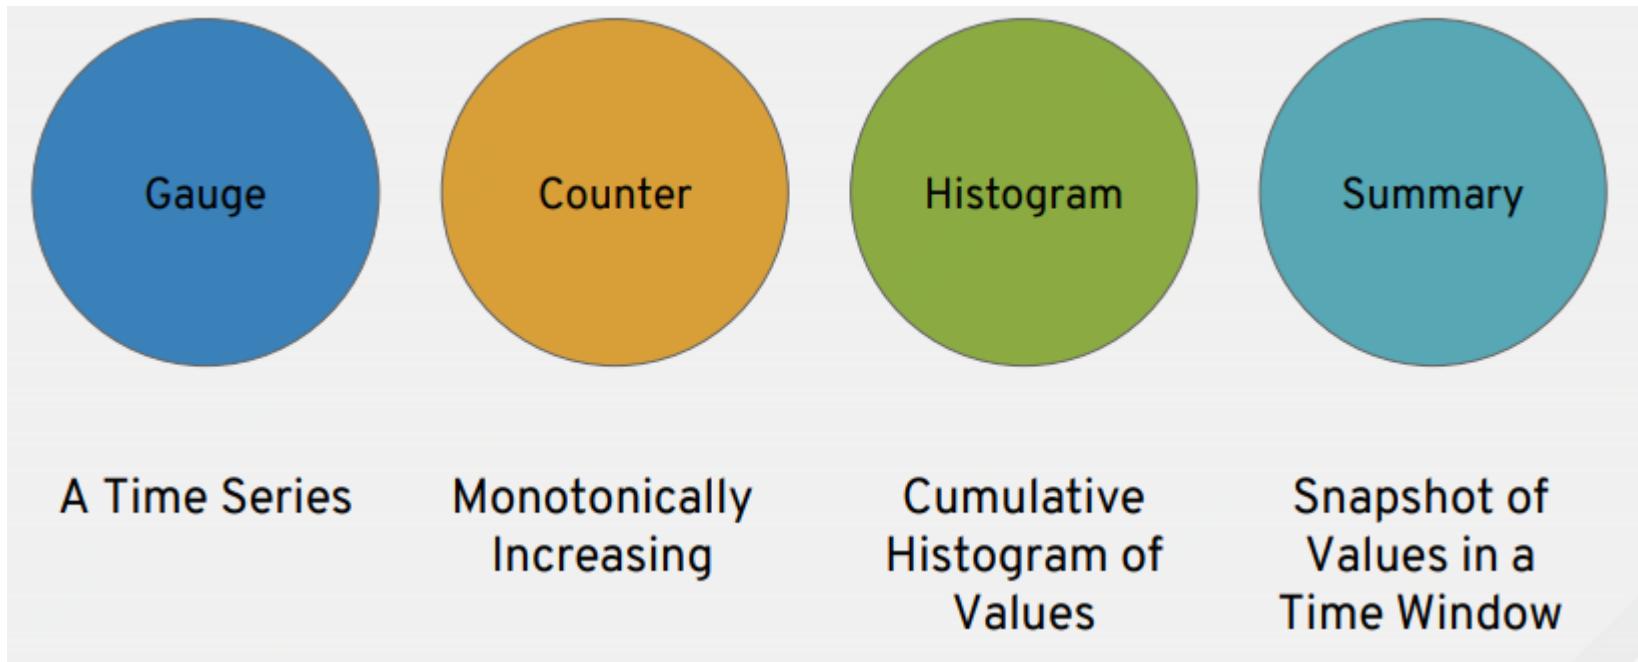
Trace分析

Case	Py	Input	Output
cluster_predict_Total	cluster_predict_Total/ prediction_model.py	Machine_usage_groupby.csv	plot & RMSE & MSE
	cluster_predict_Total/ test_mse_predict.py	Machine_usage_groupby.csv	plot & RMSE & MSE (error)
cluster_predict_compare	Train_40000/COV+GRU+GRU+D+D+DROP+E50/ COV+GRU+GRU+D+D+DROP+E50.py	Machine_usage_groupby.csv	Plot & RMSE & MSE
	Train_40000/Bidirectional_Swish/ bidirectional_swish.py	Machine_usage_groupby.csv	Plot & RMSE & MSE
	Train_20000/COV+GRU+GRU+D+D+E50/ COV+GRU+GRU+D+D+E50.py	Machine_usage_groupby.csv	Plot & RMSE & MSE
	Train_20000/COV+LSTM+LSTM+D+D+DROP+E50/ alibaba_lstm_predict.py	Machine_usage_groupby.csv	Plot & RMSE & MSE
	Train_20000/GRU+D+D+E50/ GRU+D+D+E50.py	Machine_usage_groupby.csv	Plot & RMSE & MSE
	Train_20000/GRU+D+E50/ GRU+D+E50.py	Machine_usage_groupby.csv	Plot & RMSE & MSE
	Train_20000/GRU+GRU+D+D+DROP+E50/ GRU+GRU+D+D+DROP+E50.py	Machine_usage_groupby.csv	Plot & RMSE & MSE
	Train_20000/GRU+GRU+D+D+E50/ GRU+GRU+D+D+E50.py	Machine_usage_groupby.csv	Plot & RMSE & MSE

Remedy-controller

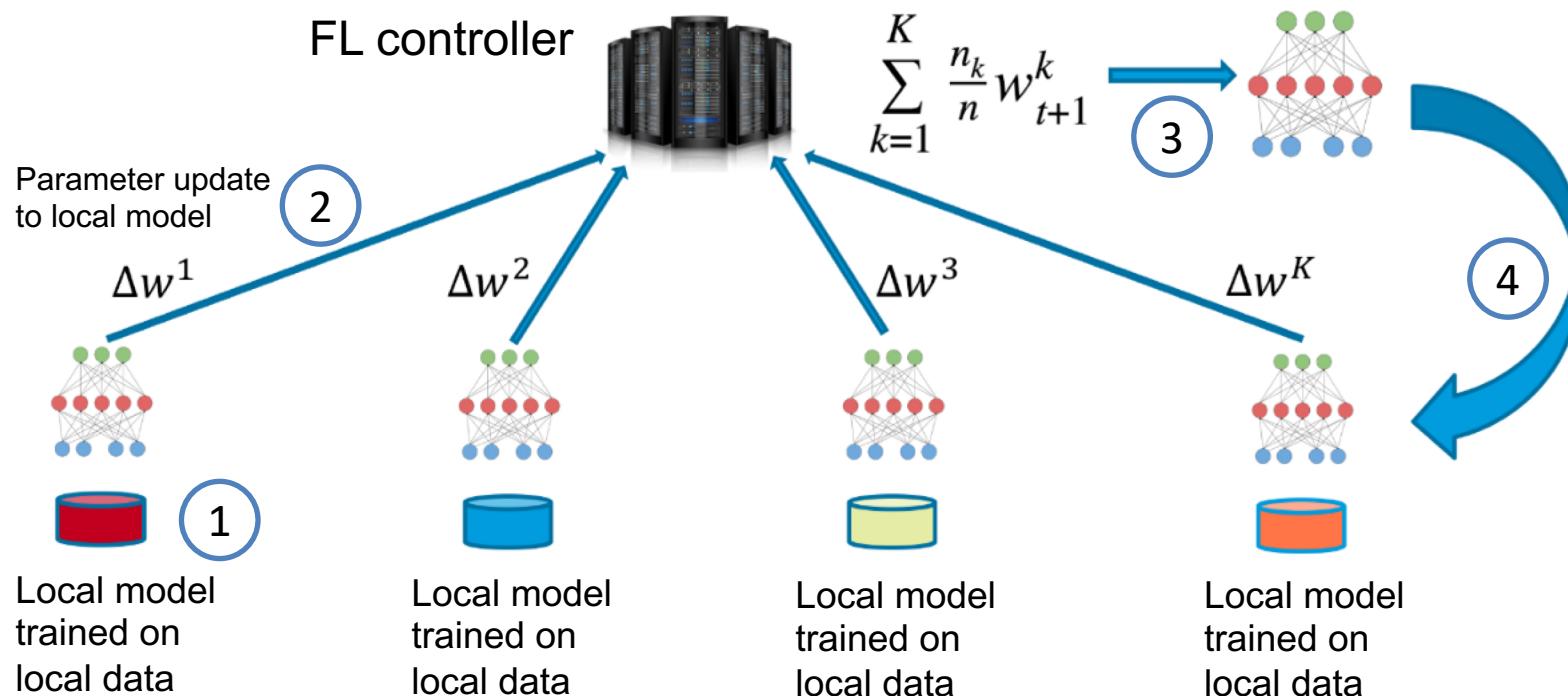
- [[draino](#)]
- [[descheduler](#)]
- [[debug-application-cluster](#)]

Prometheus Metric Types



- Anomaly detector
 - <https://github.com/AICoE/prometheus-anomaly-detector>

Federated Learning

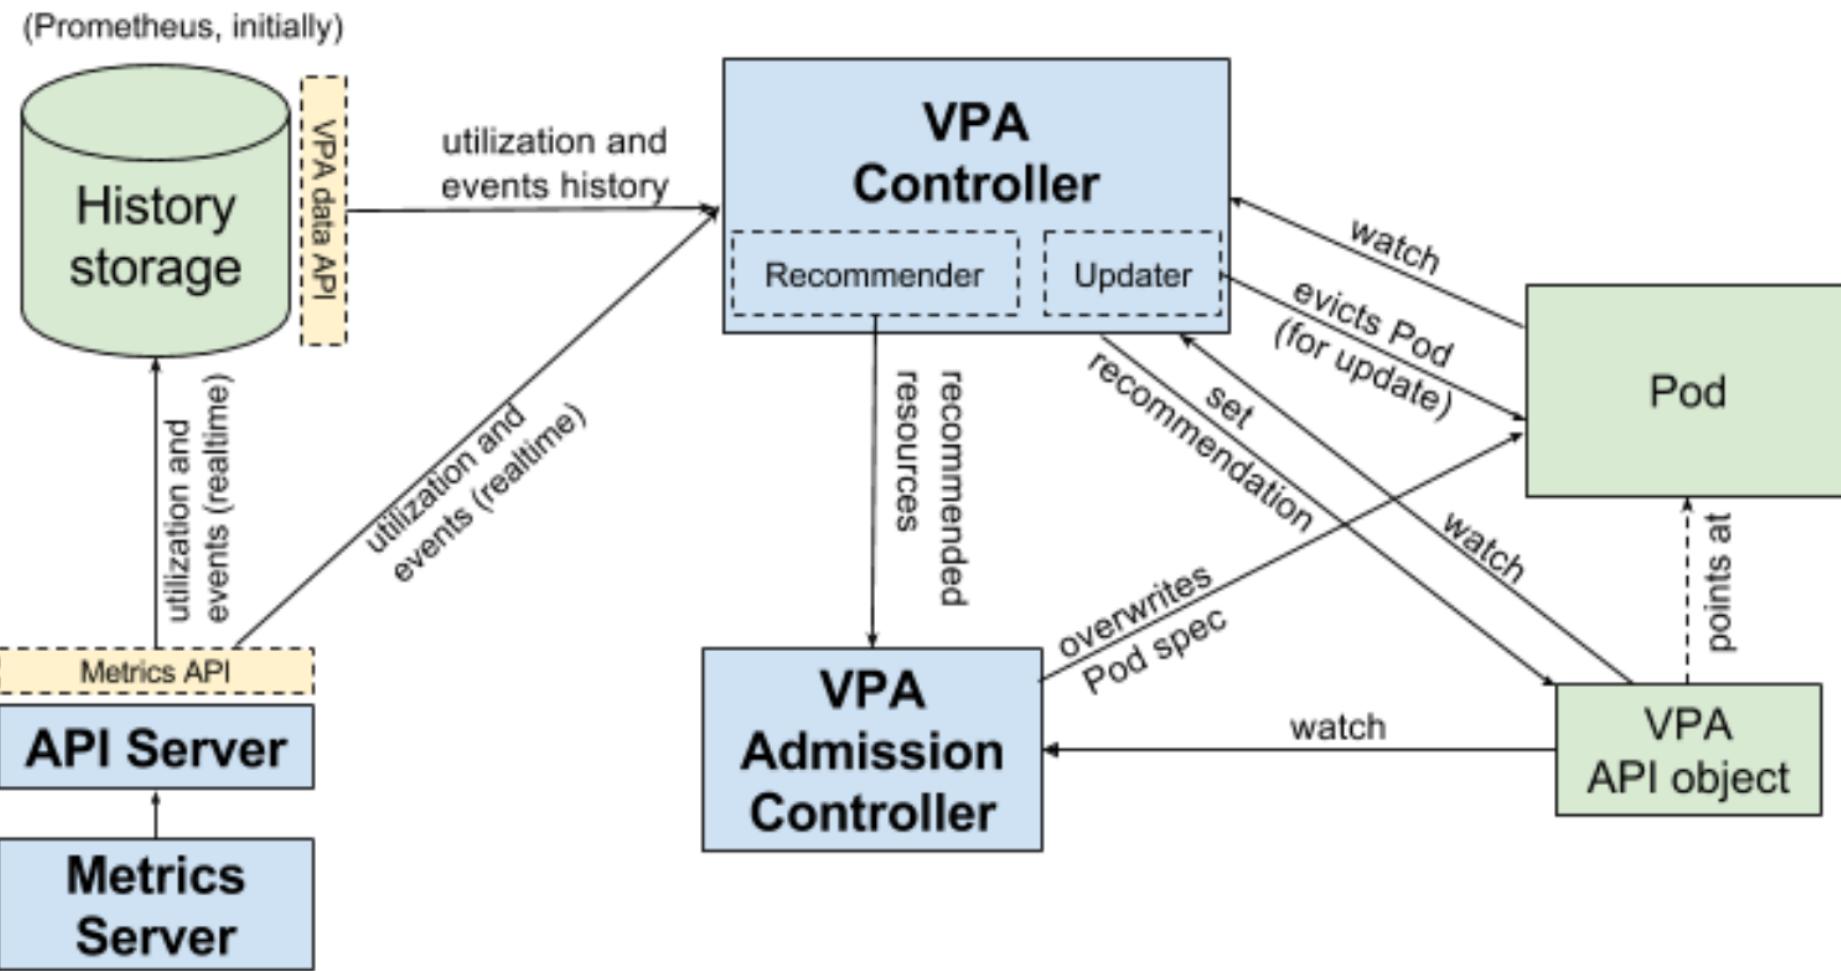


Source: https://miro.medium.com/max/1400/1*HaH611vAy2eB1e42vz3X4g.png

最优建模

- COV+GRU+GRU+D+D++DROP+E50

VPA(Autopilot)



ML platform

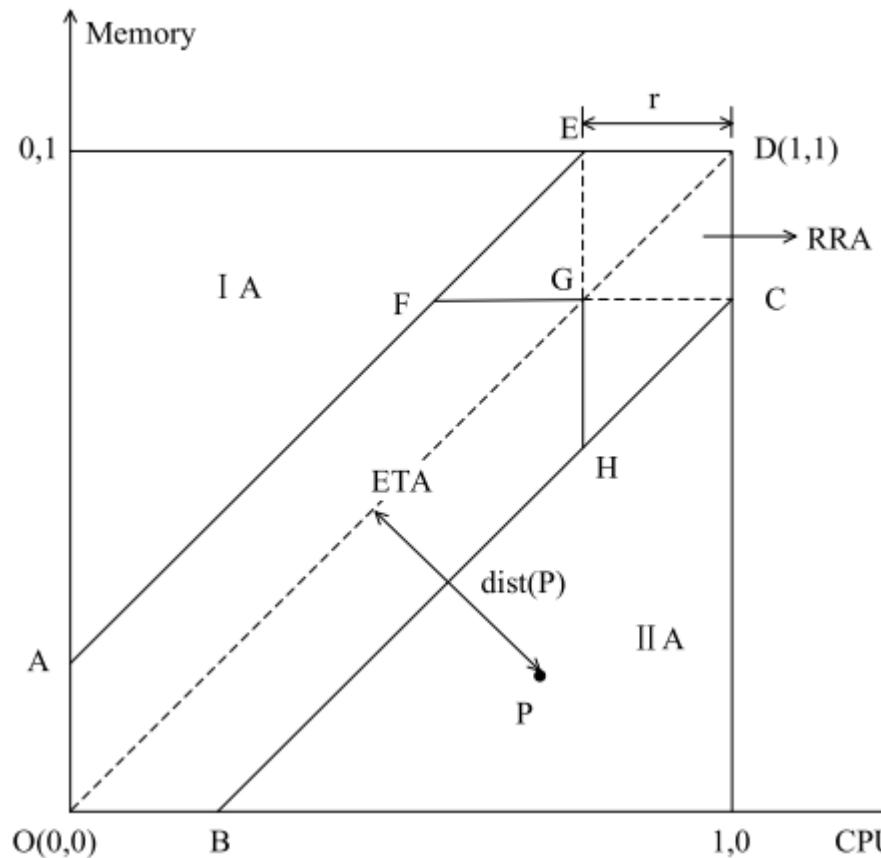
- maro on k8s

Defragmenting the Cloud

- 资源碎片化建模
 - 有CPU无MEM
 - 无CPU有MEM

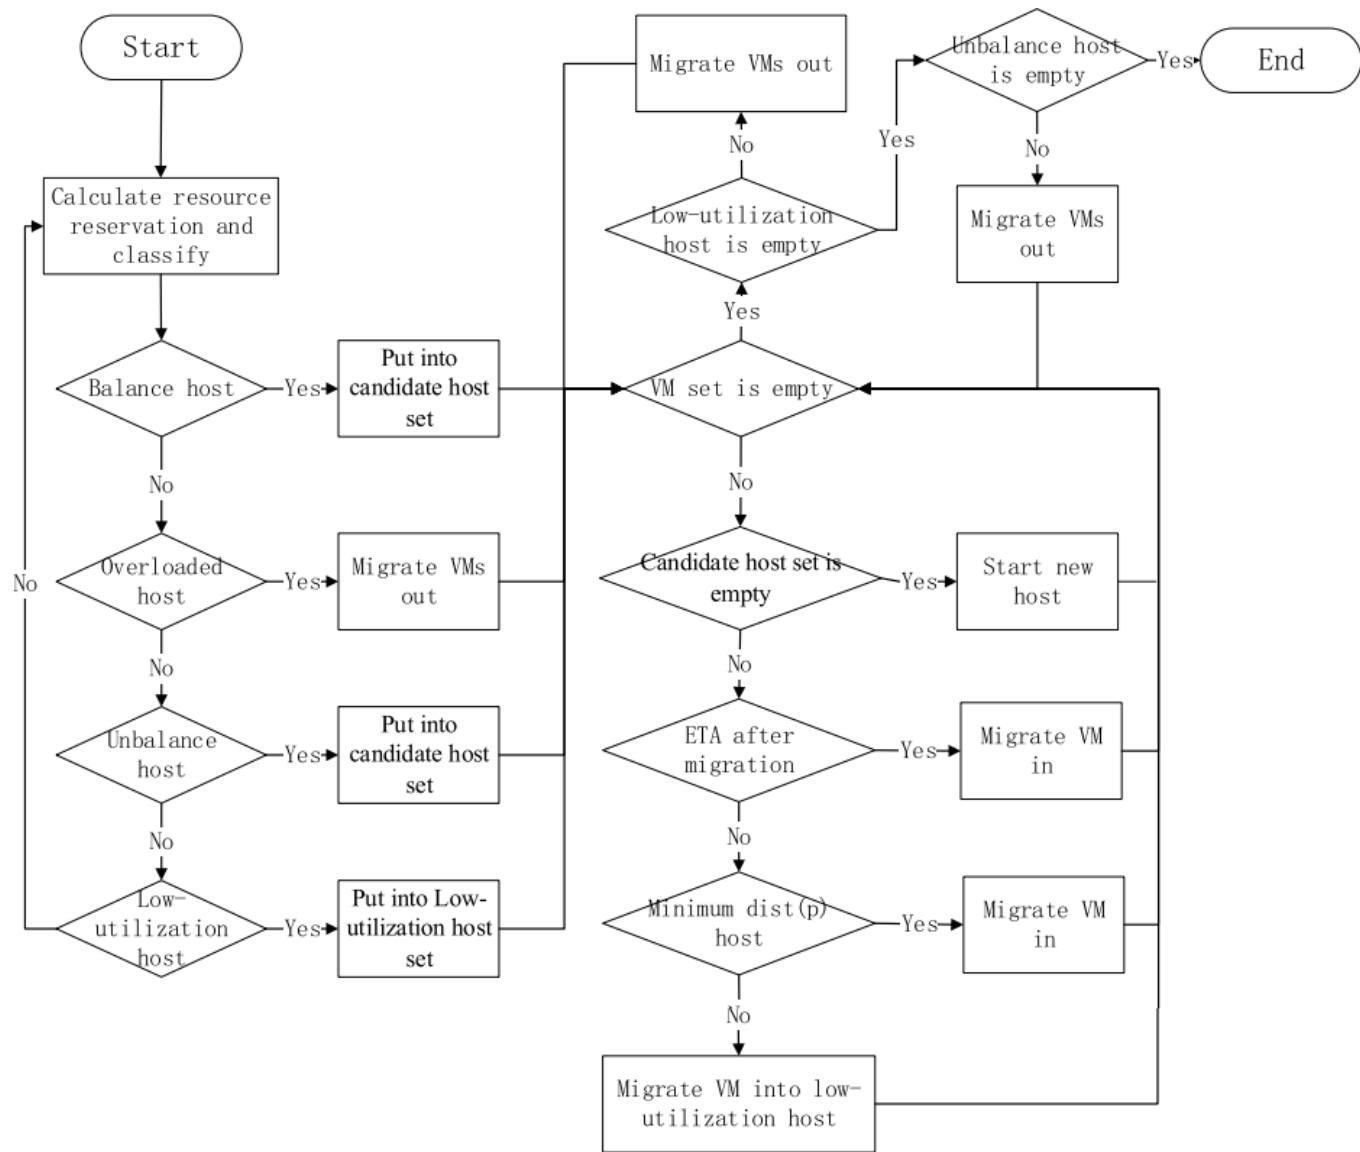
利用率描述

- RRF index = $\min\{(\text{空白}) / (\text{TOTAL})\}$



$$f(x, y) = \begin{cases} RRA, & (1-x) \leq r \& (1-y) \leq r \\ ETA, & dist(p) \leq \frac{\sqrt{2}}{2}r \\ IA \text{ or } IIA, & \text{otherwise} \end{cases}$$

动态迁移流程



rms all	Not.ready	Ready	optimes
23628	54*	23574	11-16

jcs独部	jks-jcs混部	jks独部	baremetal	网络不通	操作合计	RMS节点总数
8087	15402	4	78	3	23574	23574

*

SUCC,10.217.174.172

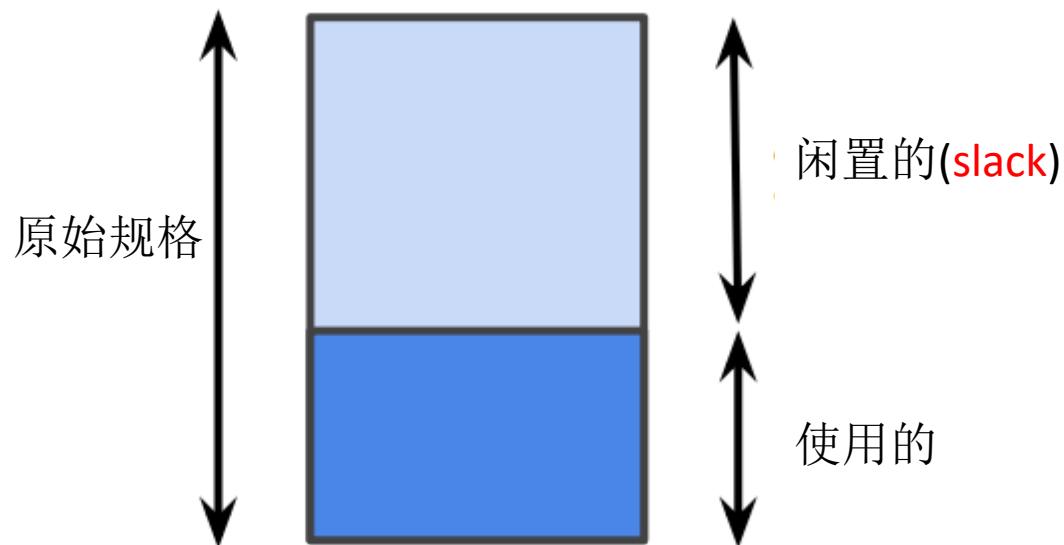
SUCC,10.217.35.215

SUCC, 10.217.94.186

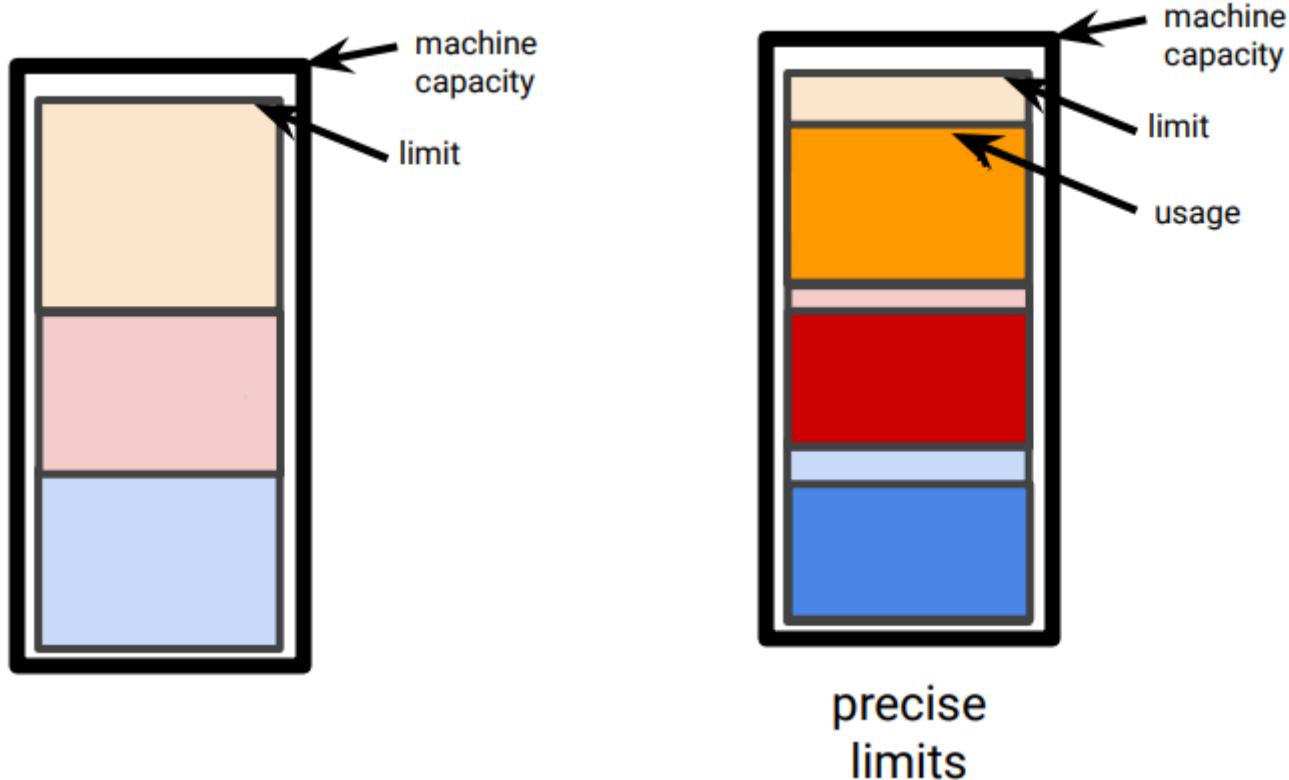
Limit预测

- 原理
- 场景
- 实验
- 测评

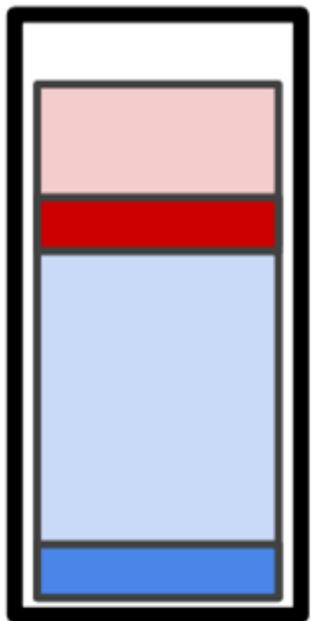
问题背景



分析

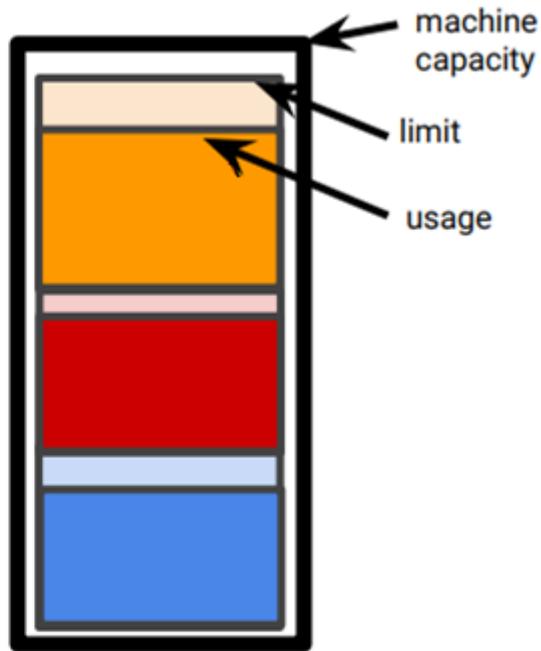


面对场景

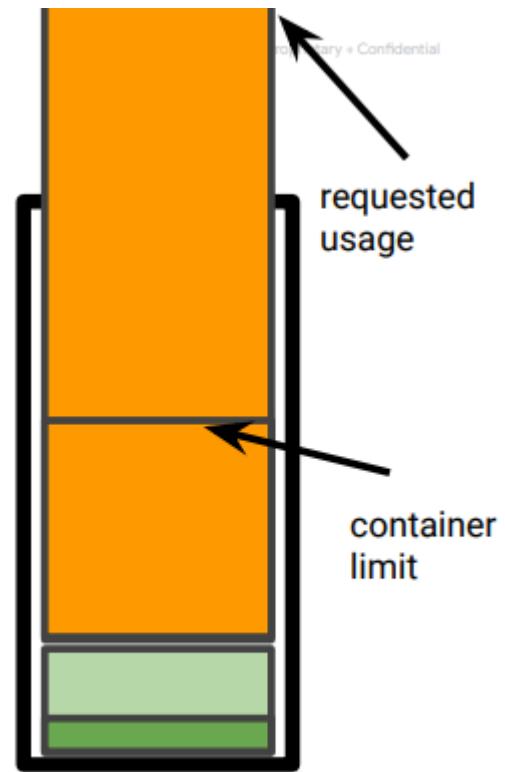


resource are
wasted
(underallocated
machine)
bad!

资源浪费



precise
limits
good!



还是**bad!** 资源超了

实际要点

- 如何准确的预测资源水位？
- 一旦预测不准确，就会：
 - 影响 Qos，CPU 可能还好
 - 内存如果预测不准，会触发 OOM
 - 这种行为不可接受

解决办法

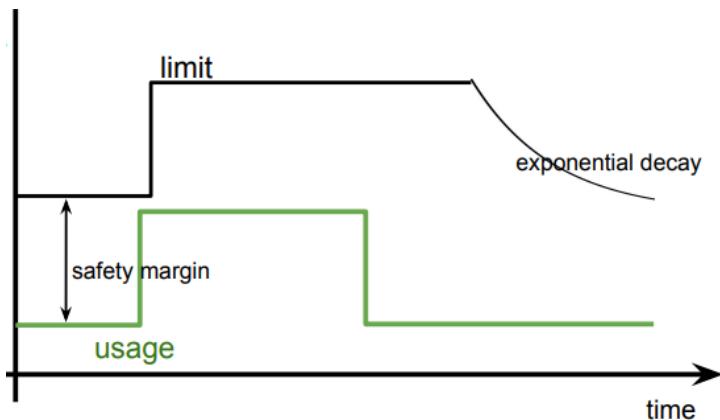
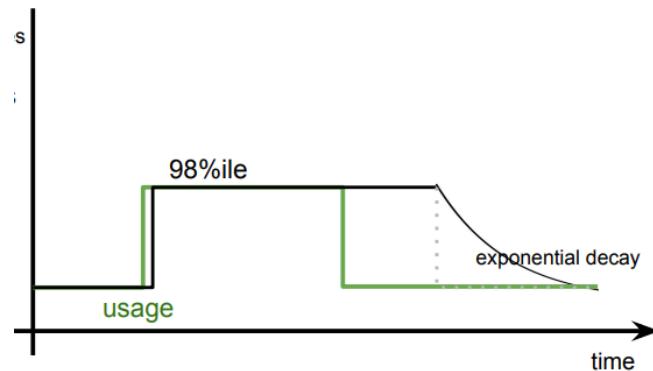
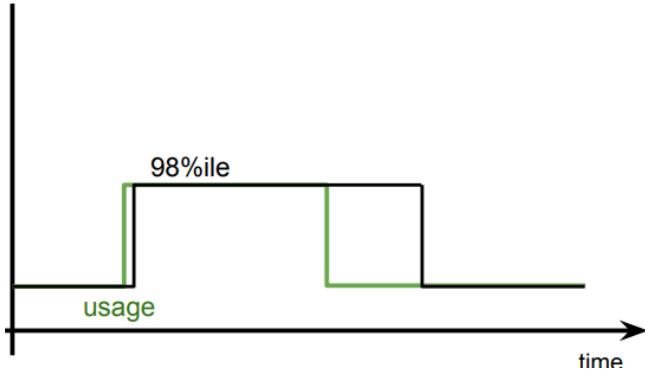
- 滑动窗口预测
 - 滑动窗口算法的作用就是根据一个窗口期内的历史数据，加权预测资源的合理limits
- 增强学习预测（辅助，未实现）
 - 提高预测的准确率
 - 引入先验

滑动窗口

- 指数衰减
 - 历史数据的权重，越久的数据权重越低
 - 半衰期： CPU=?h, MEM=?h
- 峰值预测=>Upper Bound
 - 取最近N个样本的最大值
- 均值预测=>Trend
 - 历史数据的加权平均
- J-%ile 分位值预测 => Lower Bound
 - 所有历史数据的j-%ile分位加权平均

概念实现

- 指数衰减采样 (半衰期)
- 样本统计计算 (95%ile)
- 安全边界

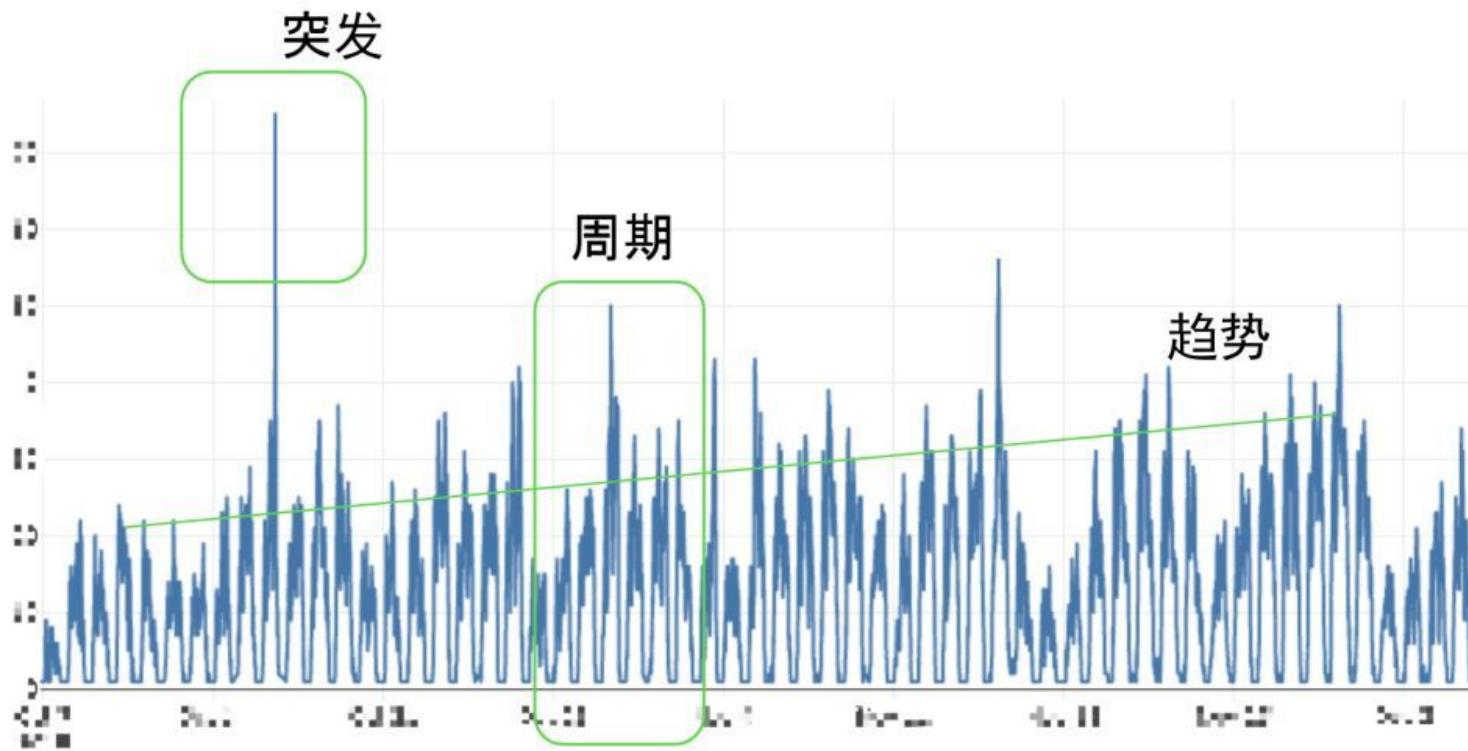


实现

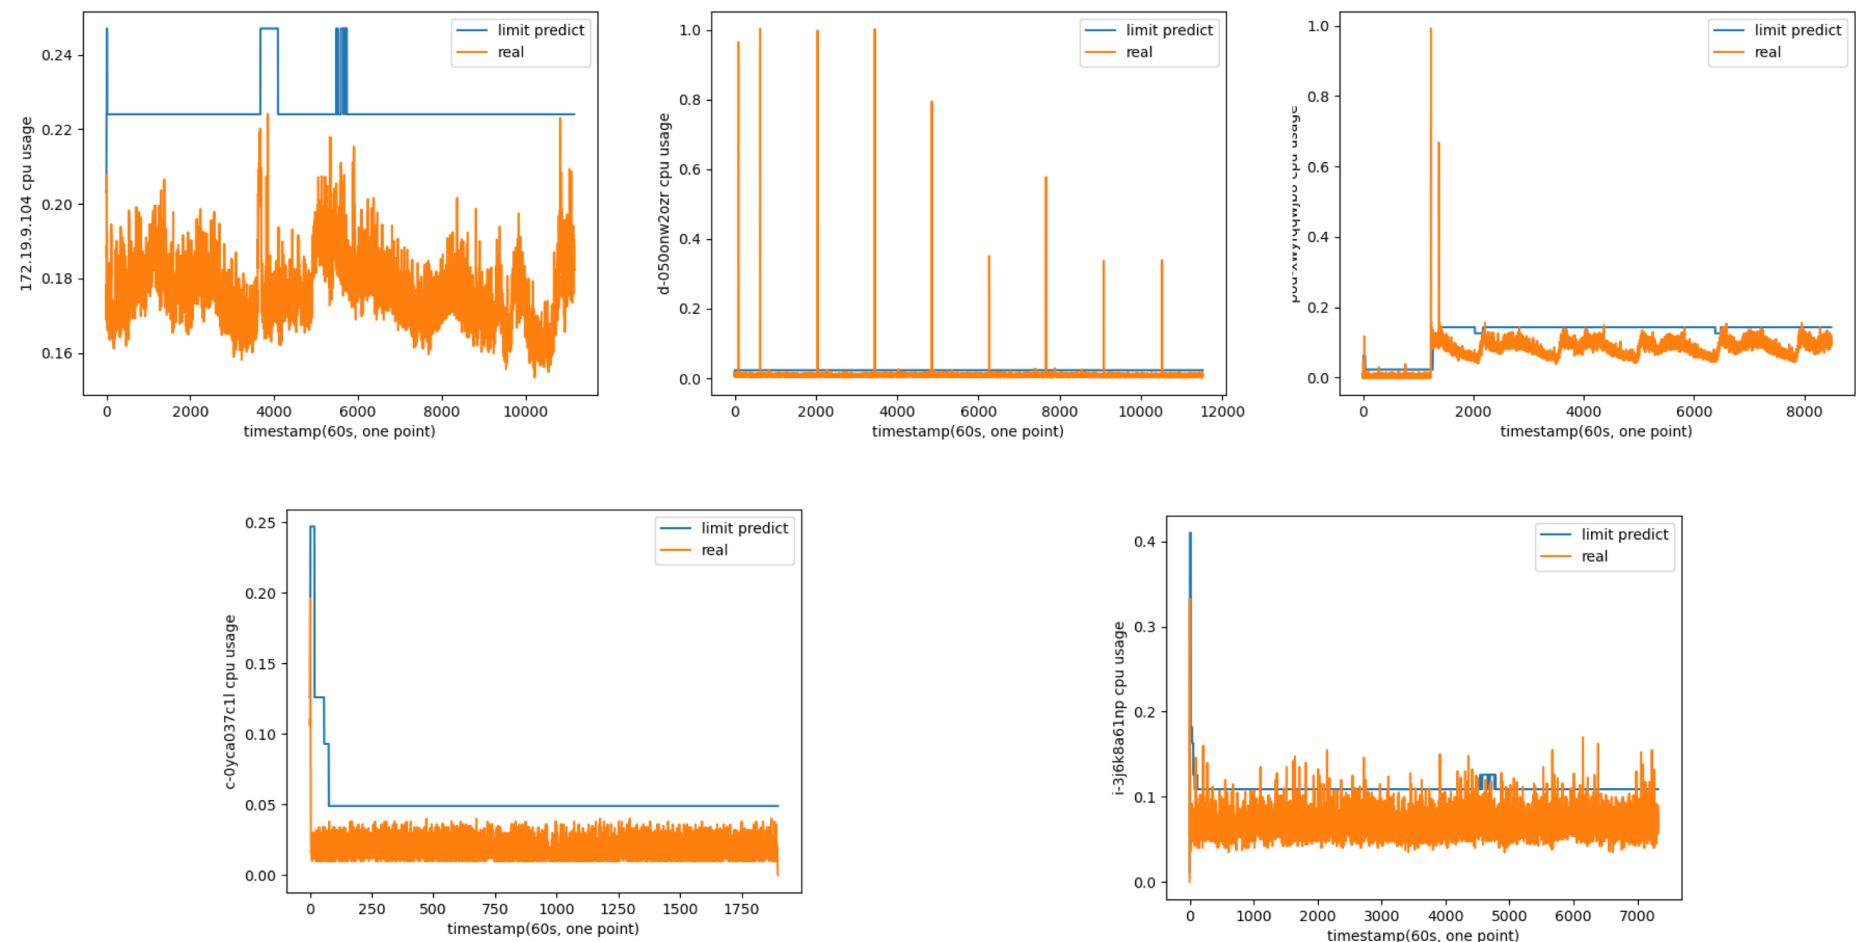
- 基于cli的模式
- 离线处理数据

```
[root@A04-R08-I137-I205-9113C72 cloud-prophet]# ./autopilot --help
Usage of ./autopilot:
-cpu-histogram-decay-half-life duration
    CPU利用率权重减半的周期, 半衰期. (default 24h0m0s)
-csv-file string
    指标数据来自监控系统, csv文件名称 (default "unkowfile")
-element-id string
    资源利用率预测的实例ID (default "unknow")
-recommendation-margin-fraction float
    预测的安全边缘余量, 百分比 (default 0.15)
-runonce-timeout duration
    一次运行的最大超时时间 (default 1h0m0s)
-sample-second-interval int
    样本的采样间隔, 单位(秒), 整型 (default 60)
-target-cpu-percentile float
    cpu预测采用的分位值, 百分比 (default 0.9)
```

数据特征

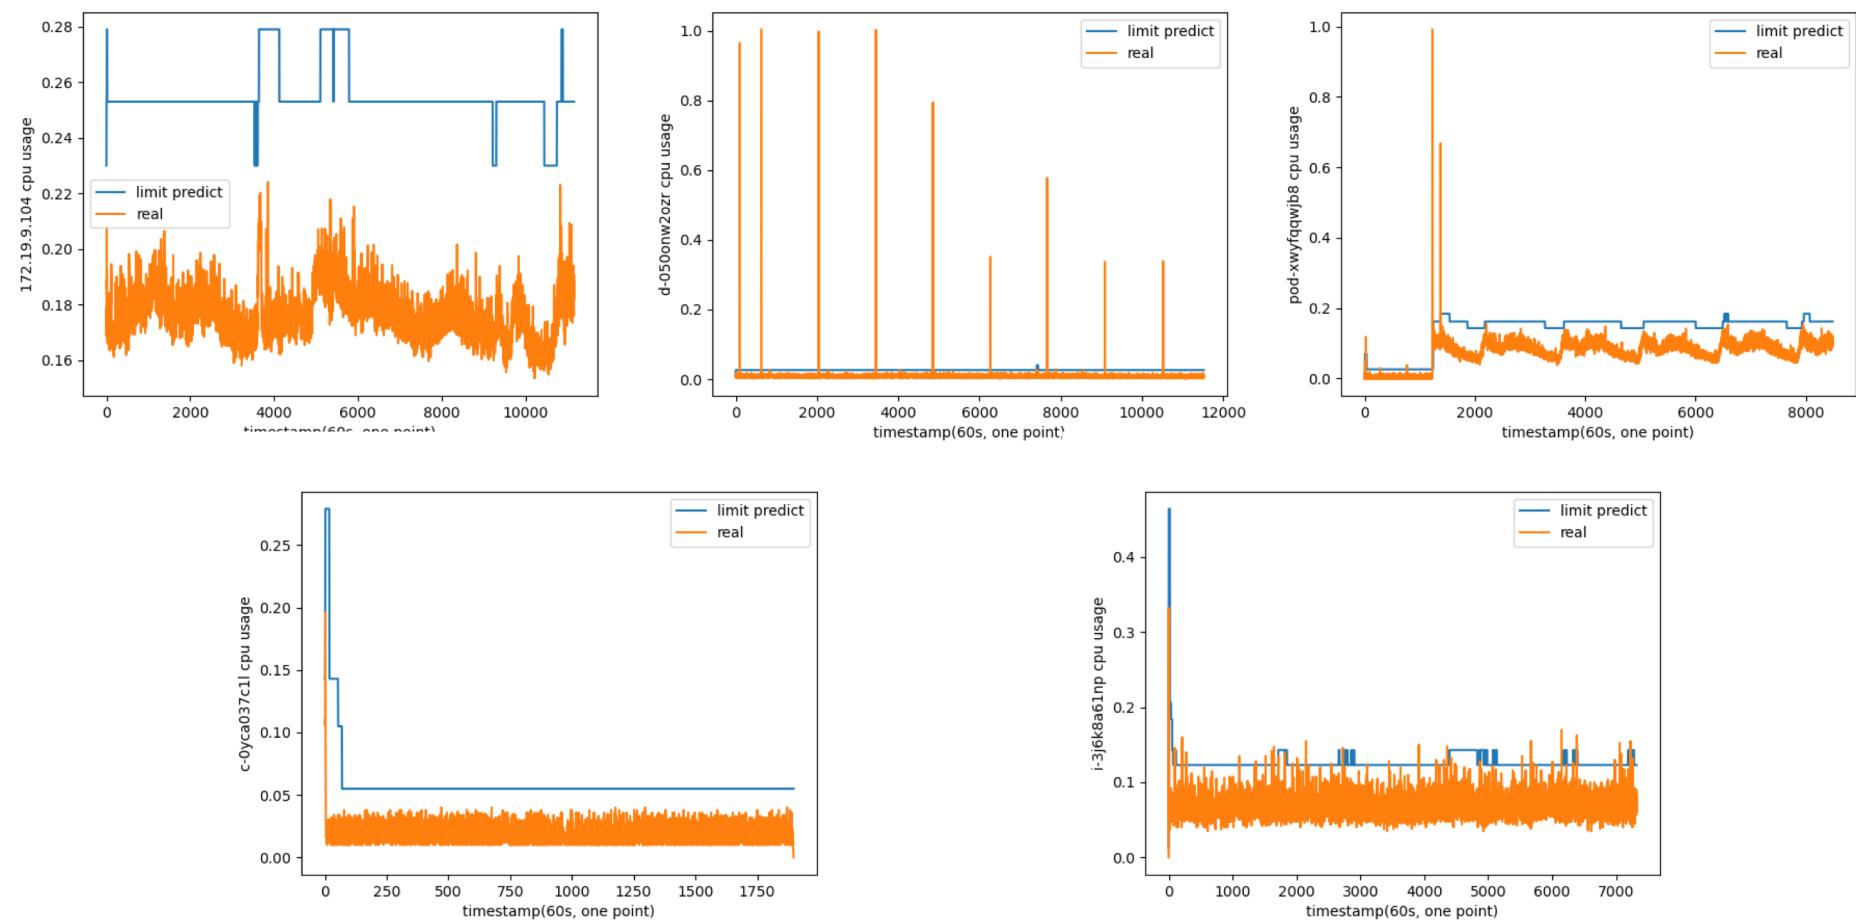


实验结果 (1)



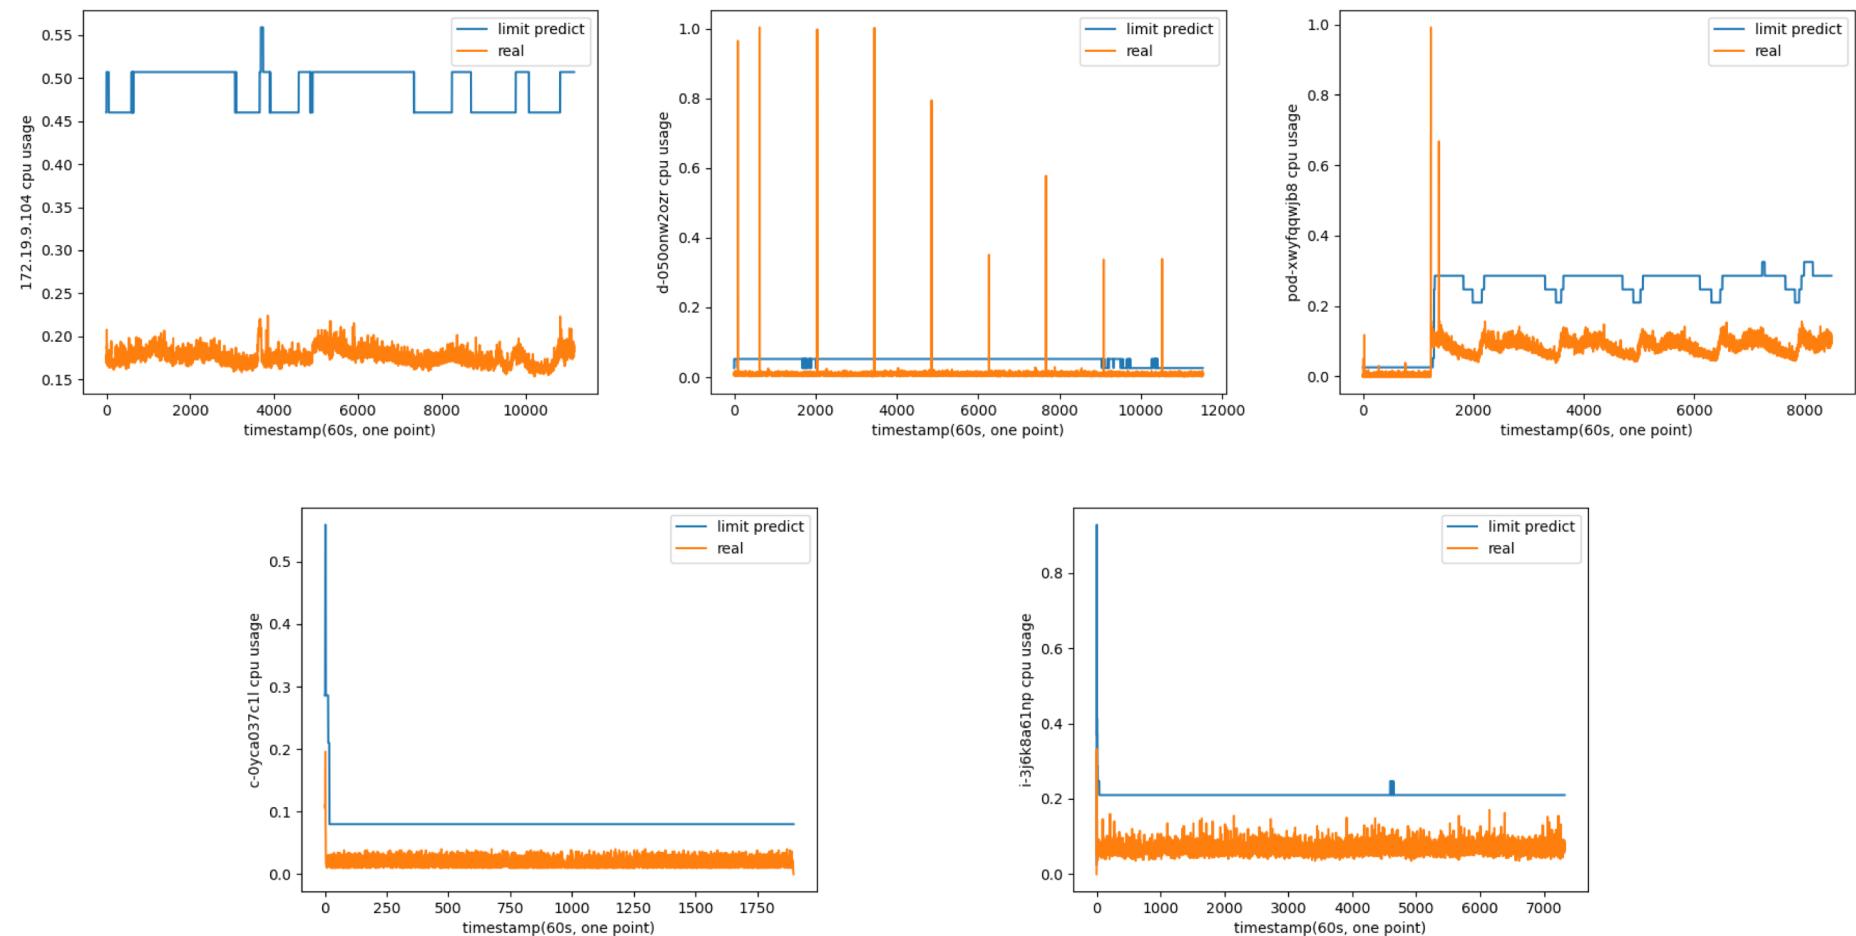
```
--sample-second-interval 60 --target-cpu-percentile 0.95 --cpu-histogram-decay-half-life  
12h0m0s --recommendation-margin-fraction 0.15
```

实验结果 (2)



--sample-second-interval 60 --target-cpu-percentile 0.95 --cpu-histogram-decay-half-life 3h0m0s
--recommendation-margin-fraction 0.30

实验结果 (3)



--sample-second-interval 60 --target-cpu-percentile 0.8 --cpu-histogram-decay-half-life 3h0m0s -
-recommendation-margin-fraction 1.6

结果分析

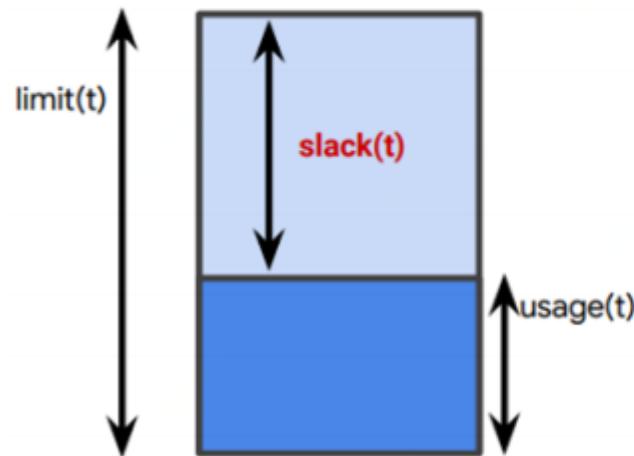
- 选取的样本具有一定的代表性
 - 涵盖了所有jdcloud资源实例（host,vm,nc,pod,docker）
 - 突增的资源消耗，无明显规律的资源消耗，见（1,2）
 - 比较固定的时间变化模式，见（3）
 - 比较稳定的资源利用，见（4, 5）
- 方法具有
 - 处理流式时序数据的能力
 - 接近实时
 - 准确性（可用）
 - 参数化，可调配

参考测评 (google)

absolute slack:

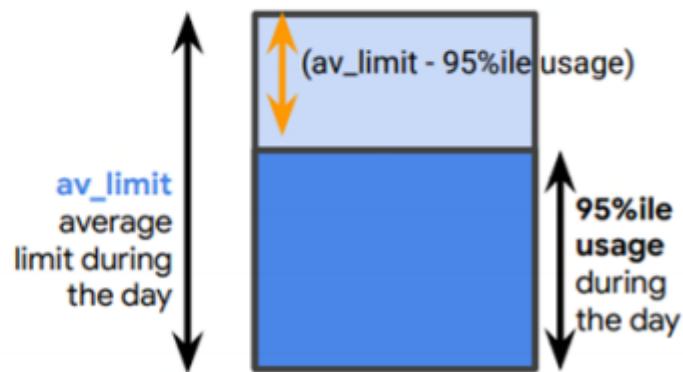
$$\int \text{slack}(t) dt = \int \text{limit}(t) dt - \int \text{usage}(t) dt$$

unit: capacity of a single (largish)
machine



relative slack:

$$(\text{av_limit} - 95\text{-ile usage}) / (\text{av_limit})$$



思想可以借鉴，实施不太匹配
我们的指标数据已经百分比规格化，而不是按时间片模式

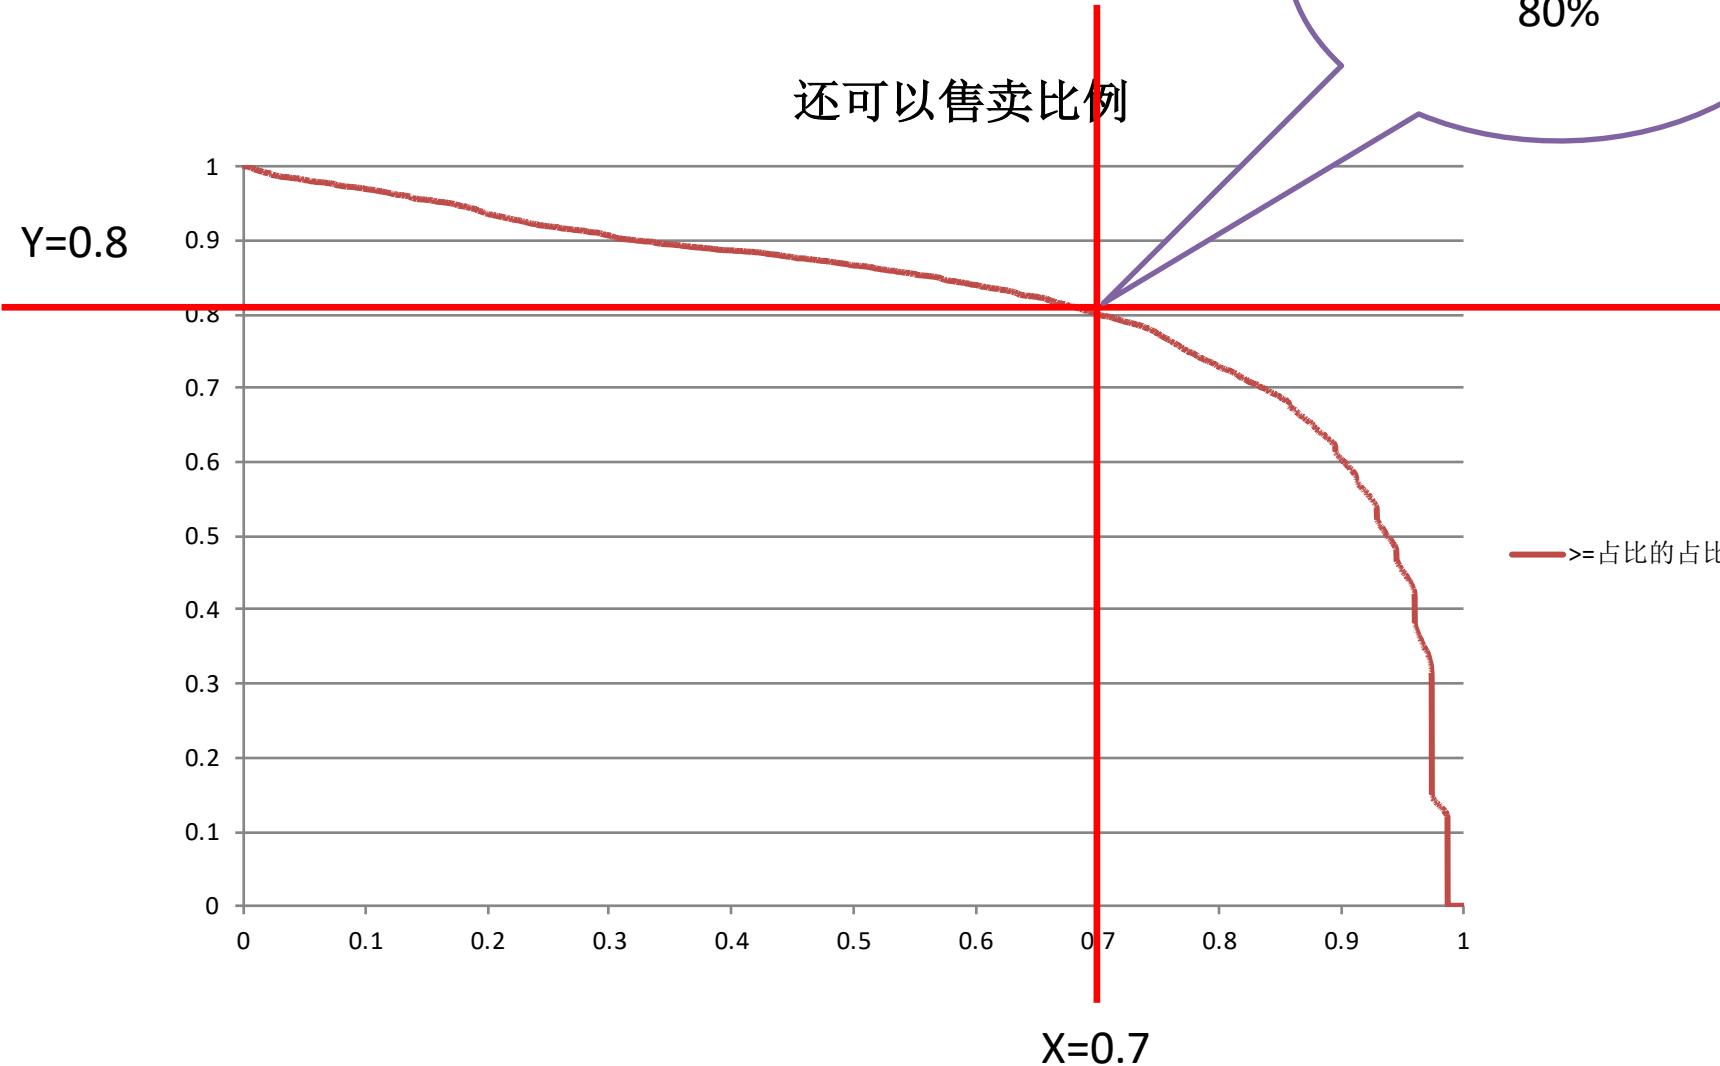
综合测评

- 给定一个线上实例
 - 对应唯一规格
 - limit为规格的定义，静态的
- 引入predicted limit
 - 没使用的资源slack
 - (predicted) limit – used
 - 越小越好：表征已分配但没使用到的资源
 - 可以卖的
 - $1 - (\text{predicted}) \text{ limit}$
 - 越大越好：超出原本的售卖能力(超卖)
- 样本源
 - 部分华北general_2实例（数目~7000）
 - 最大采集近一周的数据
 - cpu利用率原始值，10s间隔

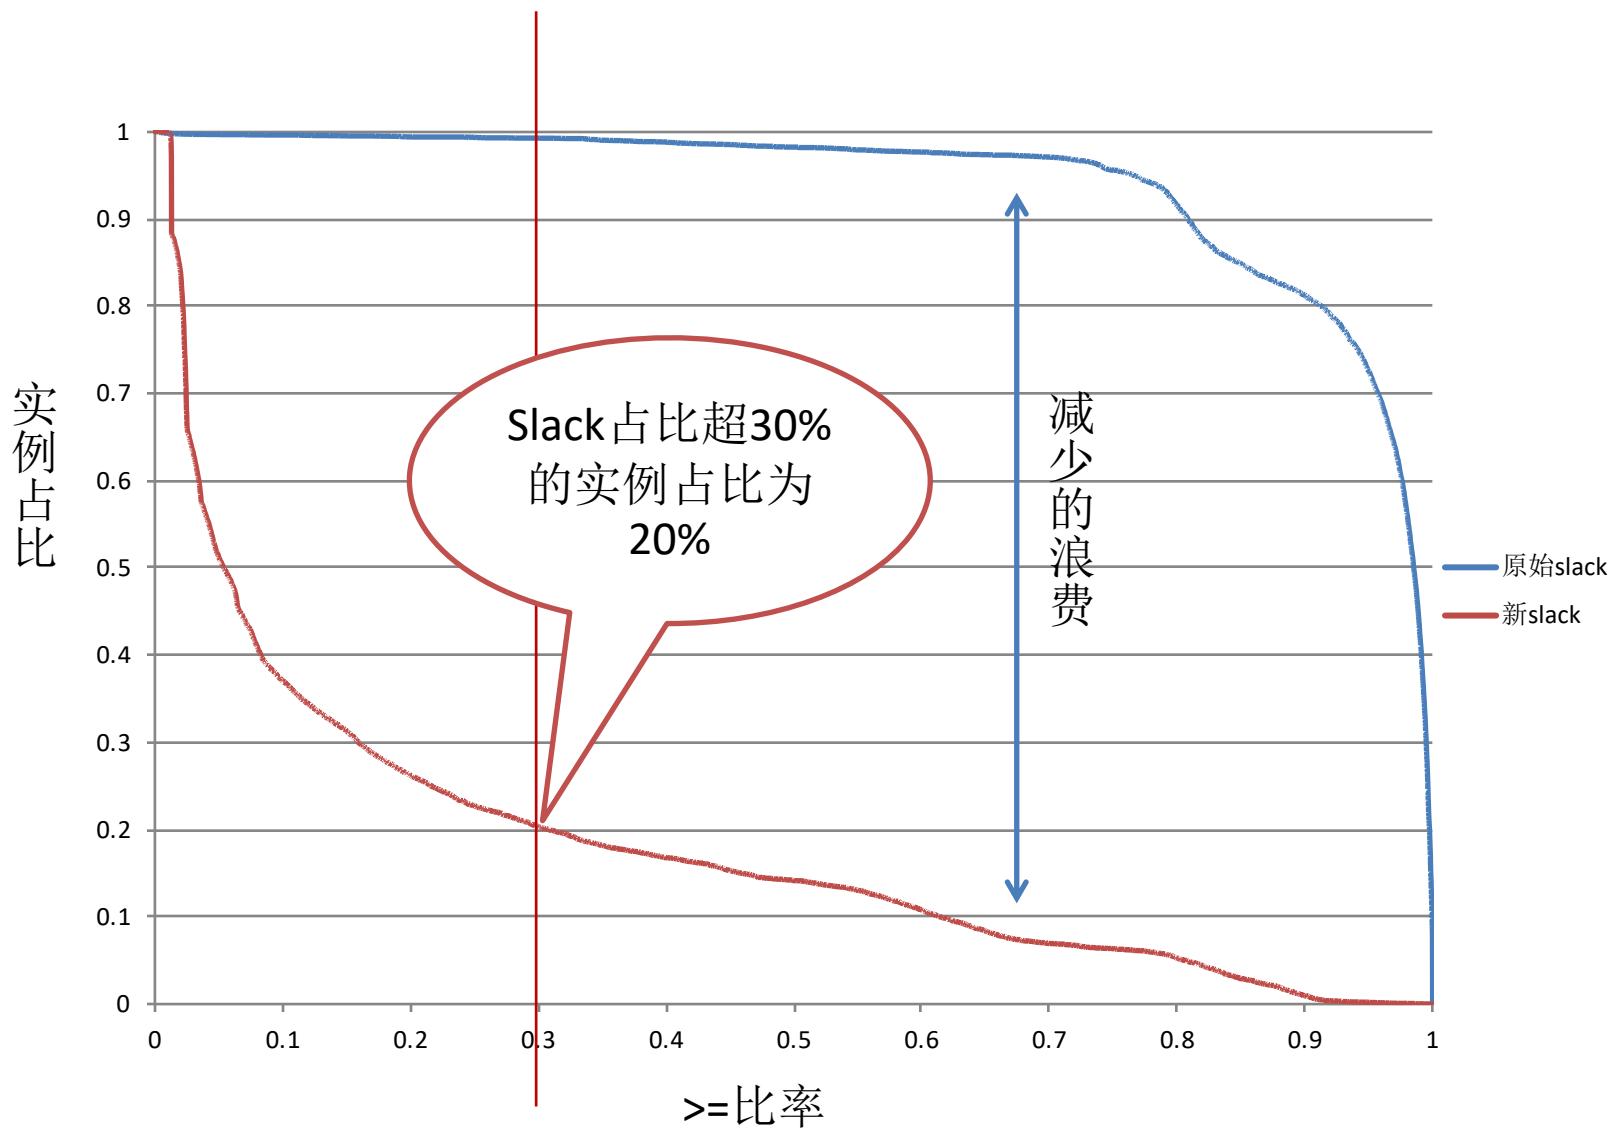
超卖能力

可超卖>70%的
已售出规格占
80%

还可以售卖比例



资源浪费 (slack)



原地增收（理想情况）

- 不增加新机器的原地增收
- 按照上面~7000实例的规格
- 按照predicted limit划分
- 新增可售卖的按规格加权CPU资源

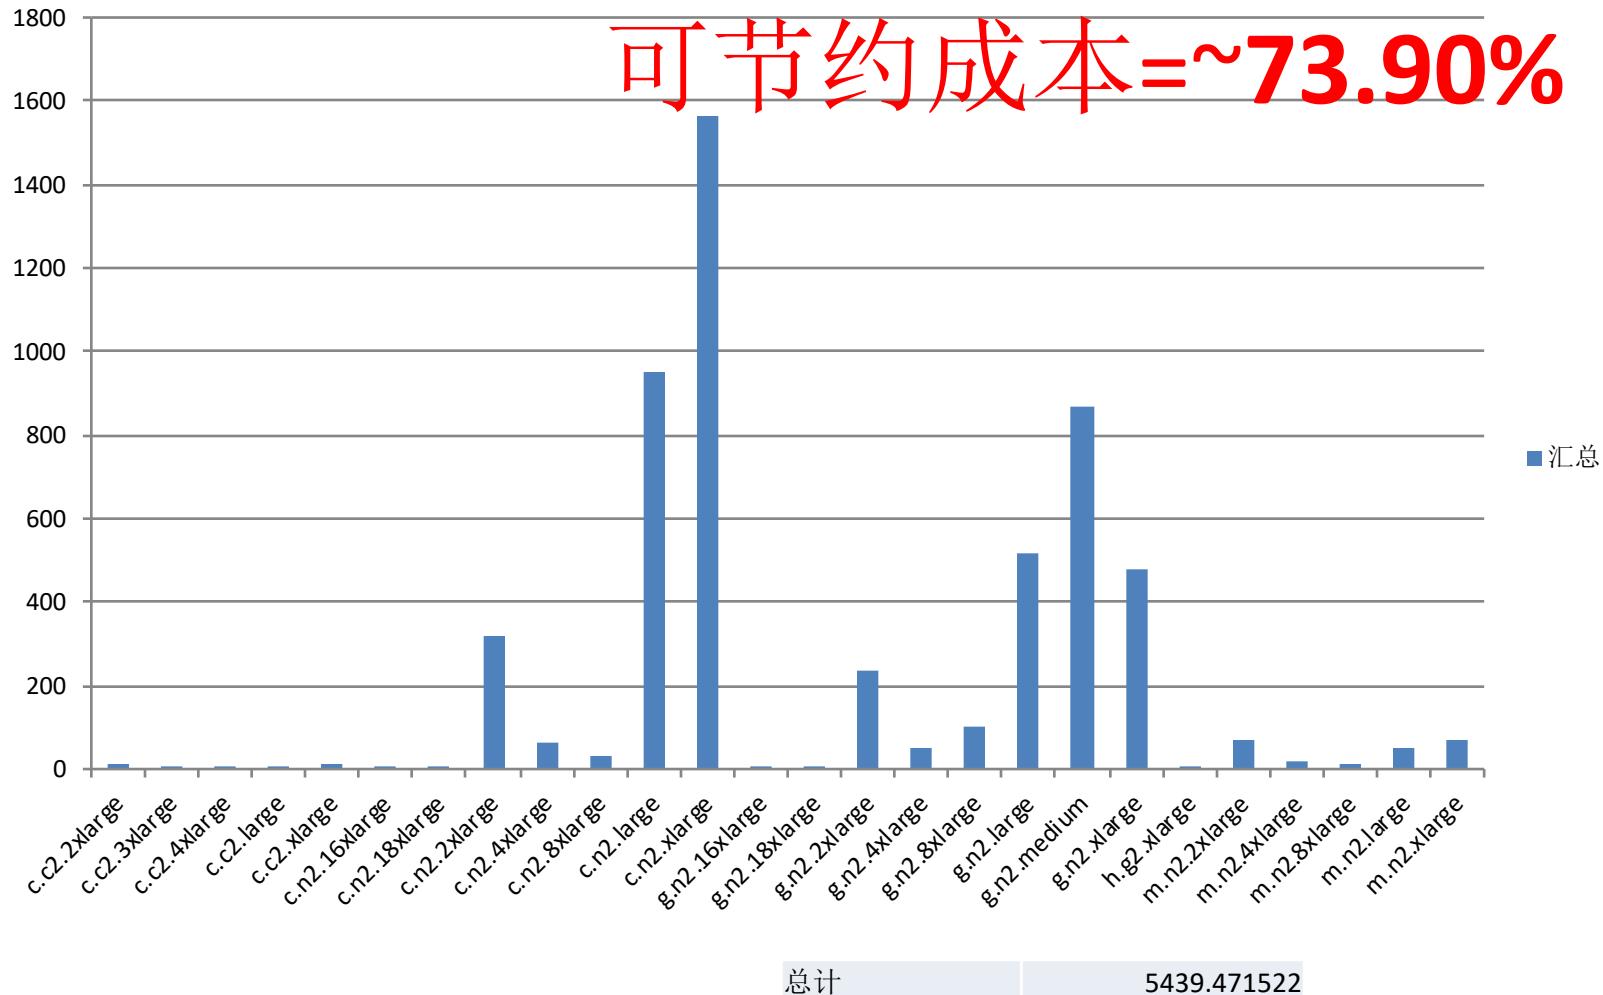
– 新增可售卖

~ 7000

$$\sum_{n=1}^{\sim 7000} \left(\text{instance}_{n, flavor} \right) * (1 - PredictLimit(n))$$

原地增收

新增可售卖规格数目 / ~7000实例



其他

- 思路：从卖规格到卖QoS
- 超卖：挖掘超出原有售卖能力的售卖
- 参考模型：
 - 从同一物理机上的allocations到allocation-set
 - 资源 recommender + updater
- 前提条件：
 - 数据规模
 - 预测
 - 实时性
 - 准确性

参考

- “Autopilot: workload autoscaling at Google”
 - <https://dl.acm.org/doi/pdf/10.1145/3342195.3387524>
- “预测实验对比”
 - <https://git.jd.com/iaas-sdn/jvirt-doc/blob/master/cluster/2020V8/资源利用率提升/数据驱动/预测实验.md>

反碎片化实践方案

- 问题&背景
- 模型&算法
- 评估

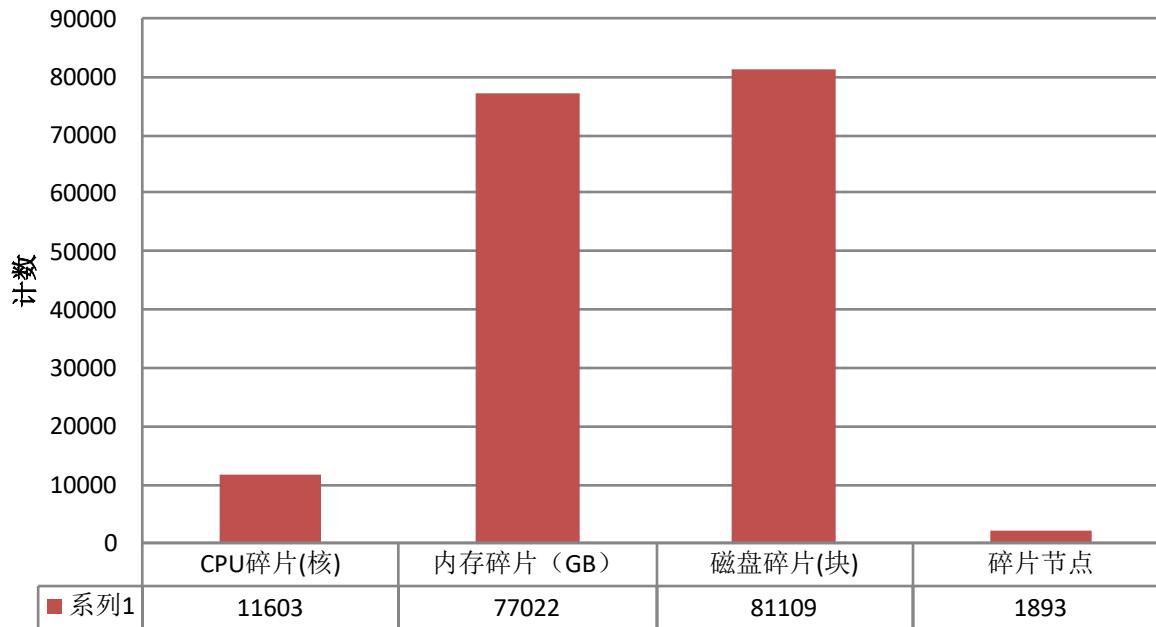
问题描述

- 有资源无法调度
 - 节点无MEM，其他维度资源都是碎片
 - 节点无CPU，其他维度资源都是碎片
 - 暂不考虑使能的影响
 - ~~节点/服务 disable~~
 - ~~节点故障~~

现状

- 样本
 - 华北通用二代机器~6000+台线上计算节点
 - 其中384台节点关闭调度

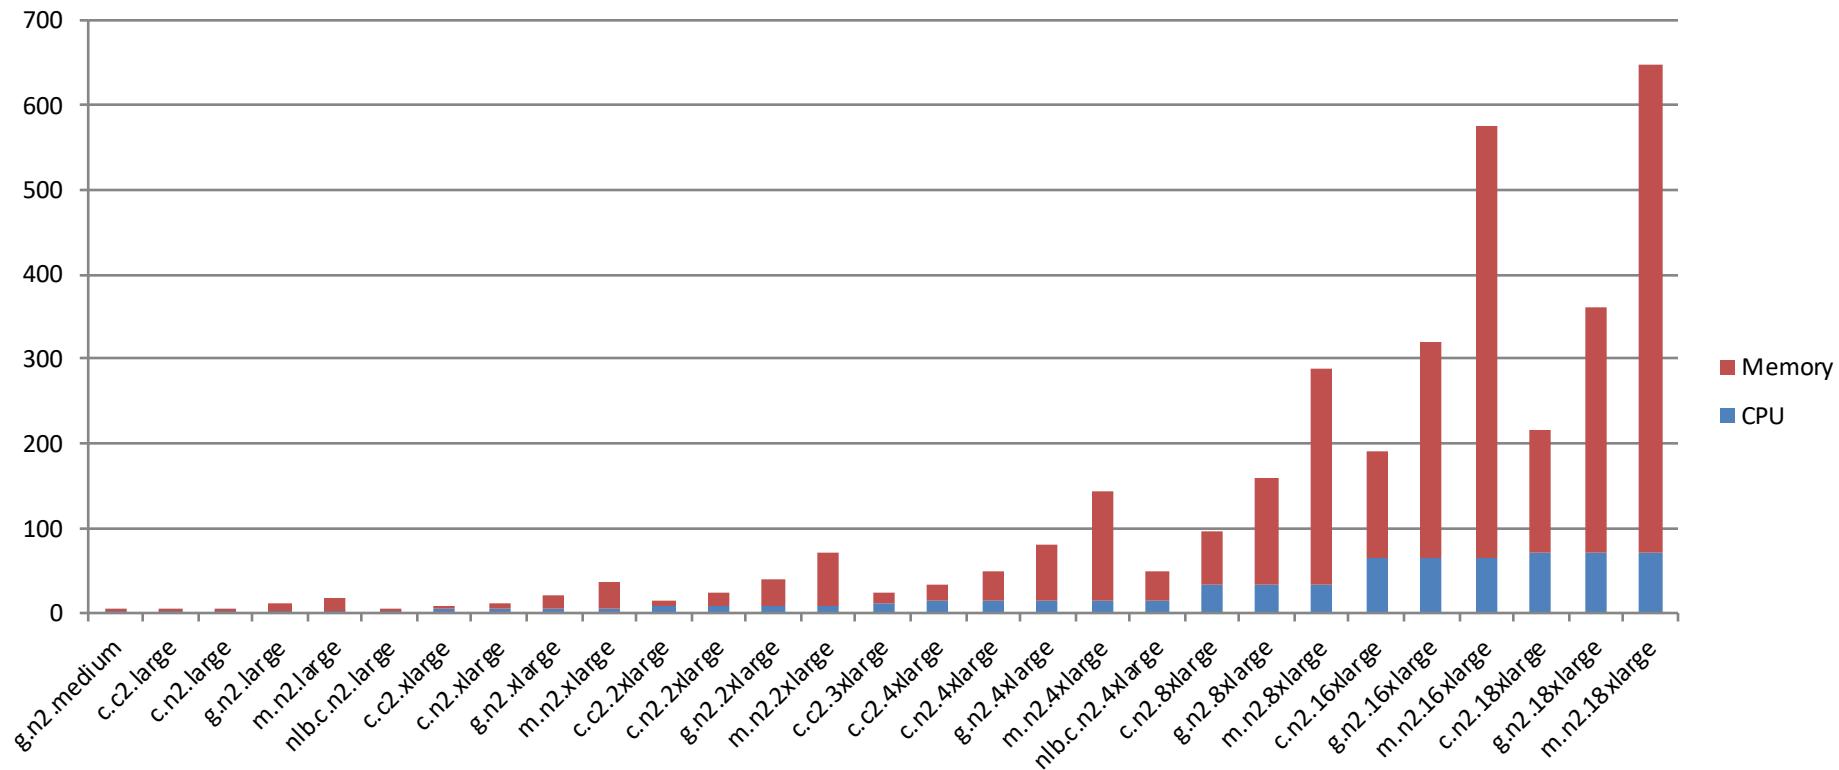
碎片资源



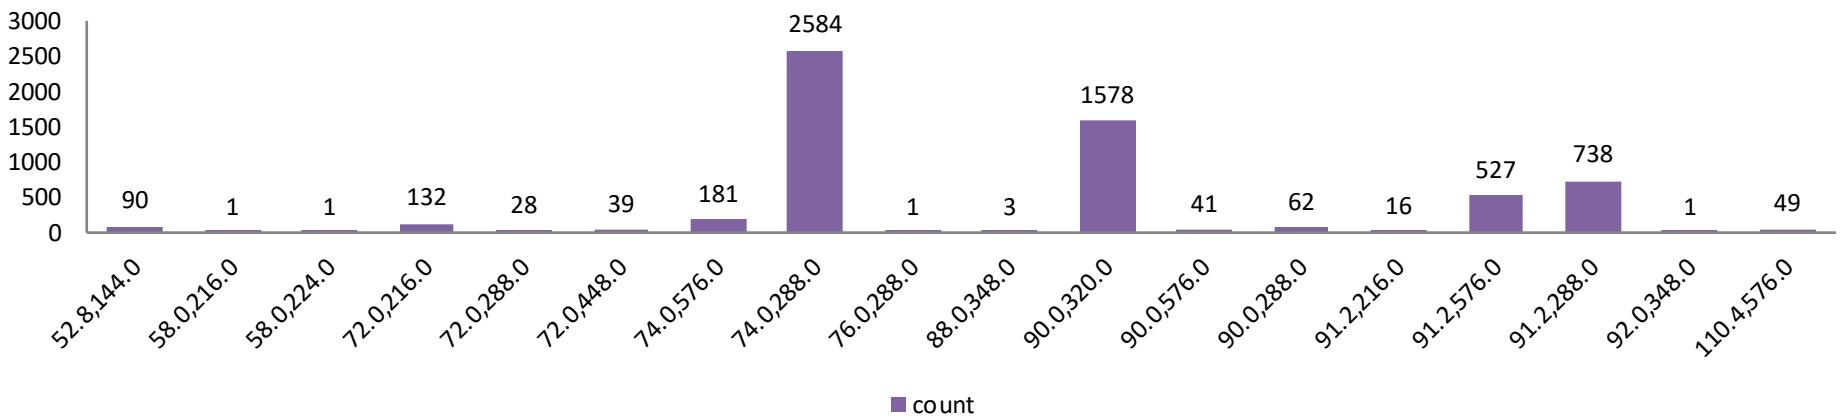
数据集

- 规格分布: ~20+ flavors
- 节点分布: ~6000+ hosts
- 实例分布: ~50000+ instances (vm, nc, pod)

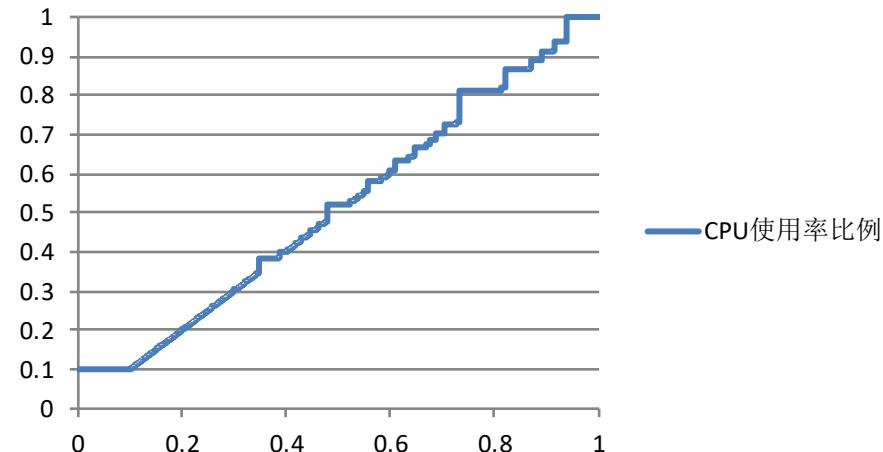
规格



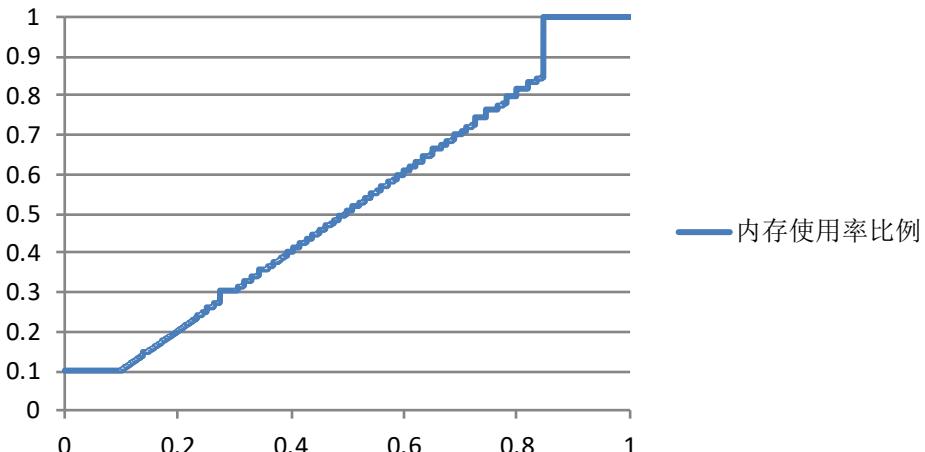
节点



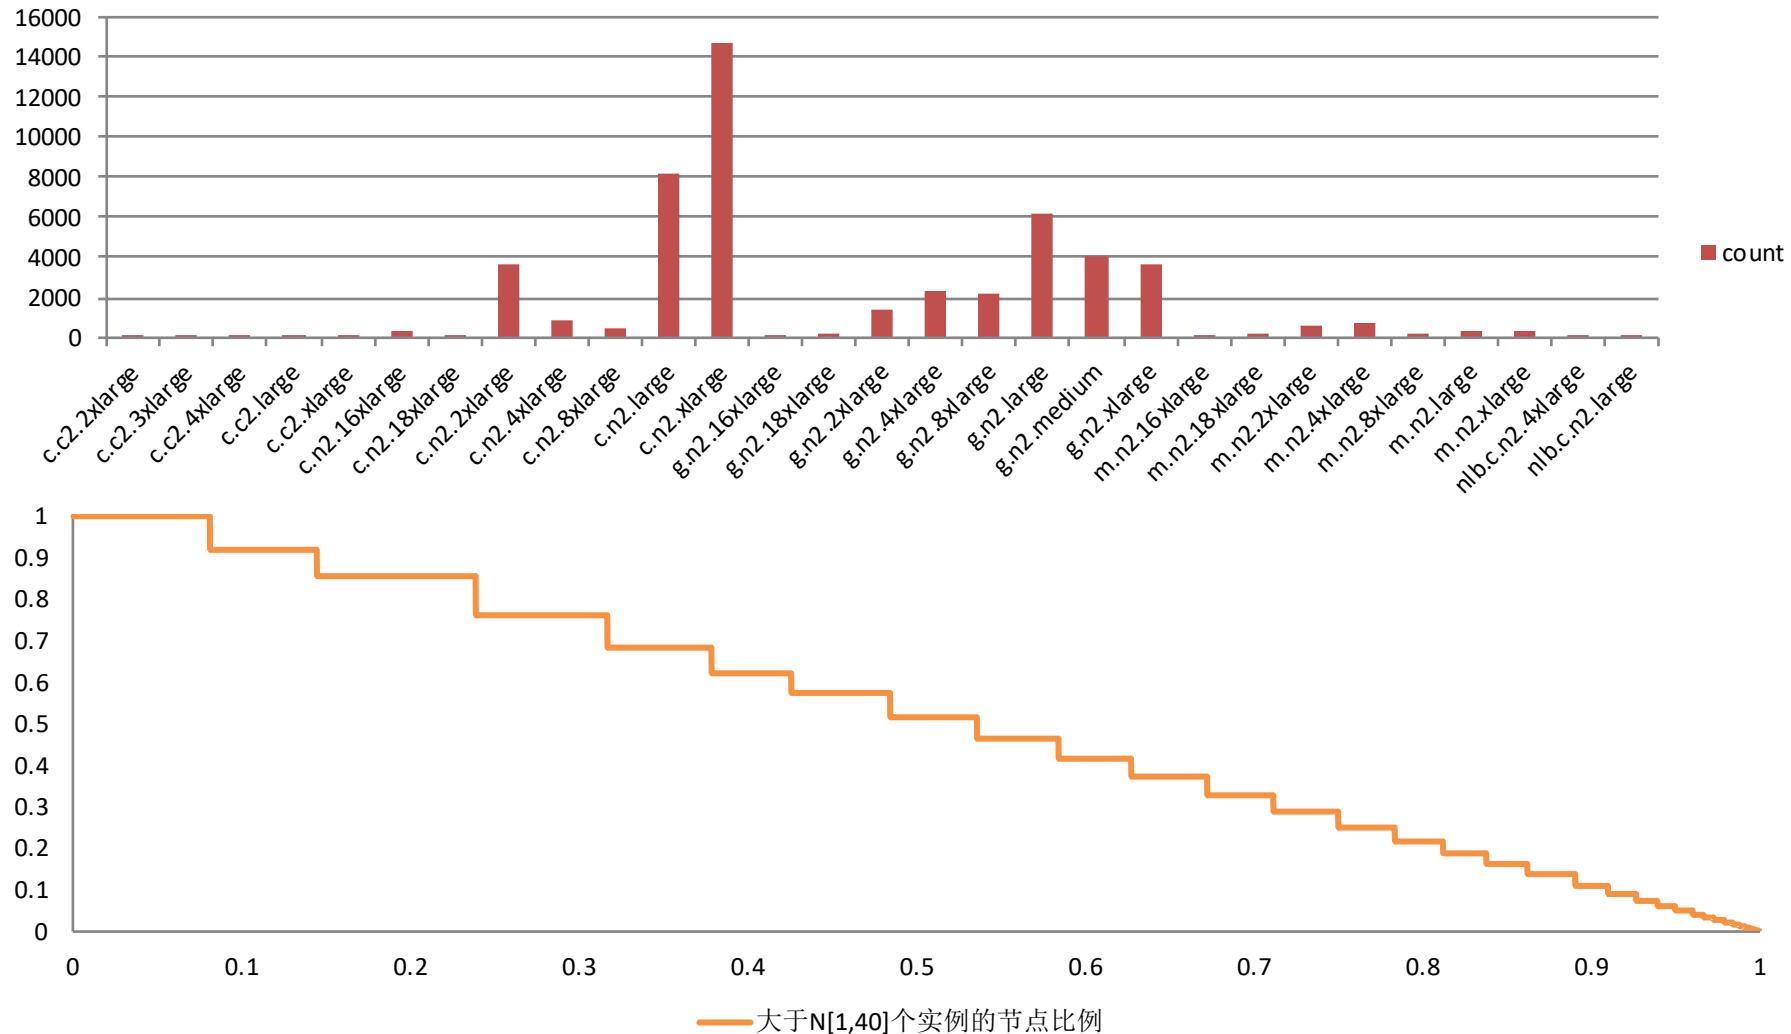
CPU 使用率比例



内存使用率比例



实例



Bin Packing

主流算法	描述
Next fit	满足则装当前打开的1个bin, 不满足开1个新bin放入
Next k fit	打开上次k个bin, 找一个满足, 不满足开1个新bin放入
First fit	从已有bin找第一个满足的bin, 不满足开1个新bin放入
Best fit	从已有bin找最合适的bin, 不满足开1个新bin放入
OPT fit	看到所有bin, 所有package, 用最少的bin装填, 必然碎片最少

装箱参数	资源分派、放置	备注
#bins, capacity	#hosts	bin一般都一样, 节点有差异
#packages, size	#flavors	package有长宽高, 并且可旋转

综合评估

- 调度评估
 - 最优化调度
- 综合测评指标
 - 碎片程度，越少越好（计划任务所需机器少）
 - 节点状态，越均衡利用越好
 - 备机需求，越少越好
 - 迁移代价，[迁移次数，涉及节点，候选实例]

调度评估

- Classic Definition
 - A packing heuristic
 - Resource allocation demand vector “ $<$ ” machine resource vector
 - “ $<$ ”: alignment score (**AS**)
 - “**AS**” works because:
 - No over-allocation
 - Bigger balls get bigger scores
 - Abundant resources used first
- NP hard
 - 最优化方法、整数规划、约束规划、元启发式方法
 - 简化为动态规划

全局装箱任务

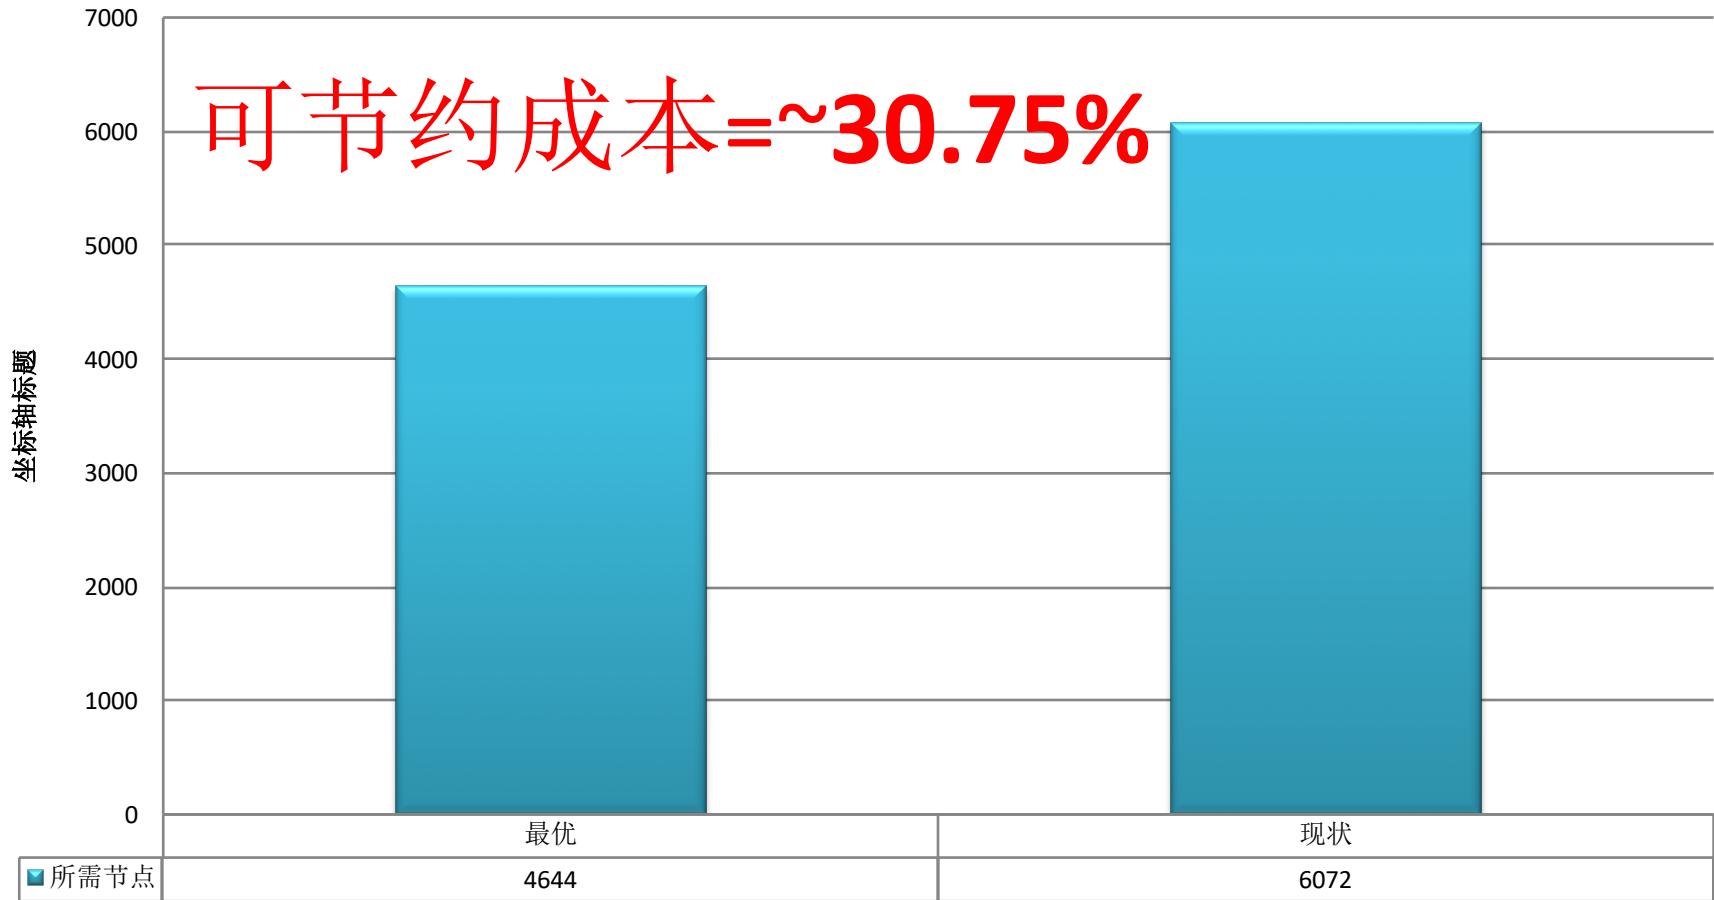
bins		packages				
capacity	count	规格	CPU	Memory	count	
52.8,144.0	90	g.n2.medium	1	4	4108	
58.0,216.0	1	c.c2.large	2	2	65	
58.0,224.0	1	c.n2.large	2	4	8136	
72.0,216.0	132	g.n2.large	2	8	6175	
72.0,288.0	28	m.n2.large	2	16	288	
72.0,448.0	39	nlb.c.n2.large	2	4	10	
74.0,576.0	181	c.c2.xlarge	4	4	77	
74.0,288.0	2584	c.n2.xlarge	4	8	14741	
76.0,288.0	1	g.n2.xlarge	4	16	3626	
88.0,348.0	3	m.n2.xlarge	4	32	297	
90.0,320.0	1578	c.c2.2xlarge	8	8	97	
90.0,576.0	41	c.n2.2xlarge	8	16	3602	
90.0,288.0	62	g.n2.2xlarge	8	32	1418	
91.2,216.0	16	m.n2.2xlarge	8	64	646	
91.2,576.0	527	c.c2.3xlarge	12	12	2	
91.2,288.0	738	c.c2.4xlarge	16	16	93	
92.0,348.0	1	c.n2.4xlarge	16	32	835	
110.4,576.0	49	g.n2.4xlarge	16	64	2339	
总节点数	6072	m.n2.4xlarge	16	128	756	
		nlb.c.n2.4xlarge	16	32	42	
		c.n2.8xlarge	32	64	427	
		g.n2.8xlarge	32	128	2175	
		m.n2.8xlarge	32	256	170	
		c.n2.16xlarge	64	128	353	
		g.n2.16xlarge	64	256	46	
		m.n2.16xlarge	64	512	10	
		c.n2.18xlarge	72	144	15	
		g.n2.18xlarge	72	288	178	
		m.n2.18xlarge	72	576	207	
		总实例数			50934	

最优化方案（理想情况）

Problem Statistics					
	PMs	Jobs	VMs	Mappings	Used PMs
Number	6072	50934	50934	0	0
	CPU		RAM		
Total Capacity		497417		2011784	
Minimum Capacity		52		144	
Maximum Capacity		110		576	
Average Capacity		81.92		331.32	
Total Requirements		362848		1294062	
Minimum Requirements		1		2	
Maximum Requirements		72		576	
Average Requirements		7.12		25.41	
Usage Percentile	0.7294644132000000		0.6432410239000000		

降本

所需节点



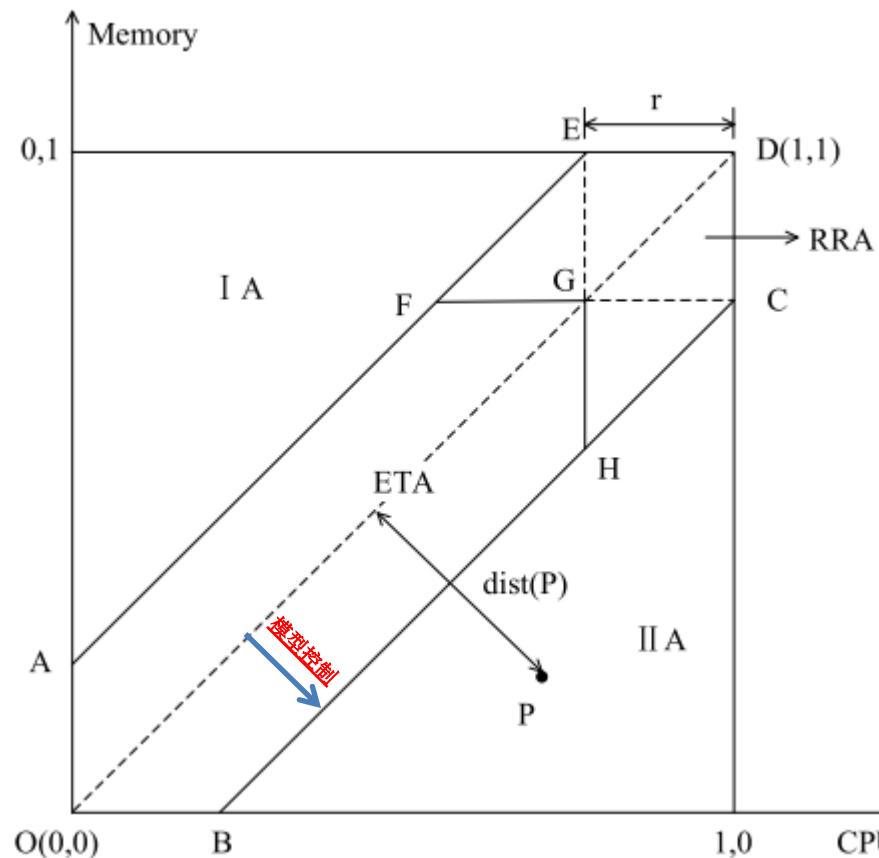
根因分析

- 规格不按比例齐整
- 节点上线有时间线
- 其他调度参数
 - 指定节点
 - Tag
 - 亲和性, HG
- 时序请求无法最优
 - 请求有先有后
 - 删除+创建

解决方法

- 节点分类
 - 空闲机器
 - 满载机器
 - 偏科机器
- 迁移计划 → **再平衡，利用率优先**
 - 调度
 - 效率
 - 目标

节点如何，为什么分类

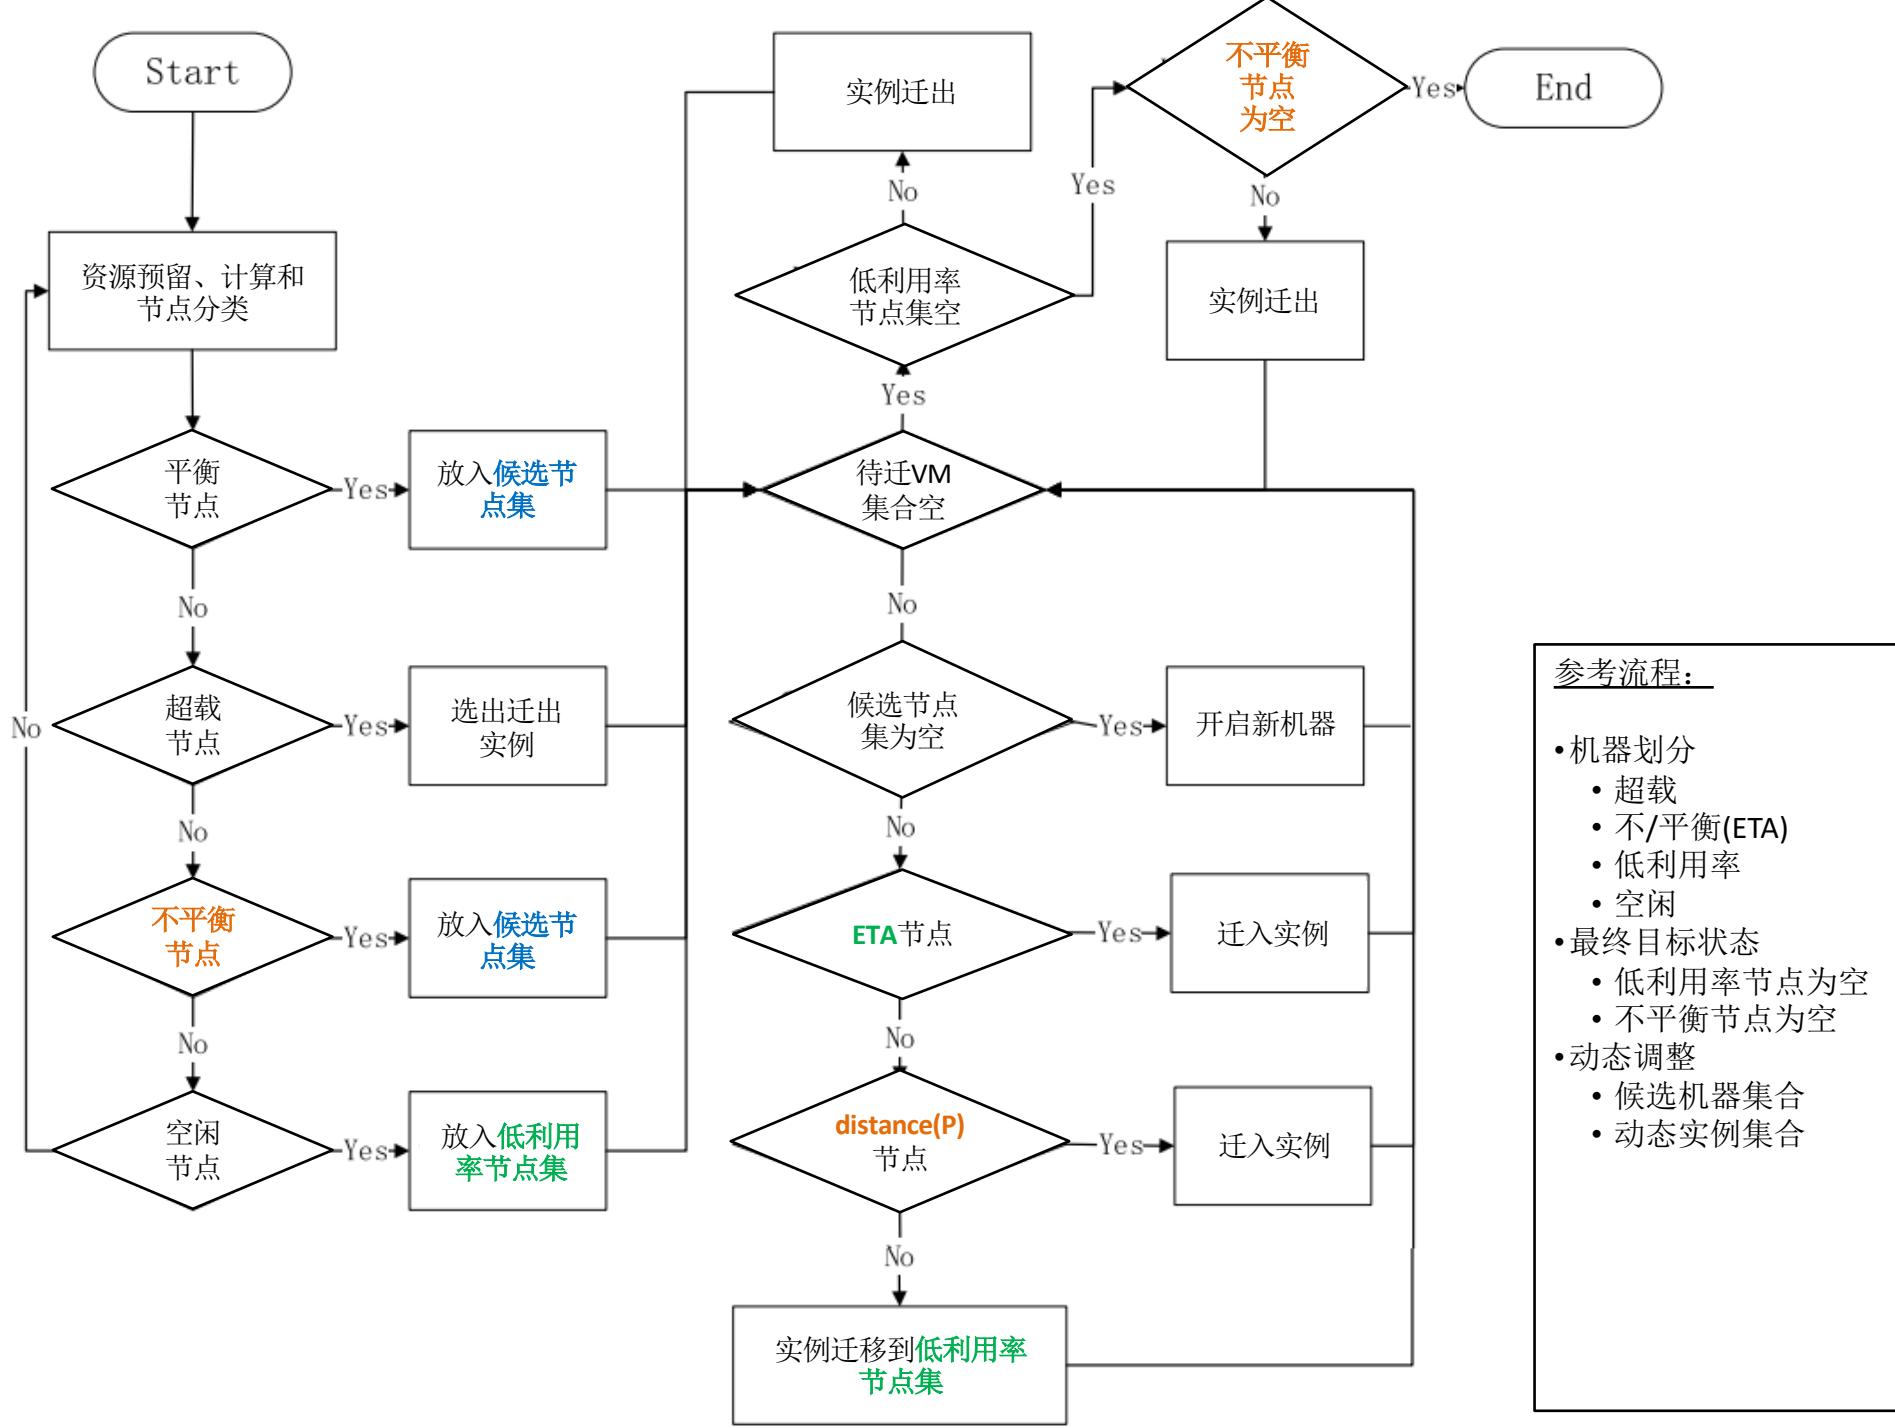


$$f(x, y) = \begin{cases} \text{(饱和, 过载)} \\ RRA, & (1 - x) \leq r \& (1 - y) \leq r \\ \text{(平衡)} \\ ETA, & dist(p) \leq \frac{\sqrt{2}}{2}r \\ \text{(偏科)} \\ IA \text{ or } IIA, & \& (1 - x) > r \& (1 - y) > r \\ & otherwise \end{cases}$$

- 基于资源利用分布对节点建模
 - 通过资源分布划分节点类型
 - 根据节点类型，判定节点迁移状态、迁移调度

物理意义

- RRA: Resource Reserved Area, CDEFGH
 - 应对负载突增
 - 保证SLA
- ETA: Equilibrium Tolerance Area, AOBHGF
 - 合理利用率
- Area I & Area II: IA, IIA
 - 不平衡区域
 - Dimension Rate



应用场景

- 动态备机遴选
 - 去掉常規备机方案
 - 664 + 21 + 58 + 44
 - 近1万/天
 - 增加售卖
- 预选目标机
 - 支持运维迁移计划
- 增加可售卖
 - 去碎片化
- 预测性维护
 - 兼顾资源利用

深入方向

- 考慮
 - 能耗（Energy consumption），也是需要钱的
 - 和利用率相关

预测性维护

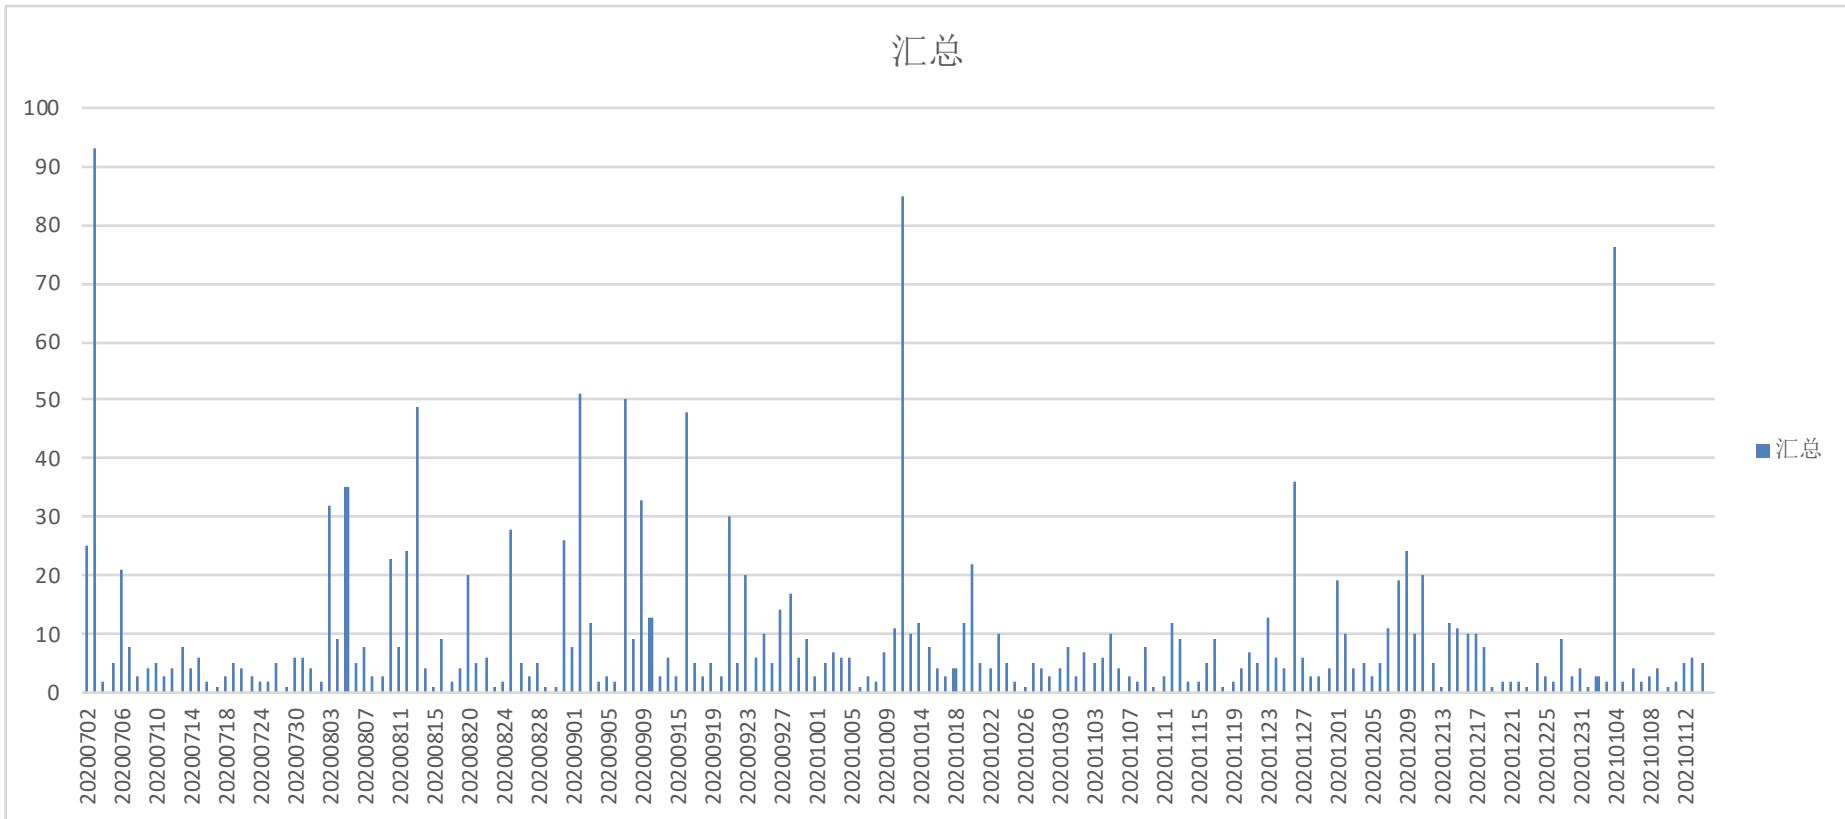
- 事后维护影响SLA
 - 代价大
 - 临时性
 - 缺乏优化目标
- 事前维护
 - 日常Routine
 - 保护SLA

问题

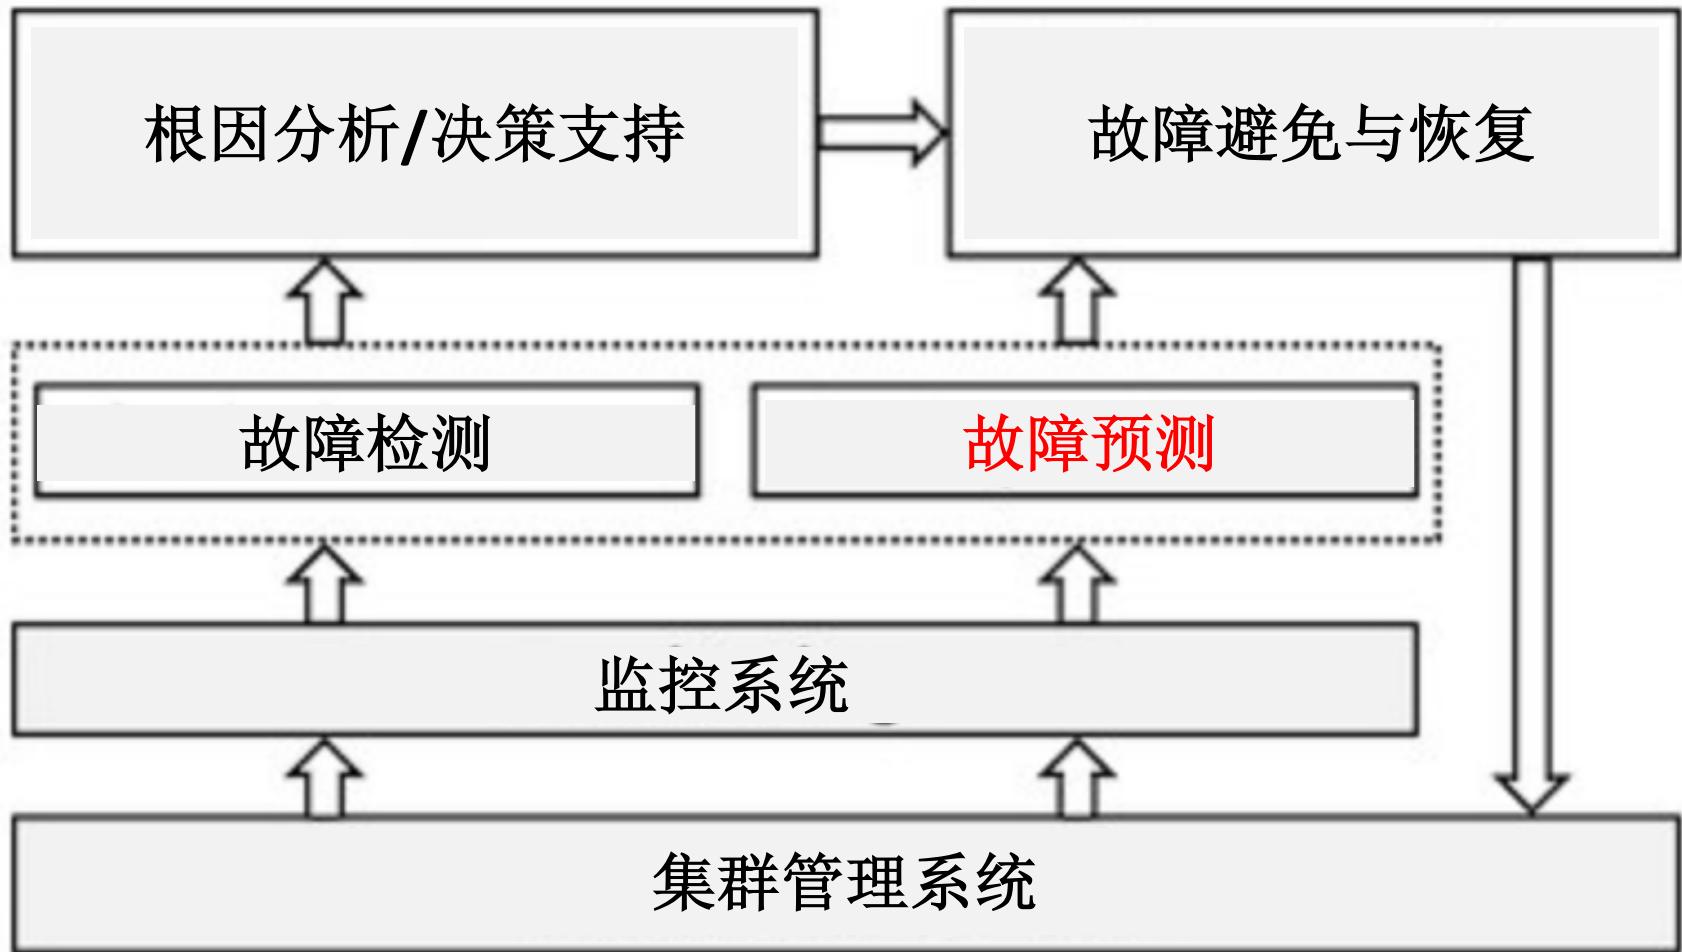
- 连续出现节点故障
- 无法做到故障避免
- 影响业务连续性

现状

- 按时间统计故障机分布
- 按时间统计影响实例分布



框架



数据

- 数据集
 - 正例
 - 收集线上迄今故障节点和故障时间点
 - 以故障点为终止时间往前搜集一周监控指标
 - Cpu, Mem, Network, Disk
 - 历史原因可能采样点比较大
 - 反例
 - 整理中同物理型号和位置节点
 - 近期无故障节点
 - 时间窗口对应
 - 分离出训练集和测试集

方法

- 指标数据滑动窗口
- 节点指标分别训练预测
 - CNN-LSTM
 - 未来的各项指标值
- 基于历史数据训练分类器
 - PCA + SVM
 - 通过指标判定故障还是非故障
- 联合所有预测数据，过分类器
 - 预测故障

提升空间

启发式资源管理

- 现状分析
- 问题

算法

- Vector bin packing
- NP-hard
- Greedy-algorithm
- Worse fit & Best fit
- First fit decreasing (FFD)
 - FFD heuristics
- Best fit decreasing (BFD)
 - BFD heuristics

优化

- Dimension extension
- Power consumption
- Resource wastage
- Server loads
- Inter-VM
- Storage/Network traffic

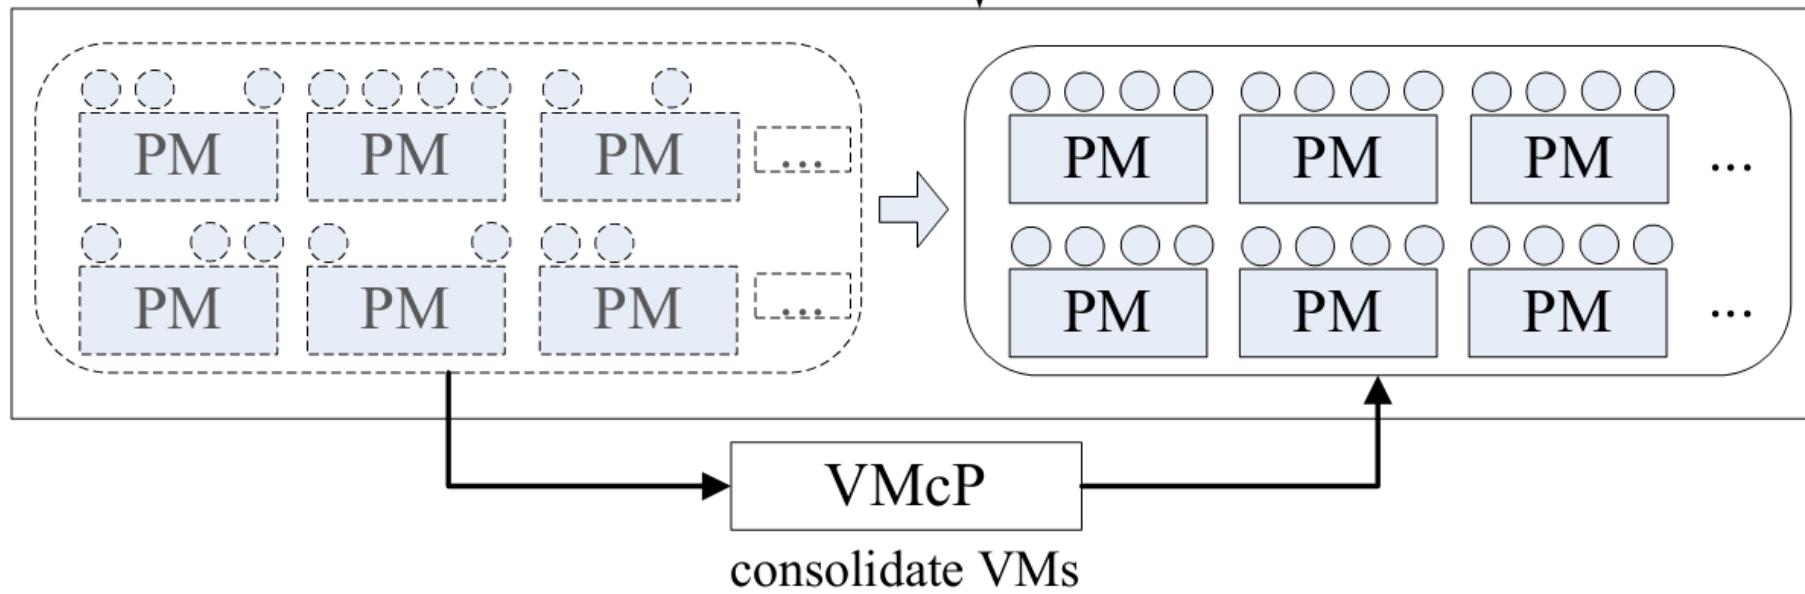
回顾资源调度

... ○ ○ ○ ○ ○ ○ ○ ○ ○ ○ ○ ○

continus VM requests
(deployment requests & removal requests)

VMiP

deploy/remove VM



两个场景：连续请求（incremental）, 动态调整(consolidate)

Vector

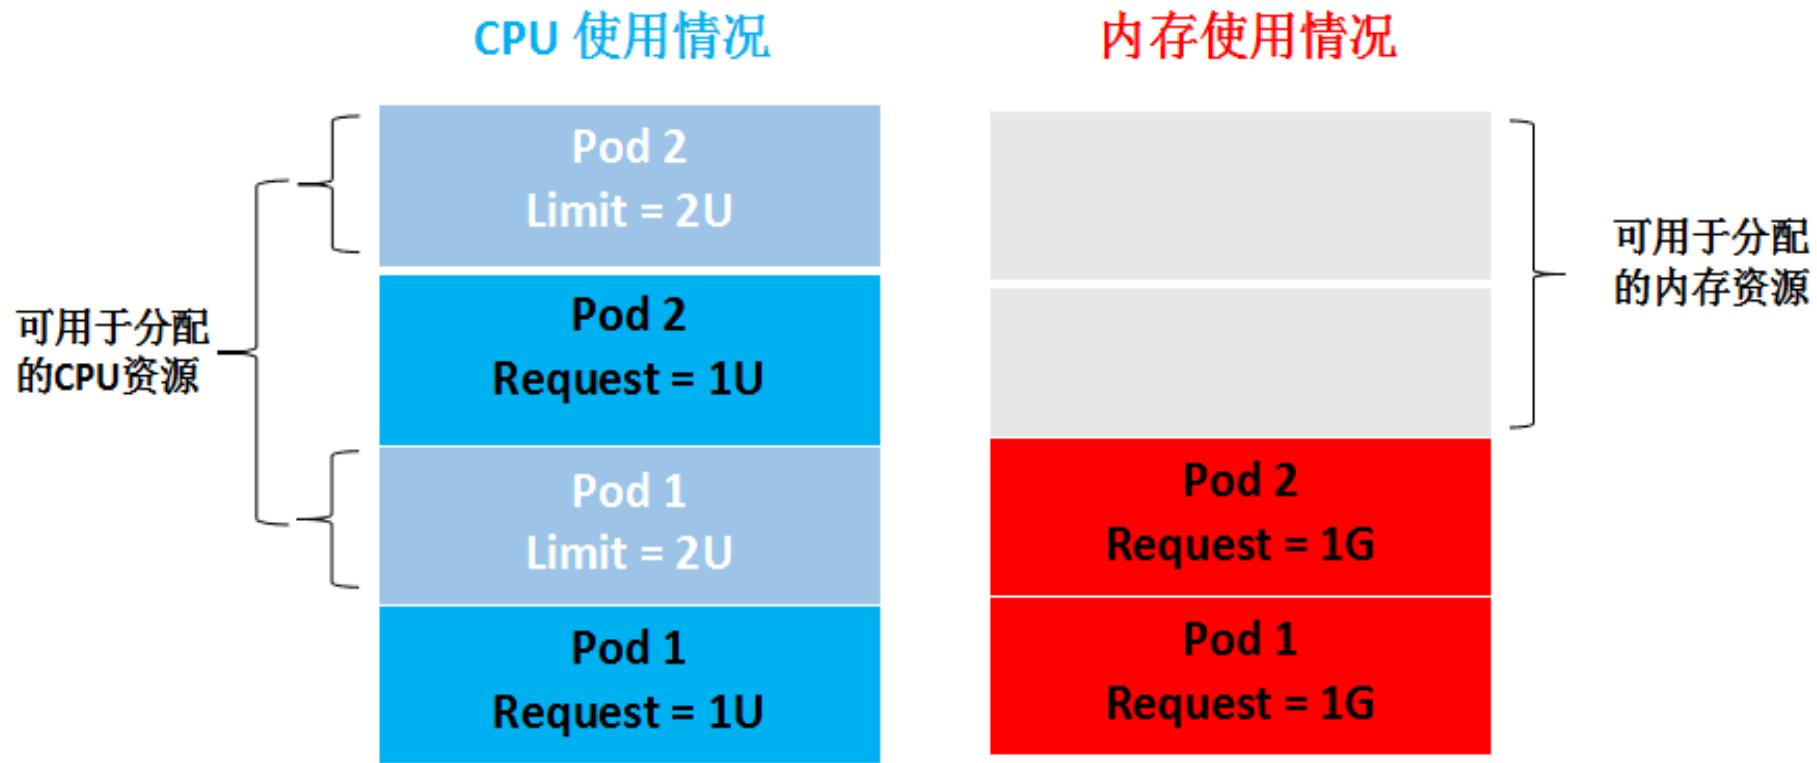
元素	规格	多维资源	使用率	下限	已用	可用	预测	上限	状态
实例	部分资源	全部资源维度	实例指标	基准性能	实际消耗	空闲	基于历史	规格	可调度、需调整、优先级、抢占
节点	机型	+机器属性	节点指标	预留	+实际消耗	+空闲	基于历史	规格	空闲、平衡、过载、能耗
集群	机型集合	+集群属性	聚合指标	备机	+实际消耗	+空闲	聚合预测	聚合规格	浪费、能耗、水位、库存、备机

1. 启发式
2. 多目标优化
3. 原生支持批量

Vector应用

- 调度
 - 节点打分，启发规则
 - 节点状态，资源不/紧张
 - 维度稀缺，不/可压缩
 - 原生支持超卖&抢占&优先级
- 评估
 - 集群状态（指标）
 - 资源浪费、闲置成本、平衡状态
 - 负载，与CPU利用率正相关

参考K8s: Request-Limit模型



原理分析

- 资源维度三元表达
 - Request (可调度)&Limit(后限制)&Usage(实际值)
- 资源空闲时
 - 节点上利用率低
 - 发挥Request=0的作用
- 资源紧张时，包括真/假紧张
 - 抢占，节点上的总资源 < 所有实例Limit的总和
 - 可压缩资源(CPU), 按照Request的比例分占CPU调度时间片
 - 不可压缩资源(Mem), 利用优先级高低驱逐(迁移)实例
 - 优先搞走Request=Limit=0
 - 其次搞走 $0 < \text{Request} < \text{Limit} < \text{Infinity}$
 - 保留 $\text{Request} = \text{Limit}$, 除非节点剩余资源未达剩余资源的要求
- 极限情况，资源紧缺处理，真紧张
 - 不可压缩资源抢占的资源回收策略
 - SLA保证，突发预留

直接映射

RoadMap	Jdcloud	kubernetes	rms调度
现状 ① 所有CPU模式 ② 所有规格类型 ③ CpuAllocRatio	Flavor(cpu)	$0 < \text{Request}(\text{cpu}) = \text{Limit}(\text{cpu})$	Bin packing
	VmAlloc(cpu)	$\text{Request}(\text{cpu}) = \text{Limit}(\text{cpu})$	
	HostAvail(cpu)	HostAvail(cpu)	
基于利用率 ①②③ +MixMode	PredictVmUtil(cpu)	Request(cpu)	Bin packing
	Flavor(cpu)	Limit(cpu)	
	HostAvail(cpu)	HostAvail(cpu)	
基于Vector ①②③ +MixMode	PredictVmUtil(cpu)	Request(cpu)	Bin packing decreasing (BPD)
	Flavor(cpu)	Limit(cpu)	
	HostAvail(cpu)	HostAvail(cpu)	

复用策略

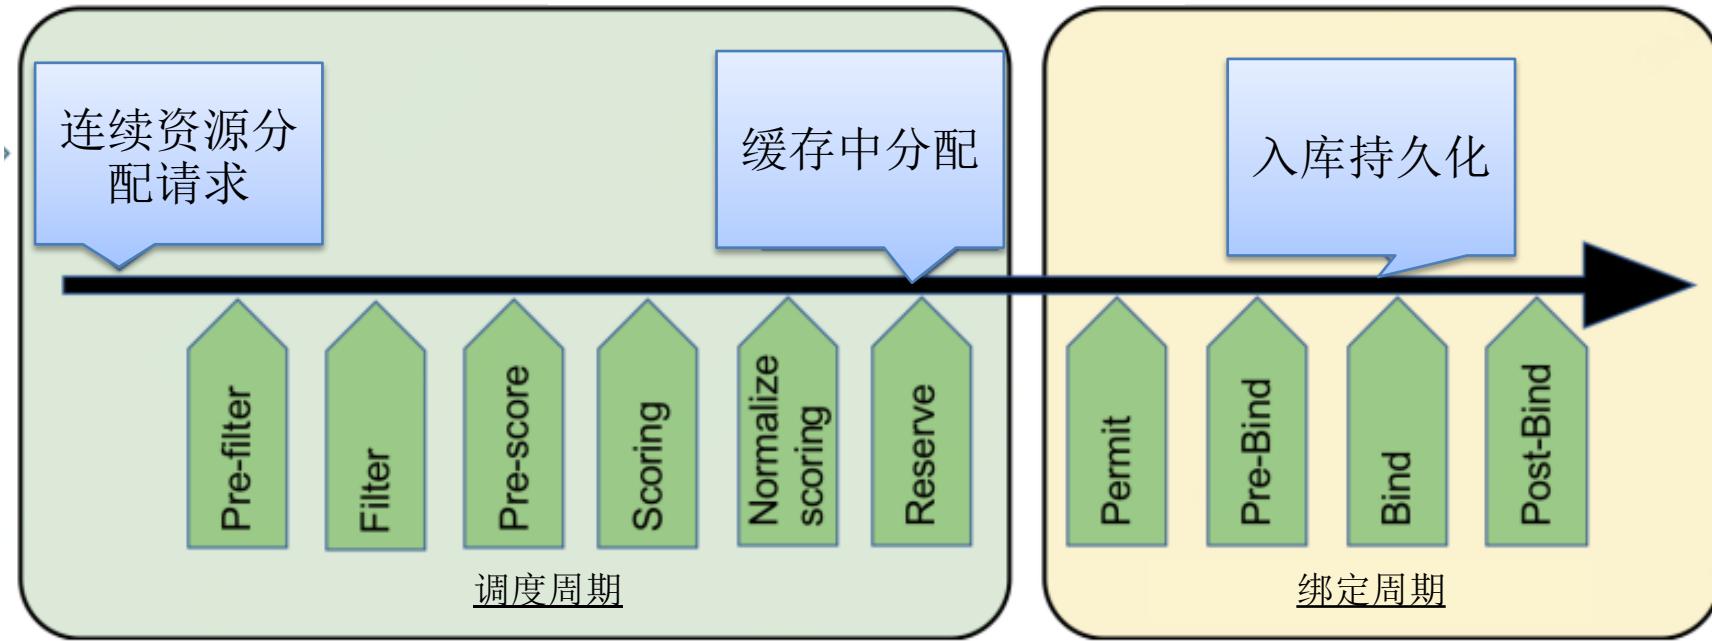
大类	小类	Context	控制	其他
实例	通用	一级调度 Request = Flavor	Limit = Flavor	
	抢占	一二级调度 Request = 0 , 离线业务	Limit = Flavor 基准	
	优先级	二级调度 Request + Limit	$\sum Limit > Host$	
	超卖	一二级调度 Request = [0, 预测Limit]		
节点	绑核	实际扣减, 在线业务		
	超卖	等价扣减, 弹性+通用业务		
	碎片	超卖 + Consolidation		
	功耗	基于利用率模型	$\min \bigcup (\text{节点} + \text{机架})$	

架构目标

- Scheduler Framework
 - 轻量化, Framework Core@RMS
 - 组件化; 接口化,
 - 内置默认调度
- 所有策略插件化
 - 抢占&预占&突发
 - 可配置
- 标签计算

框架

分配调度上下文



栏目	备注
时间窗口	两个周期：调度周期（充分利用缓存数据）、绑定周期（利用状态数据、持久化）
阶段	预选+优选+绑定
Qos	Request + Limit (涉及驱逐/GC策略，节点超卖策略，)

量化目标

- 装箱节点数
 - 算法, 约束条件 $BFD(I) \leq \frac{11}{9} OPT(I) + 1$
 - 标签合并, 调度空间
 - Incremental阶段
- 碎片化, 资源浪费率WASTE
 - 节点利用率分类, 状态转移, 代价
 - Consolidation阶段
- 能耗, 资源利用和功耗相关性
 - 空机SLEEP|WAKEUP
 - 采购优化
- 售卖率|容量管理|库存服务|经济模型
 - 运营决策, 套餐

功耗

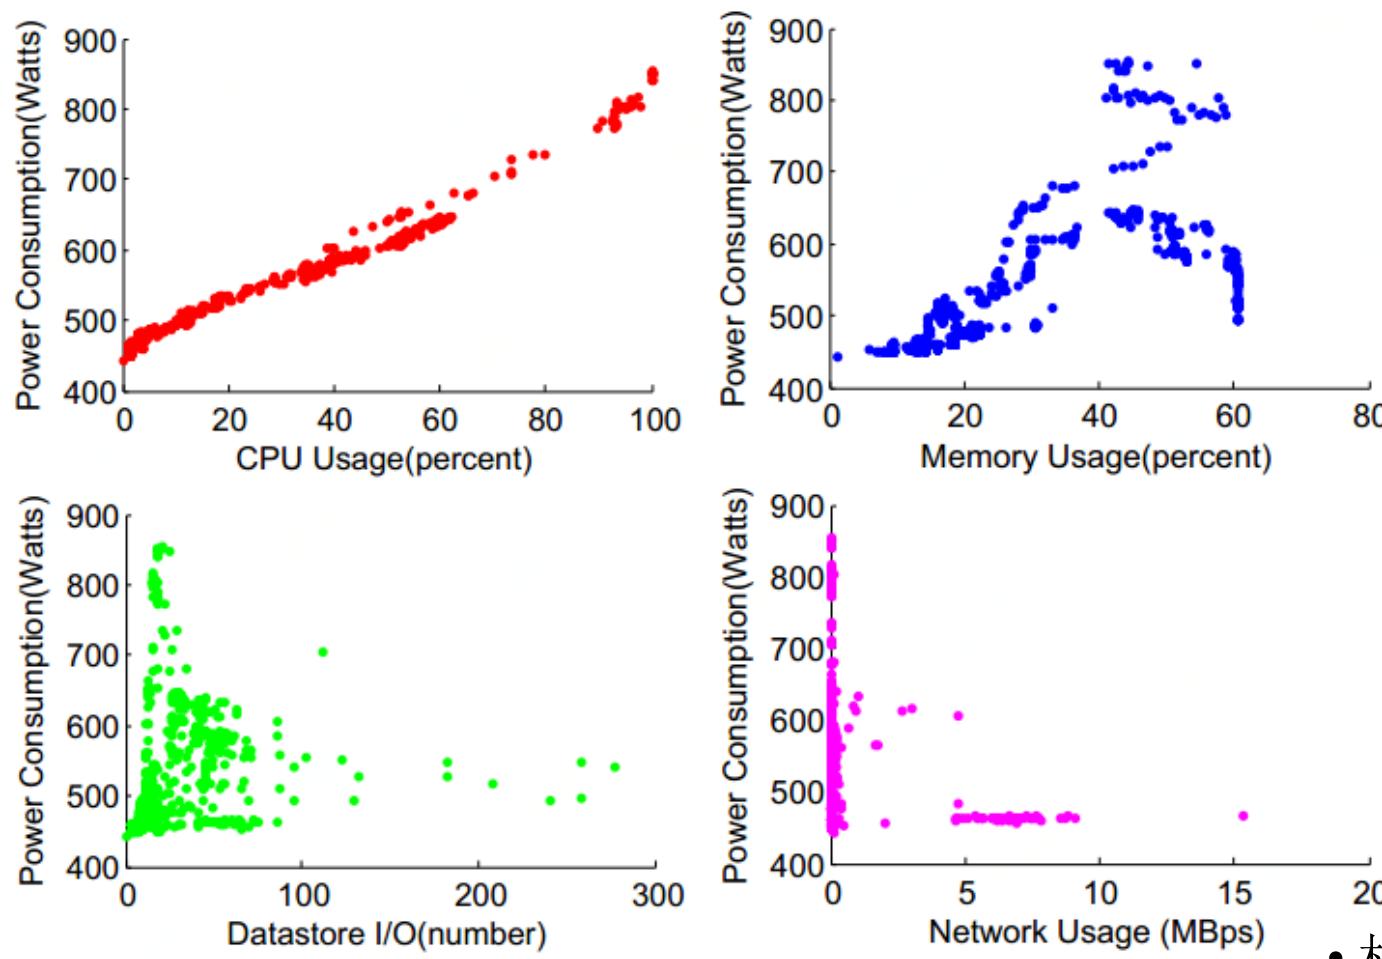


Fig. 3: Relationship between resource usages and power consumption

Table 1: VM instances in the data center

Variable a	Variable b	Correlation Coefficient(a,b)
Power Consumption	CPU Usage	0.990667
	Memory Usage	0.681875
	Datastore I/O	0.253151
	Network Usage	-0.12116

- 相关性结论[1]
 - CPU与能耗高度正相关
 - 网络与能耗无关
 - 磁盘，内存相关性偏弱
- 资源利用和能耗
 - 相关因子

功耗-约束

- 机架因素
- 供电额度
- GPU
- 四代机
- TOR拓扑

相关性结论

- Idle机器耗能 = 70% * Full机器耗能
 - Sleep状态
- 能耗(cpu利用率):
 - 功率: $P(u) = k \cdot P_{max} + (1 - k) \cdot P_{max} \cdot u$
 - 能耗: $E = \int_{t_0}^{t_1} P(u(t)) dt$
- 节点&实例Cpu利用率
 - Share, Set和Mix是统一的
 - 规格系数为权重
 - 绑核和时间片计量上是统一的
- 所有量化都统一到利用率及其预测上

可观测目标

- 可扩展
 - 架构
- 可度量
 - **量化目标** 中的衍生指标
- 可评估
 - 包含衍生指标的统计分析
- 其他
 - 无侵入性, 灵活

实现

- Scheduler Framework
 - 提供插件化和默认调度（兼容）
- 多目标优化
 - Best effort
 - 最大化利用率
 - 最小化浪费
 - 最小化能耗
 - Tradeoff
 - 效能提升优先，一级调度
 - 隔离性兼顾，二级调度

参考

- [1] Virtual machine consolidated placement based on multi-objective biogeography-based optimization

Backup

- Kubernetes Scheduler

FlagSet

File

ConfigMap

Policy

Predicates

Priorities

Extenders

Plugins

Algorithm

Queue

Schedule Pipeline

PreFilter

Filter

PostFilter

Score

Reserve

Permit

PreBind

Bind

PostBind
Unreserve

Schedule Thread

Wait Thread Bind Thread

Informer

Pod

Node

PV

PVC

StorageClass

CSINode

Storage

PDB

Preemption

RC

RS

Service

STS

Deployment

Distribution

Schedule Cache

Initial

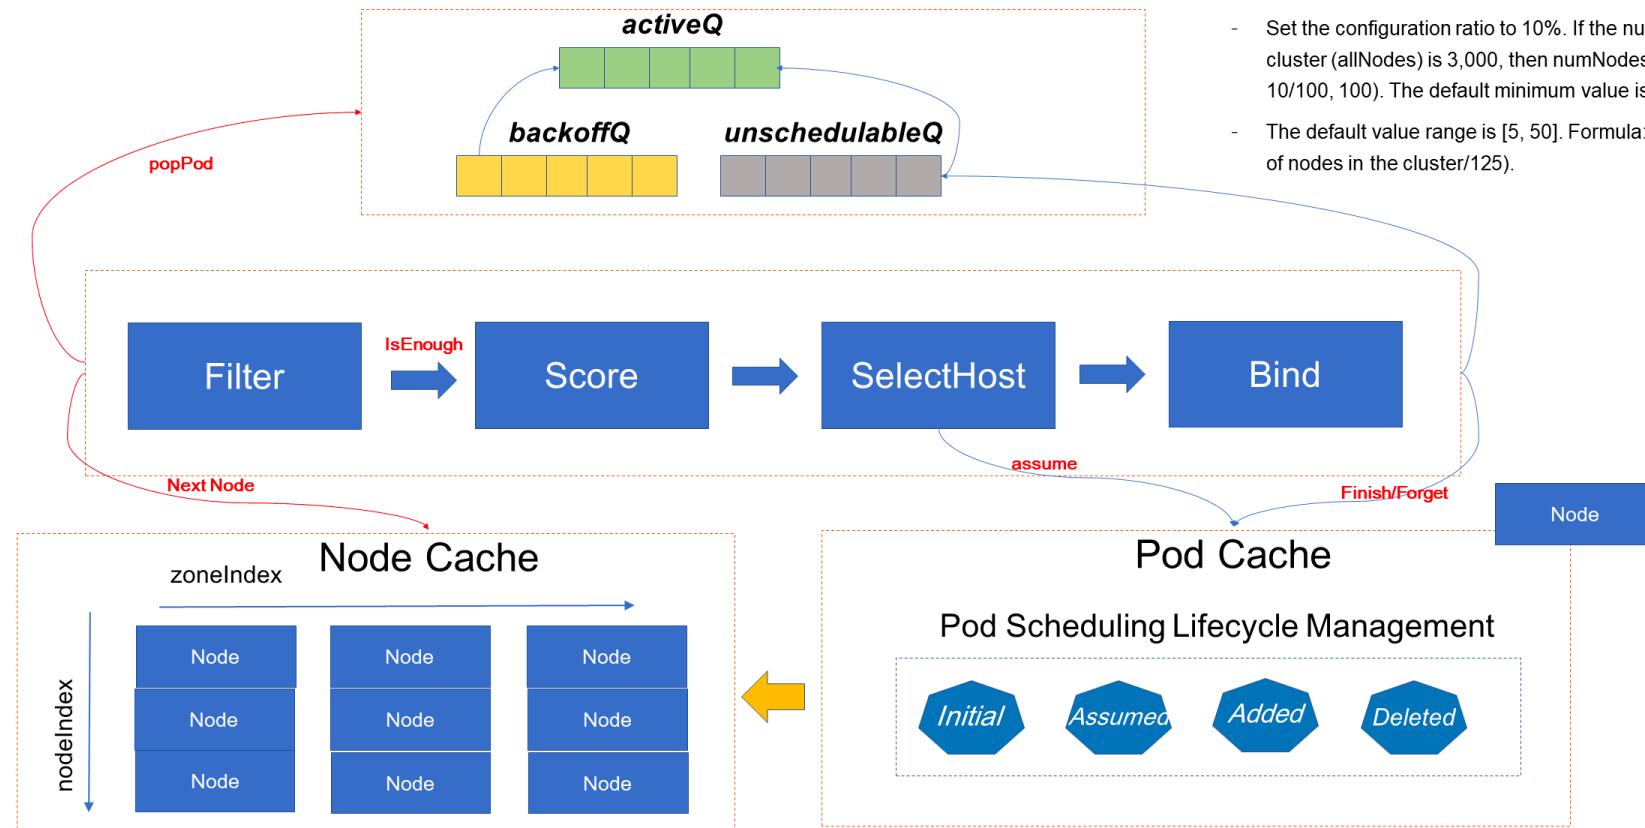
Assumed

Added

Deleted

Pod Scheduling Lifecycle Management

Scheduling Process



Configure the node sampling scope

- Set the configuration ratio to 10%. If the number of nodes in the cluster (allNodes) is 3,000, then numNodesToFind = Max(3000 x 10/100, 100). The default minimum value is 100.
- The default value range is [5, 50]. Formula: Max(5, 50 - Number of nodes in the cluster/125).

- LeastRequestedPriority: prioritizes distribution.
- MostRequestedPriority: prioritizes stacking.
- BalancedResourceAllocation: indicates the fragmentation rate.
- RequestedToCapacityRatioPriority: indicates the score of the ratio.

- Concepts used in the resource usage formula
 - Request: resources that have been allocated to nodes
 - Allocatable: resources that can be scheduled to nodes.
- Score of preferential distribution
 - $(\text{Allocatable} - \text{Request}) / \text{Allocatable} * \text{Score}$
- Formula for calculating preferential stacking
 - $\text{Request} / \text{Allocatable} * \text{Score}$
- Fragmentation rate
 - $\{1 - \text{Abs}[\text{CPU}(\text{Request}/\text{Allocatable}) - \text{Mem}(\text{Request}/\text{Allocatable})]\} * \text{Score}$
- Additional information:
 - Based on the preceding scenarios, a score of 0 is given in all scenarios where $\text{Request} \geq \text{Capacity}$.
- Specified ratio
 - MapConfig of a score with the specified ratio (0:x, 1:x, 2:x, ..., 99:x , 100:x)
 - Formula: $\text{Request} / \text{Allocatable} * \text{MaxUtilization}$

Priorities – Pod Distribution

- Solved problems
 - The pod distribution feature allows you to spread matched pods across different topologies for deployment.
- SelectorSpreadPriority
 - Collect statistics on nodes.
 - TopoPods = Selector condition for the workload of the controller with Exists Pod matching Income Pod
 - $(\text{Sum}(\text{TopoPods}) - \text{TopoPods})/\text{Sum}(\text{TopoPods})$
- ServiceSpreadingPriority
 - An official source indicates that ServiceSpreadingPriority is very likely to replace SelectorSpreadPriority.
 - Collect statistics on nodes.
 - TopoPods: compliant with the selector condition specified by the service where pods are located
 - $(\text{Sum}(\text{TopoPods}) - \text{TopoPods})/\text{Sum}(\text{TopoPods})$
- EvenPodsSpreadPriority
 - topologyKey indicated by spec
 - TopoPods = spec-compliant labelSelector
 - Score calculation formula
 - NodeTopoPods = cumulative TopoPods at the node level
 - MaxDif = $\text{Max}(\text{NodeTopoPods}) - \text{Min}(\text{NodeTopoPods})$
 - $\text{NodeTopoPods} - \text{MaxDif}/\text{MaxDif}$

- **NodeAffinityPriority**
 - TopoPods: cumulative instances of affinity compliance
 - $\text{TopoPods}/\text{Max}(\text{TopoPods}) \times \text{MaxPriority}$
- **ServiceAntiAffinity**
 - Evenly distribute the pods in a service based on a label value of the node.
- **NodeLabelPrioritizer**
 - Preferentially distribute pods to a node with or without a specific label.
- **ImageLocalityPriority:**
 - Perform scheduling based on image affinity.

Configure the Kubernetes Scheduler

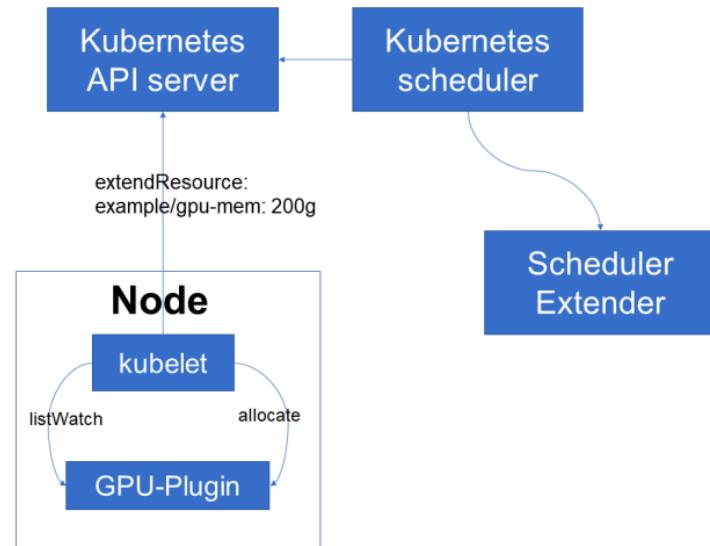
- Start the Kubernetes scheduler
 - Default configuration (--write-config-to)
 - Configuration file (--config)
- Configuration file description
 - schedulerName: scheduling based on Pod.SchedulerName
 - algorithmSource: used to configure algorithms
 - hardPodAffinitySymmetricWeight: used to configure the affinity weight
 - percentageOfNodesToScore: the threshold for exit when the ratio of the filtered nodes to all nodes is equal to this value
 - bindTimeoutSeconds: the timeout period of the binding phase

```
# Startup with the default configuration
kube-scheduler \
    --kubeconfig=/etc/kubernetes/kubeconfig/scheduler.kubeconfig \
    --v=3
# Startup with the specified configuration file, which is recommended when custom
parameters are required.
kube-scheduler \
    --config=/home/admin/scheduler/conf/scheduler-policy.config \
    --v=3
```

Scheduler Extender

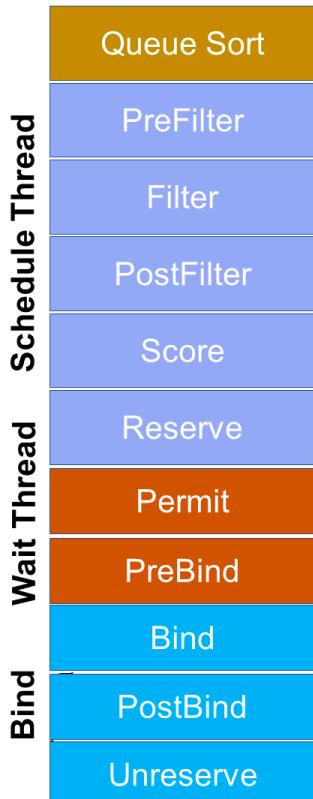
- What can we do?
 - Keep the original scheduling code and directly start the service as a plugin to be called by the Kubernetes scheduler (similar to webhook).
 - Support injection of predicate, preempt, priority, and bind.
 - A type of extended resource, which can be bound to only one extender.
- Configuration procedure
 - Configuration file description
- Example
 - Apply for GPU memory. Only the extender knows the memory size of each GPU, so add the extender filter.

```
kind: Policy
apiVersion: v1
extenders:
- urlPrefix: "http://xxxx:xxxx/scheduler-gpu-extender"
  filterVerb: filter
  weight: 1
  enableHttps: false
# 调度器传递nodenames列表，而不是nodeinfo 列表
nodeCacheCapable: true
# 调用extender服务报错或者网络不可达的时候，是否可忽略extender
ignorable: false
managedResources:
- name: "example/gpu-mem"
# resourceFit阶段是否忽略这个资源的校验
| ignoredByScheduler: false
```



Scheduler Framework

- Extension point usage
- Concurrency model

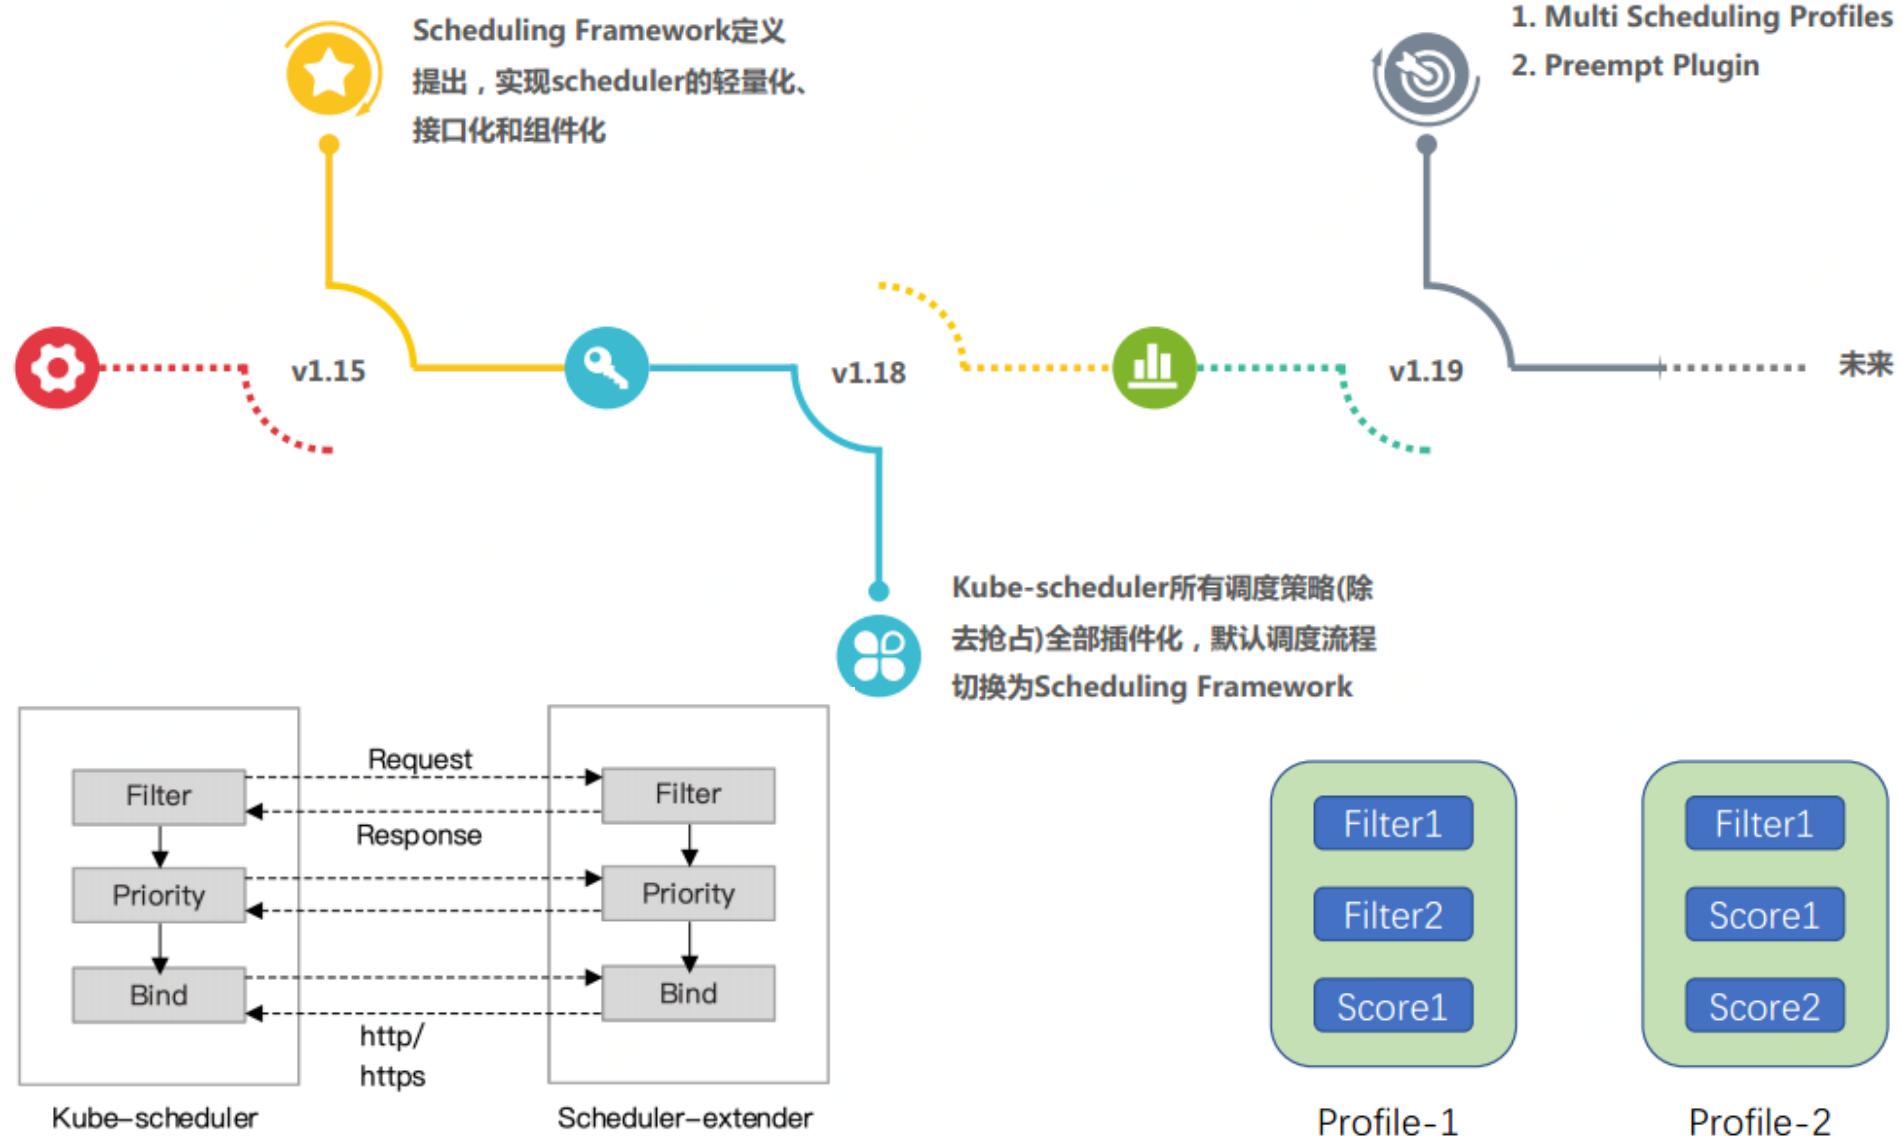


- QueueSort: supports the sorting of custom pods.
- PreFilter: preprocesses pod requests.
- Filter: creates custom filter logic.
- PostFilter: processes logs and metrics or preprocesses data before the score phase.
- Score: creates custom score logic.
- Reserve: records the memory usage of resources through stateful plugins.
- Permit: supports the Wait, Deny, and Approve actions, which can be used as Gang insertion points.
- PreBind: performs operations, such as attaching disks to nodes, before nodes are bound.
- Bind: processes a single pod through a single BindPlugin.
- PostBind: implements the logic after successful binding.
- Unreserve: implements rollback when an error is returned in any phase from Permit to Bind.

问题

- Label 划分资源池
 - 不利于资源利用率
- 调度性能
 - Qps
 - 成功率 (race condition)
 - 批量 (dp)

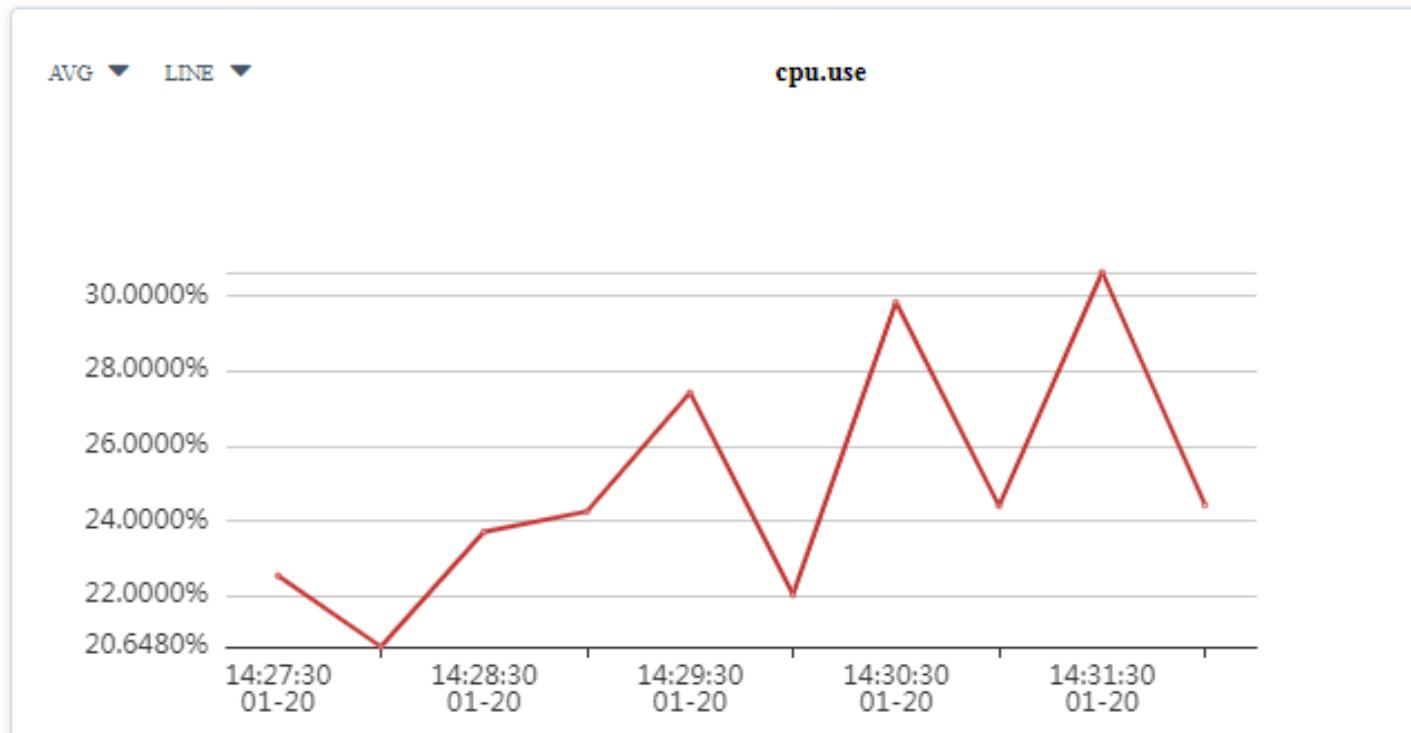
现状



Our现状

节点&实例

- 10.217.35.88
- 2021-01-20 14:27:43 - 2021-01-20 14:32:43



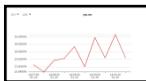
i-zg7l9bs3mf
i-7n345sywim
i-uqmi1k87gq
i-fx4mvsycow
i-09mnv9uiju
i-57bfupdns0
i-77dqm7dyph
i-r3zxvccb95z
i-j8mr1w2tqm
i-rul8xuniqh
i-e180y4gl3s
i-iwa93w60s3
i-18c109qsx9
i-0wt7gkumiq
i-d2aytdu16k
i-egwt92d5ok
i-rtp6knj9np
i-6dtdntq4sl
i-b3oc3w8dlc
i-fukjbbow8d
i-9vflikgvxs
i-hzudsnlnh6
i-gmjzhe1cyi
i-vtqb6skyf3
i-vhqbz9h8r8
i-xb1pqdqguw
i-910ju1jhdd

物理机: 10.217.35.88

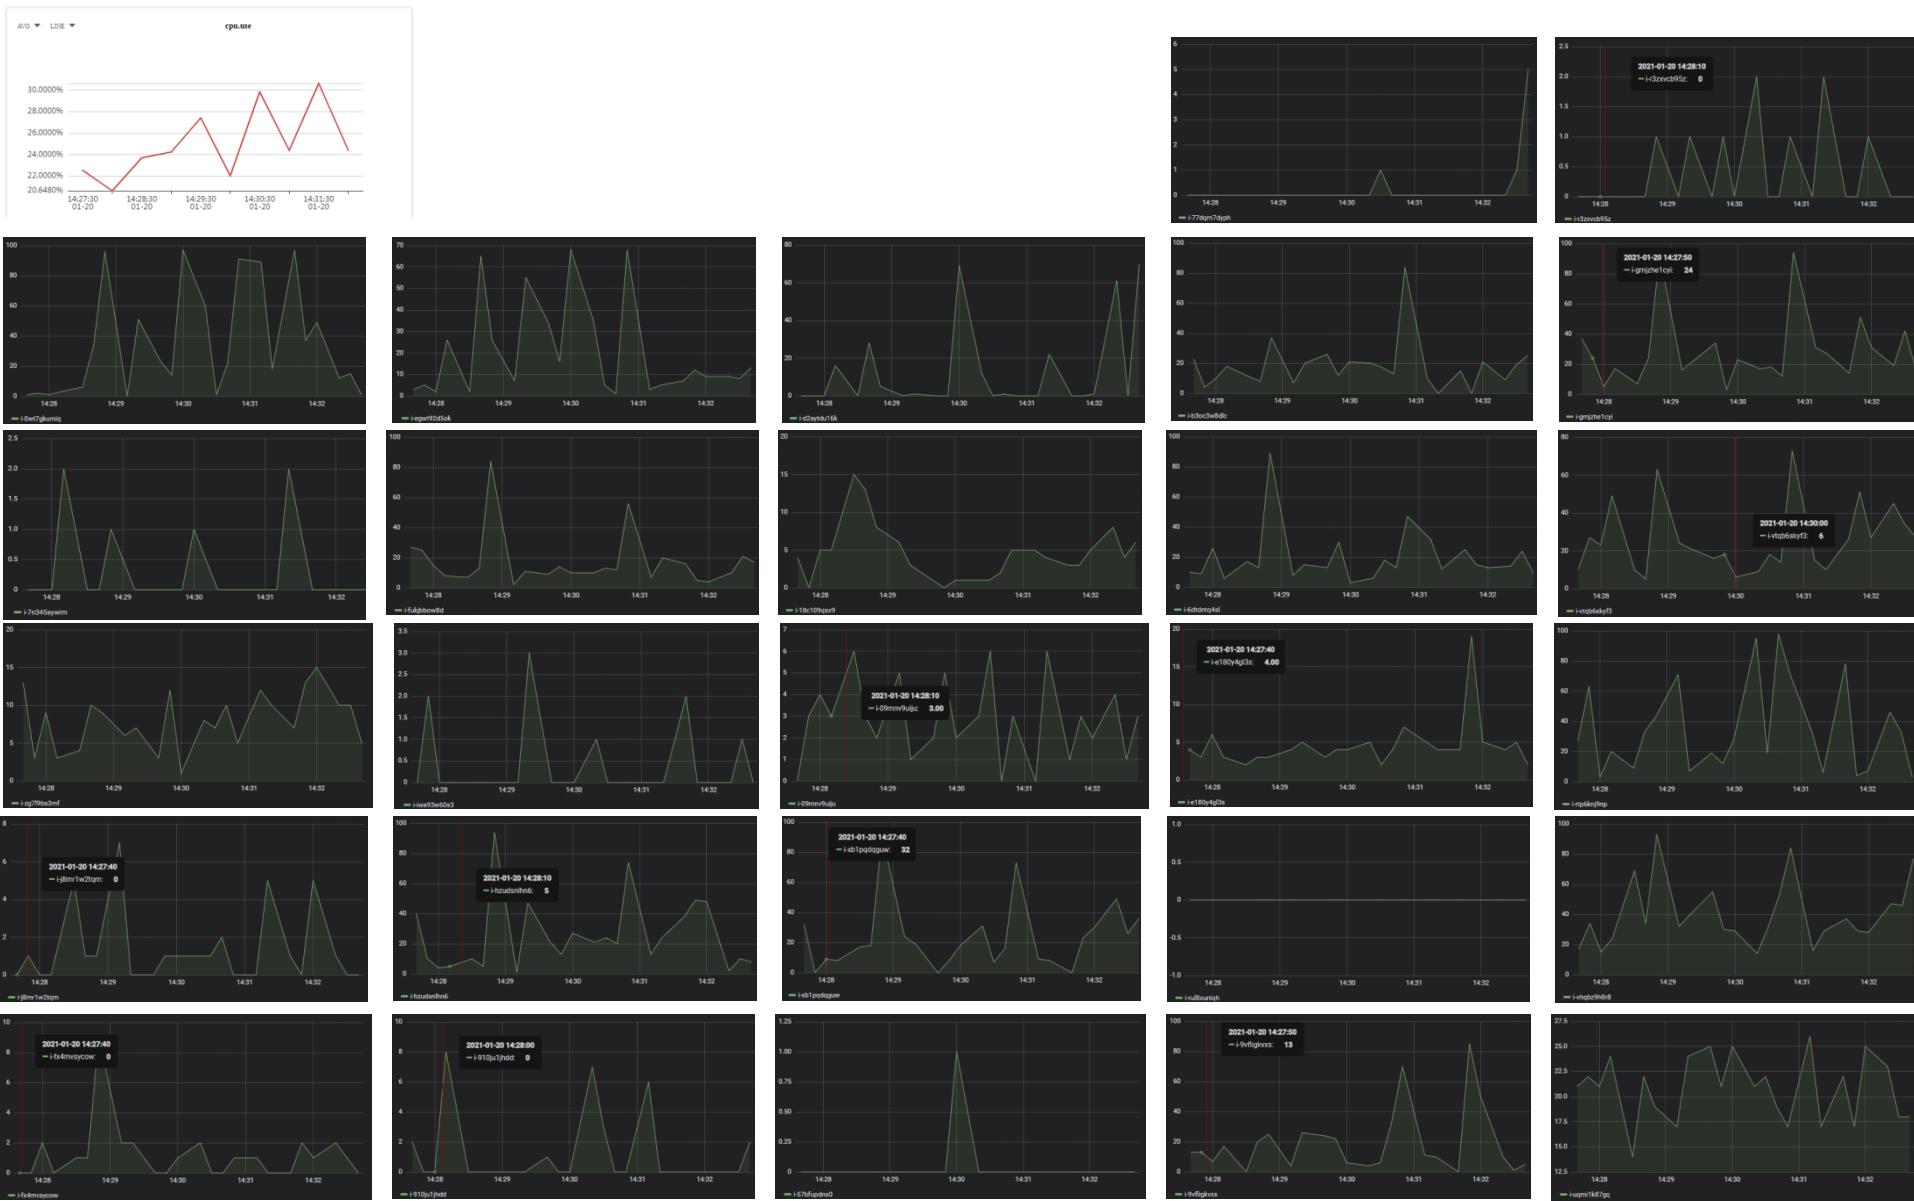
时间段: 2021-01-20 14:27:43 - 2021-01-20 14:32:43

资源: 已全部售卖

物理机:



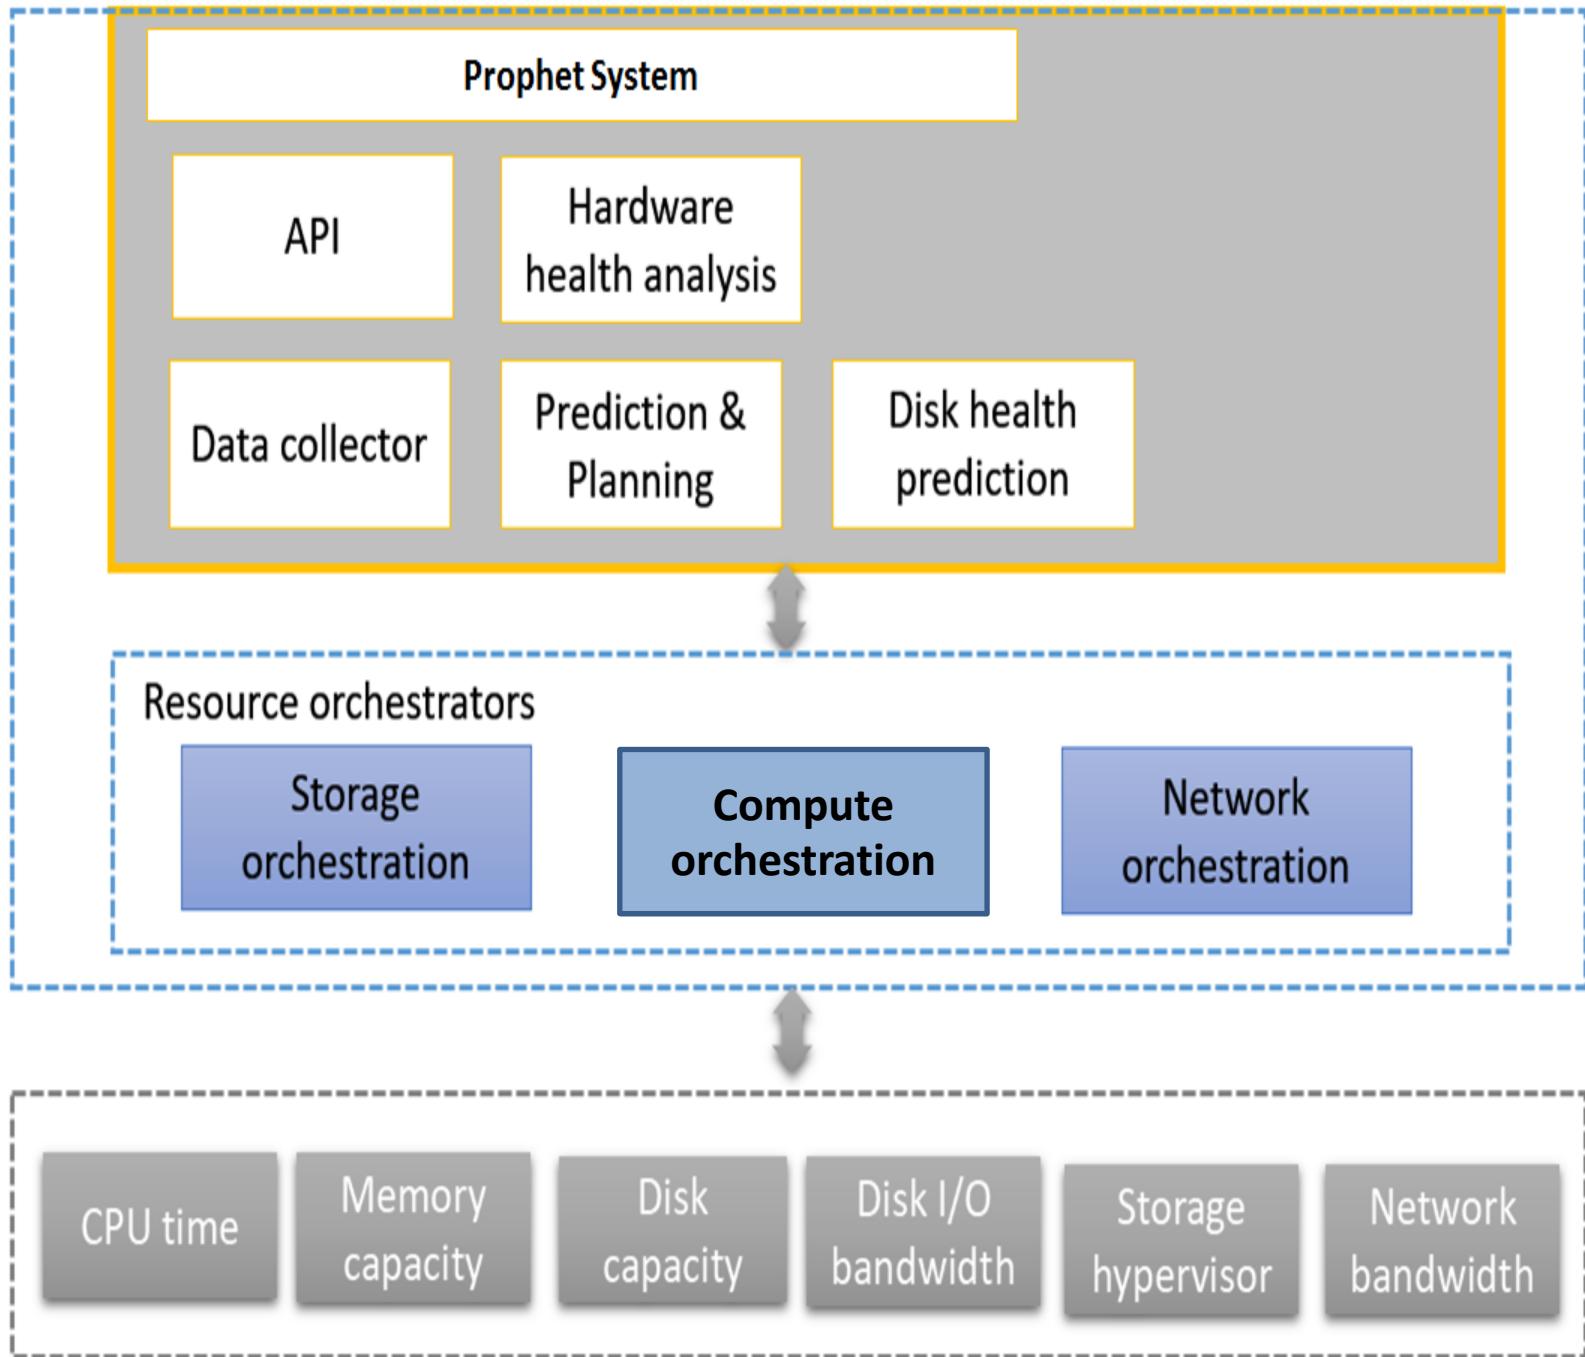
实例:



Proposal

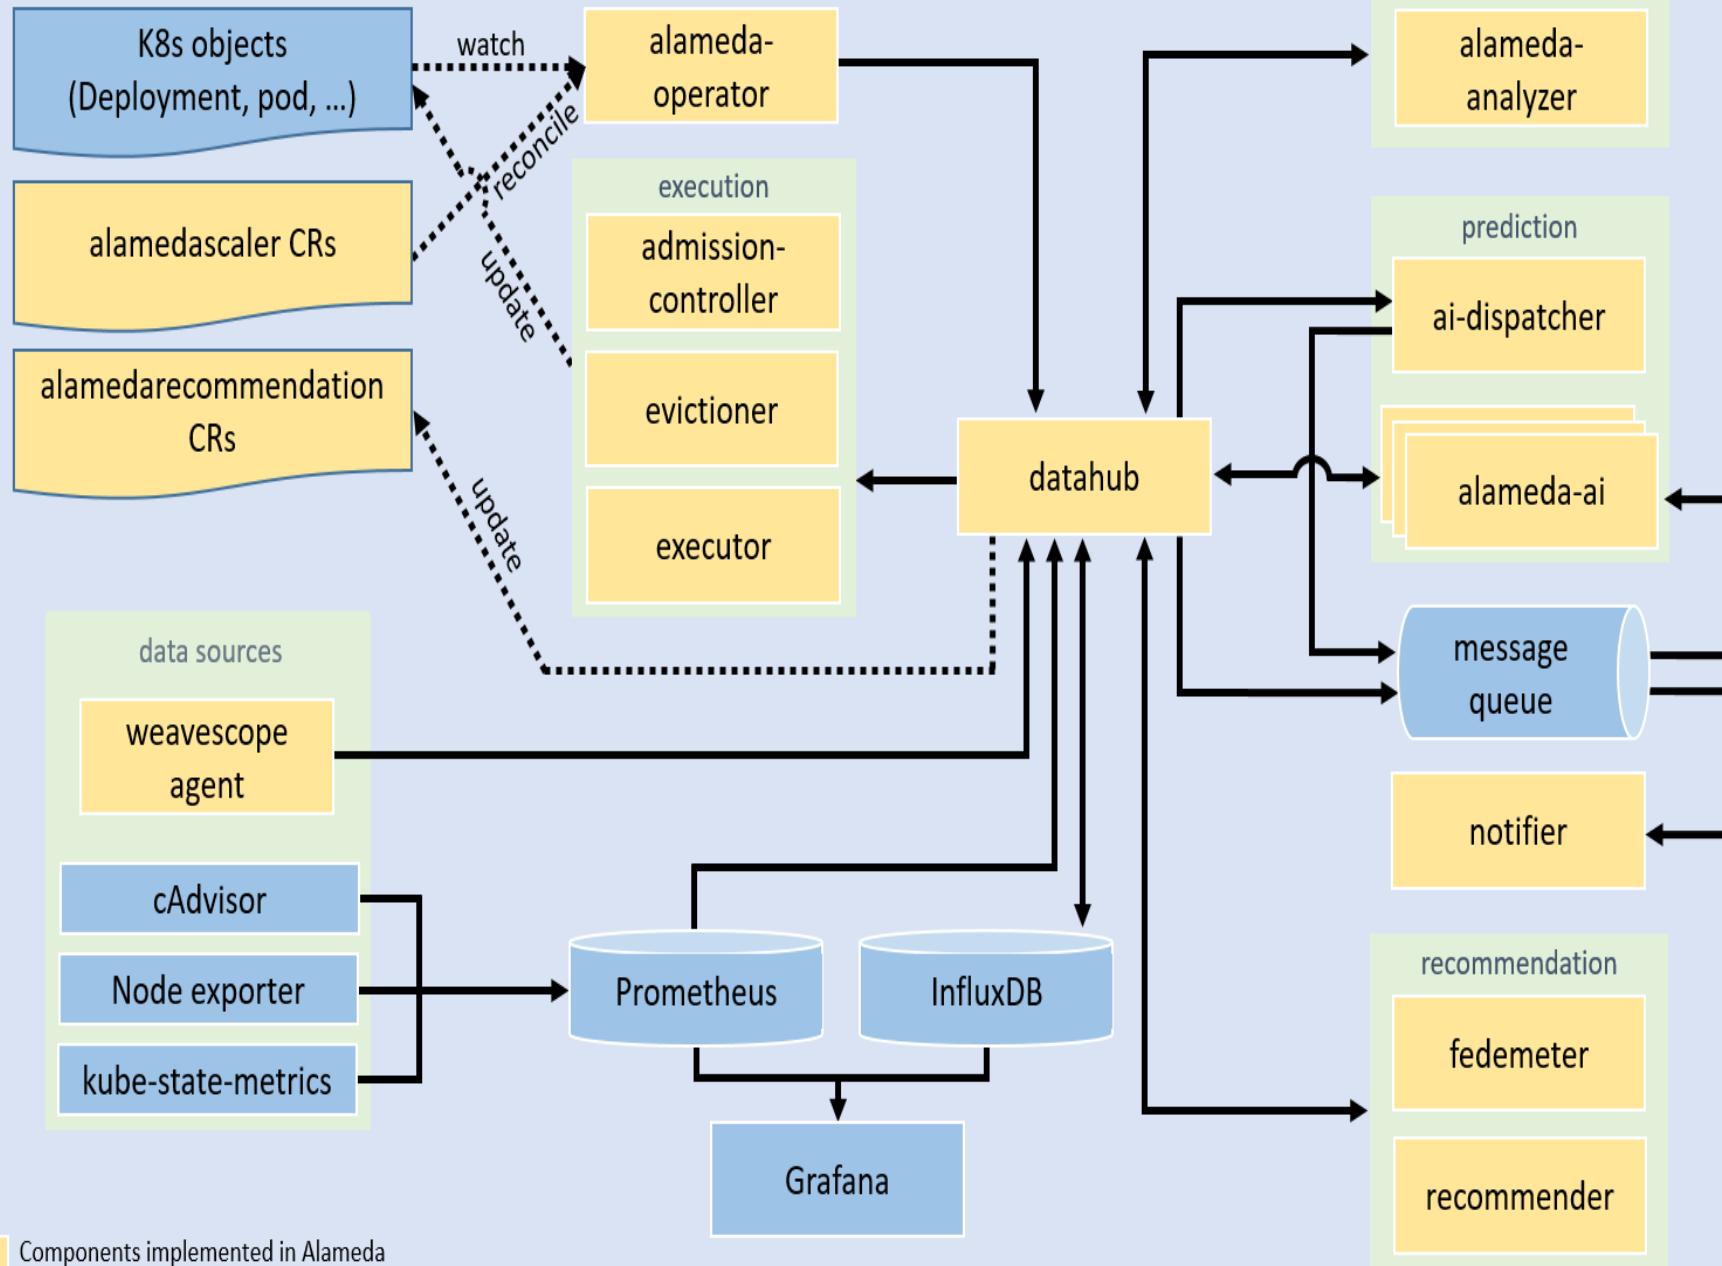
- XXXX is a prediction engine that foresees future resource usage of cluster from the cloud layer down to the instance level. We use machine learning technology to provide intelligence that enables dynamic scaling and scheduling of your instances - effectively making us the “brain” of cluster resource orchestration. By providing full foresight of resource availability, demand, health, impact and SLA, we enable cloud strategies that involve changing provisioned resources in real time.

Jvirt
Ecosystem

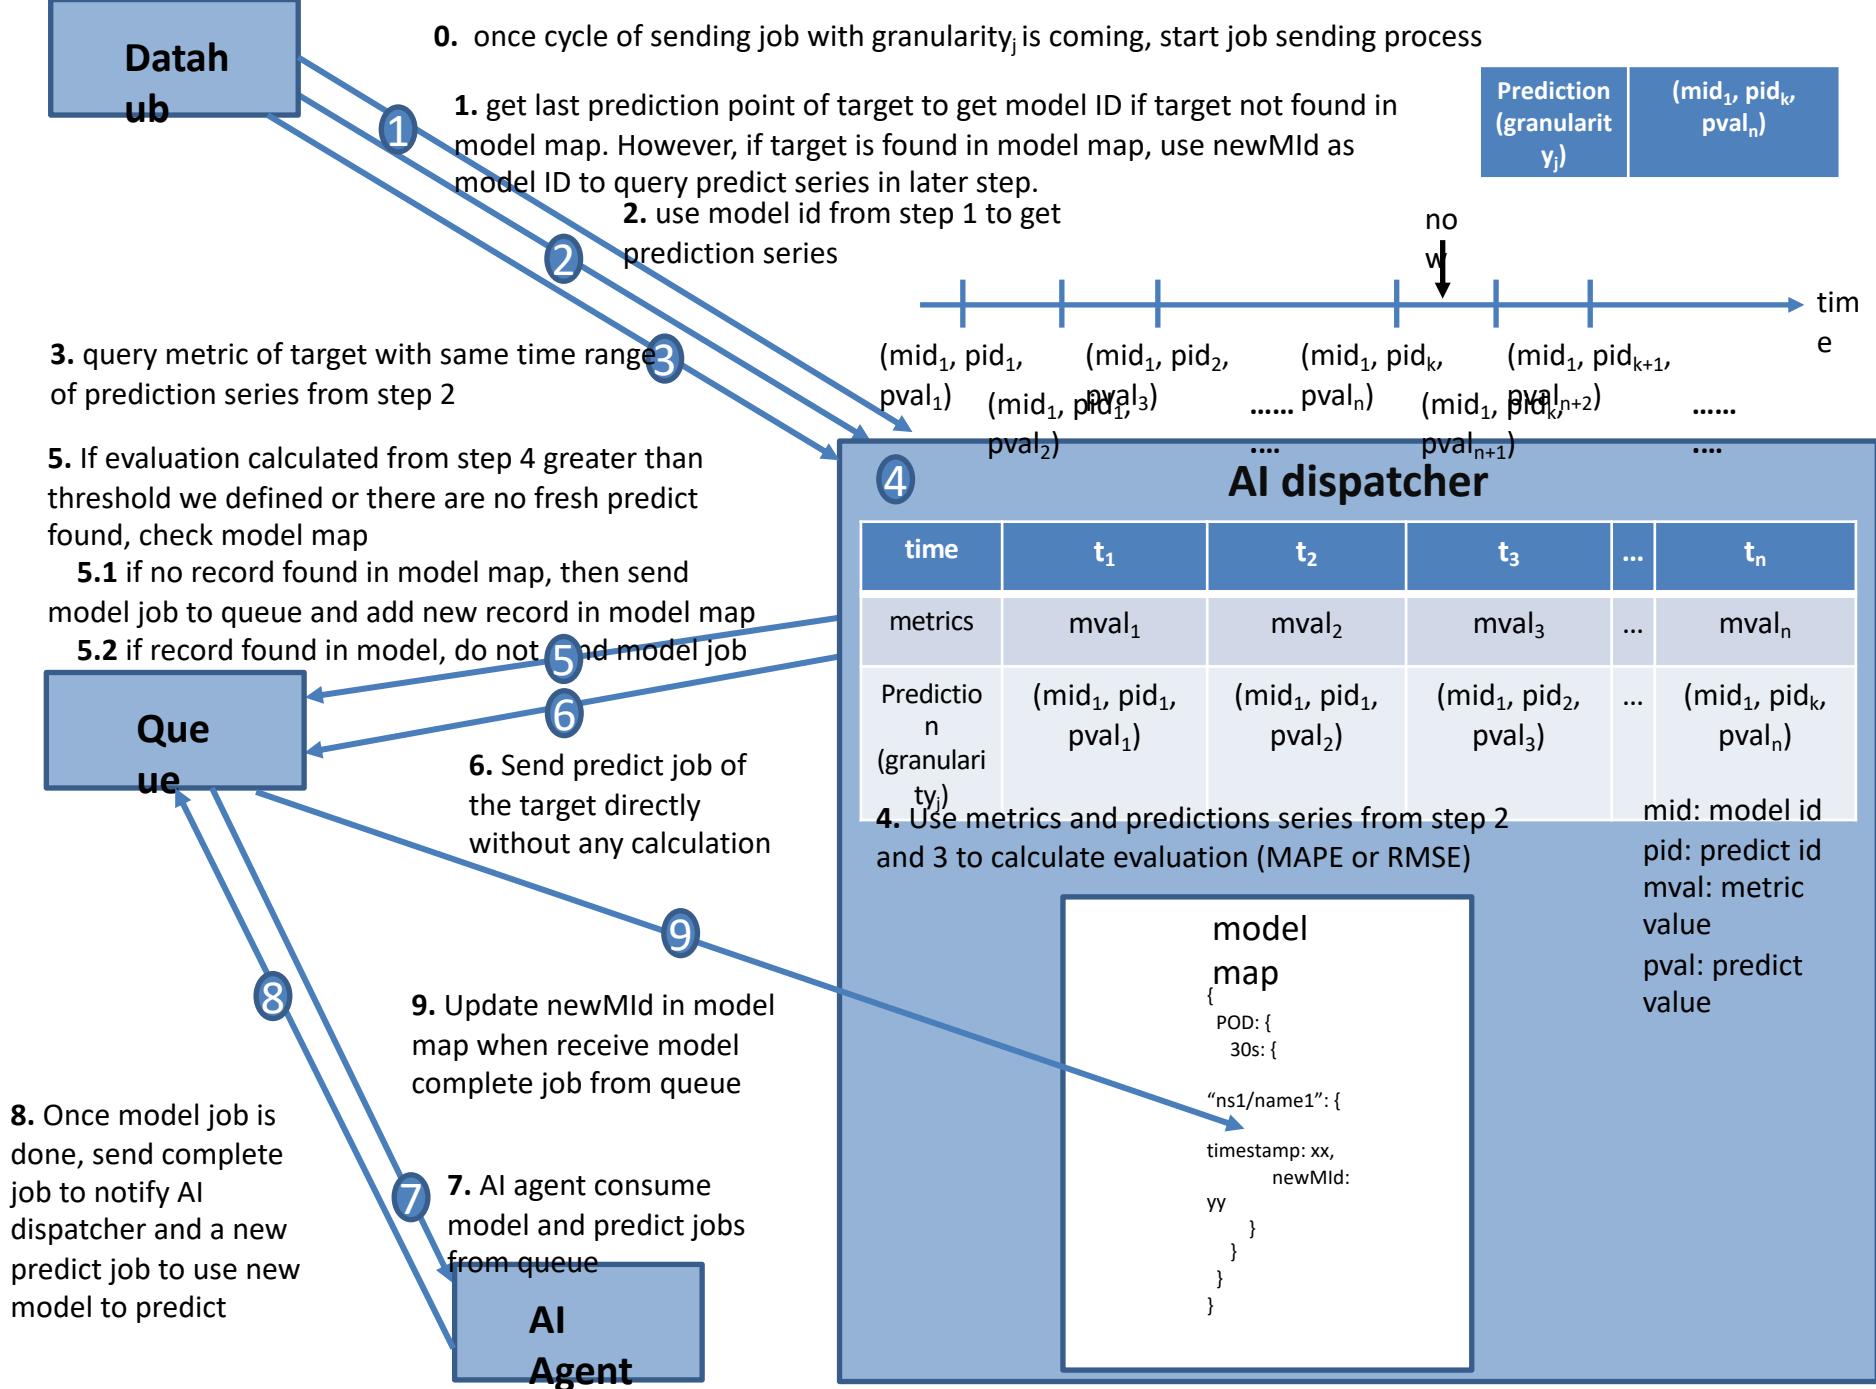


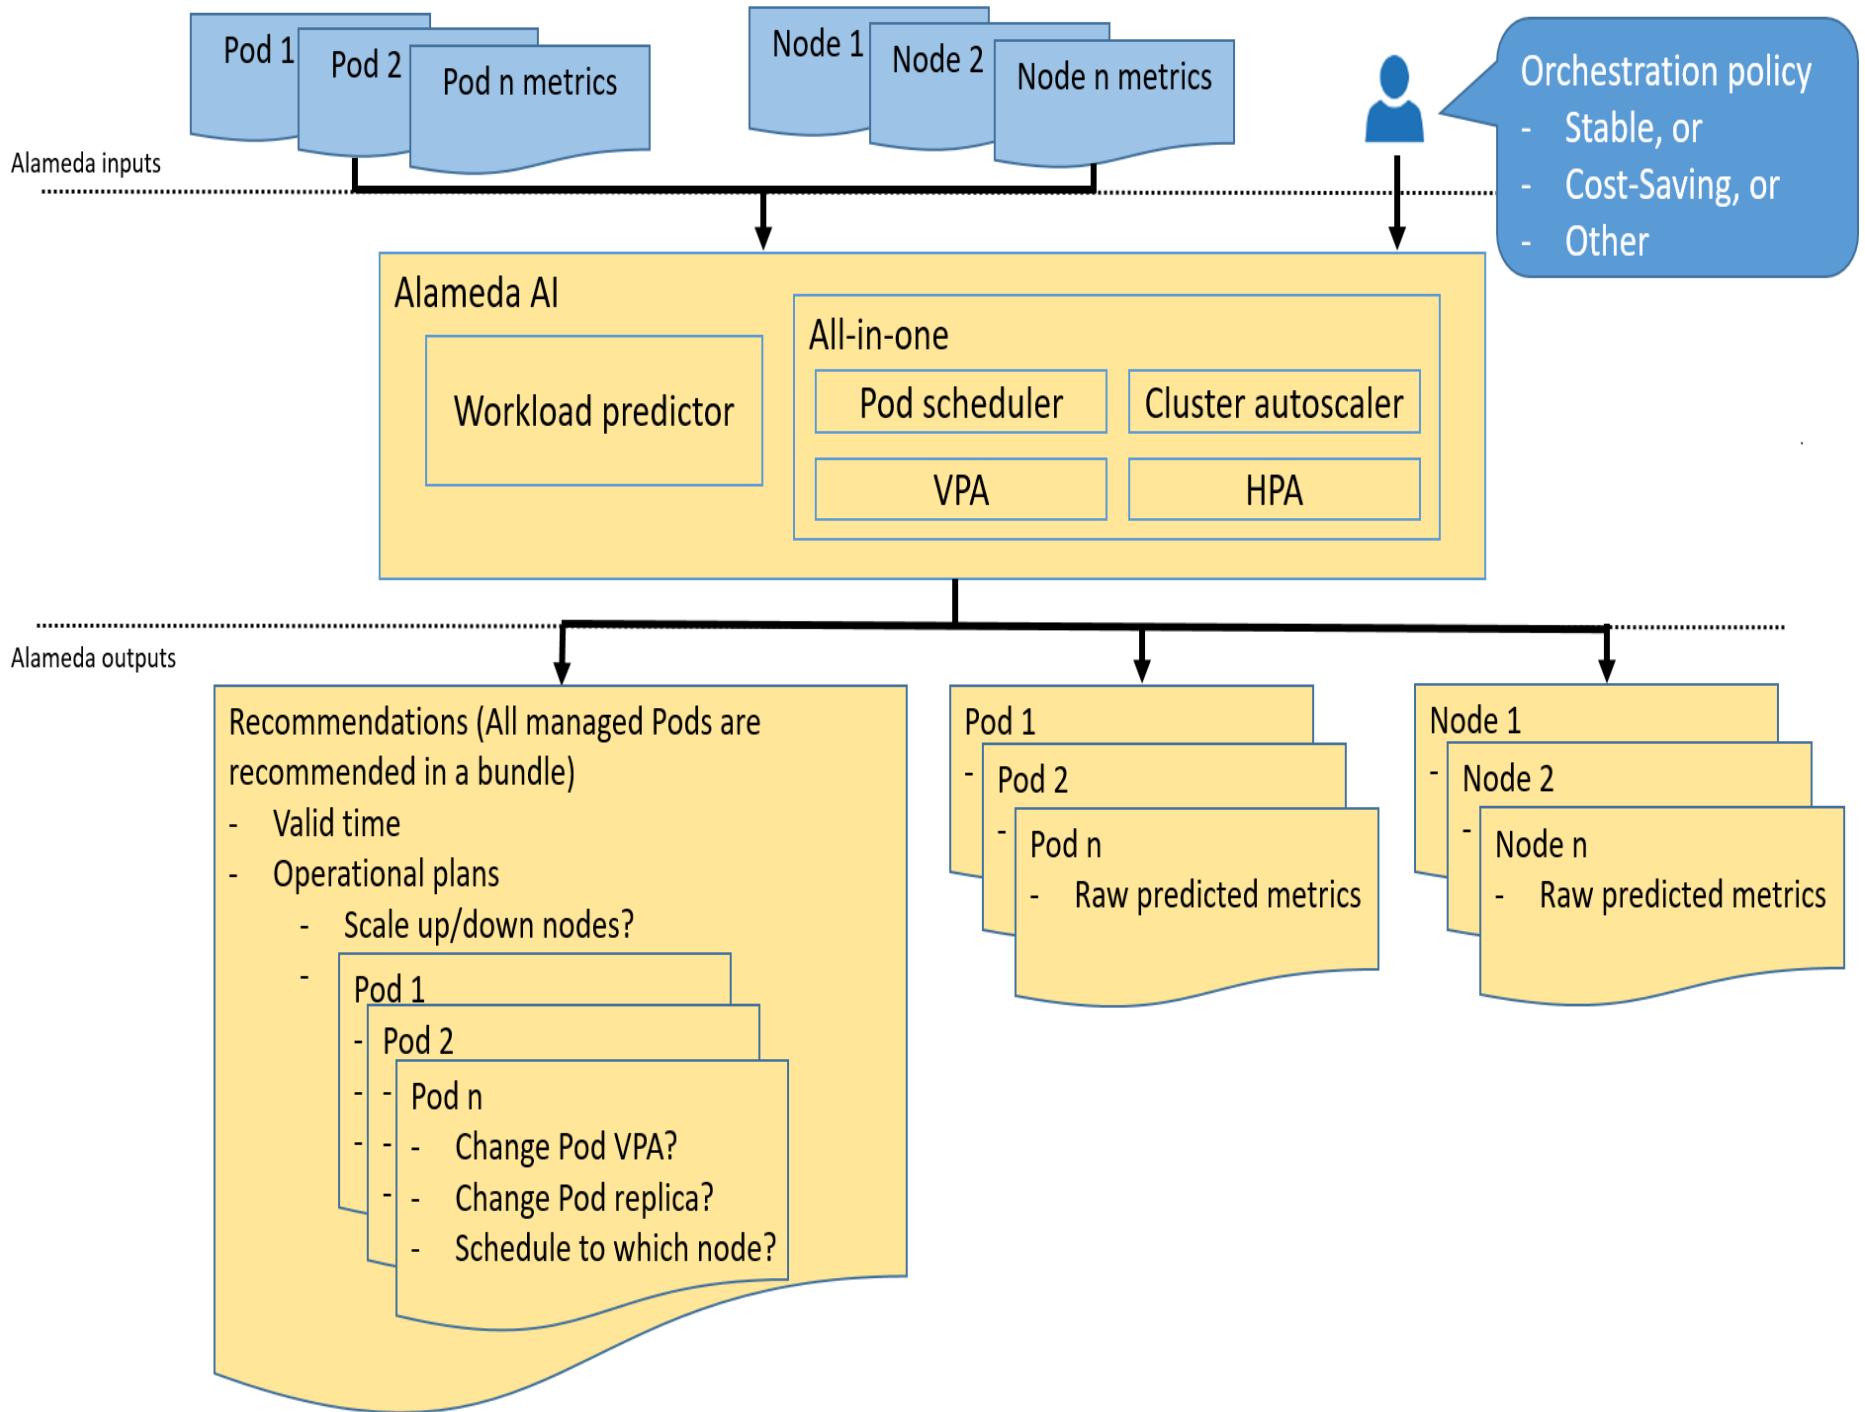


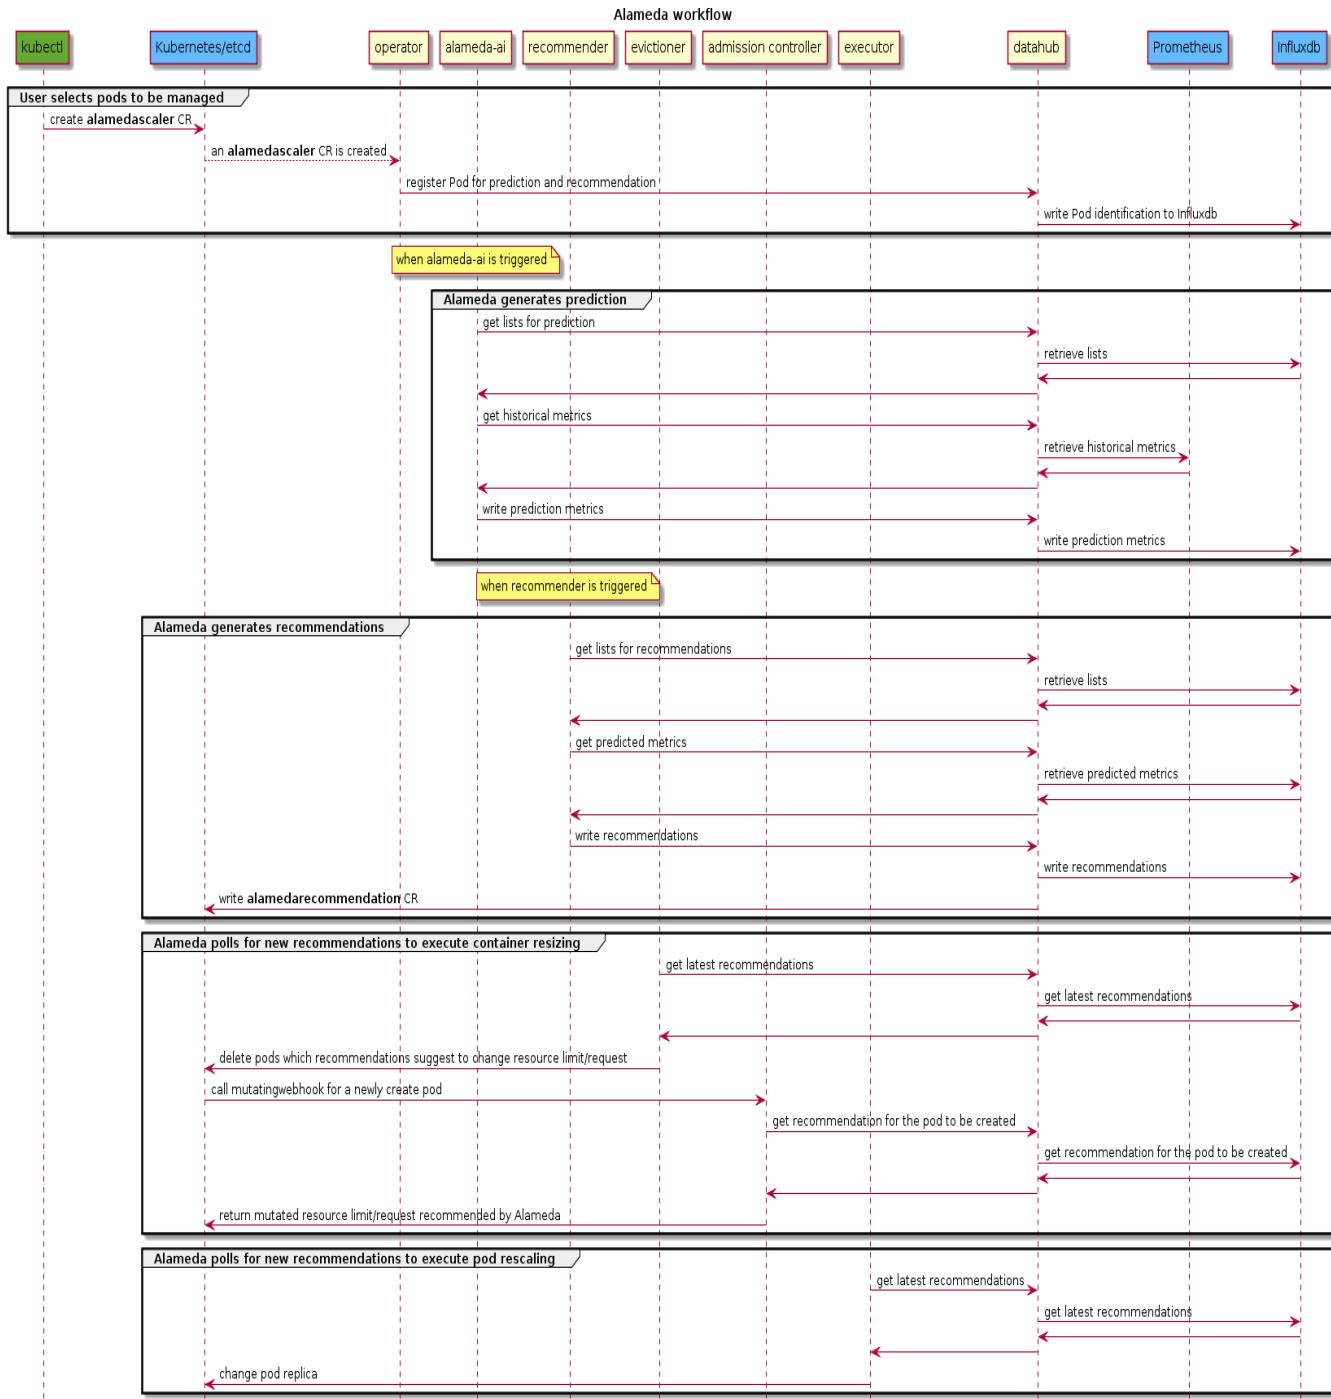
K8s cluster



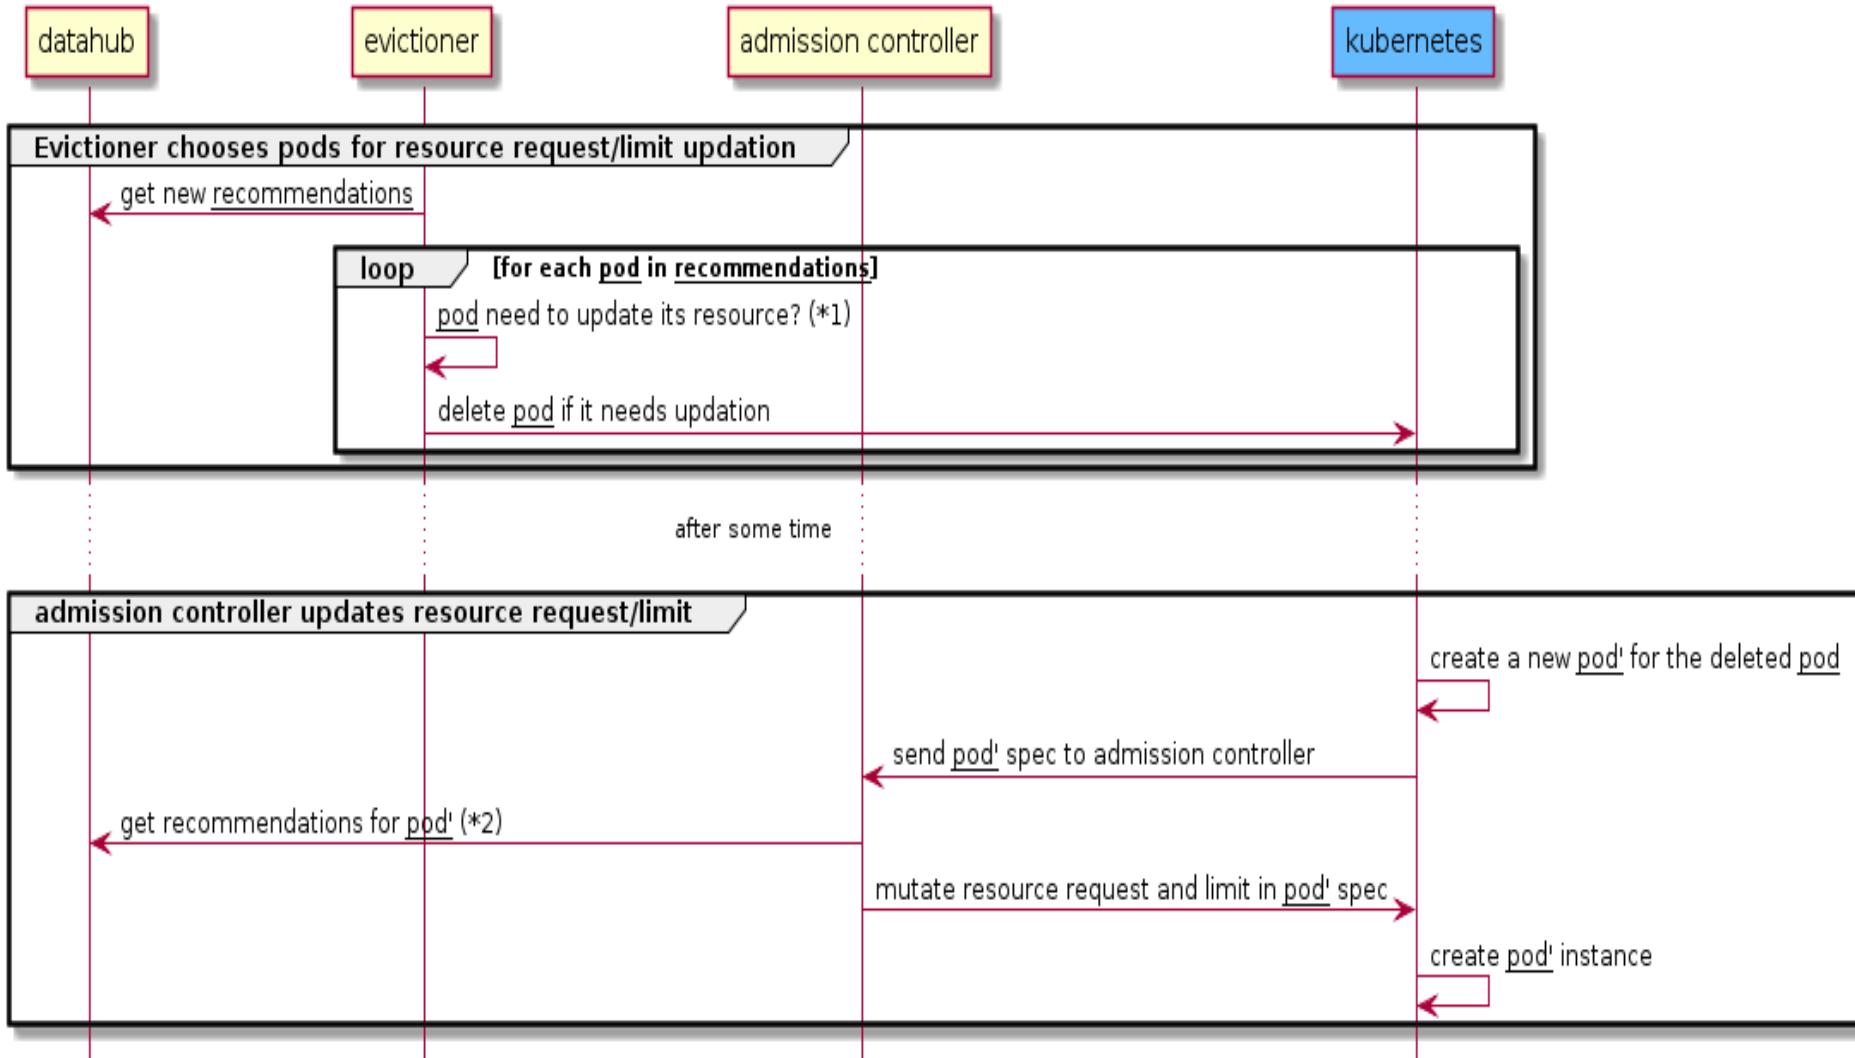
Yellow boxes represent components implemented in Alameda



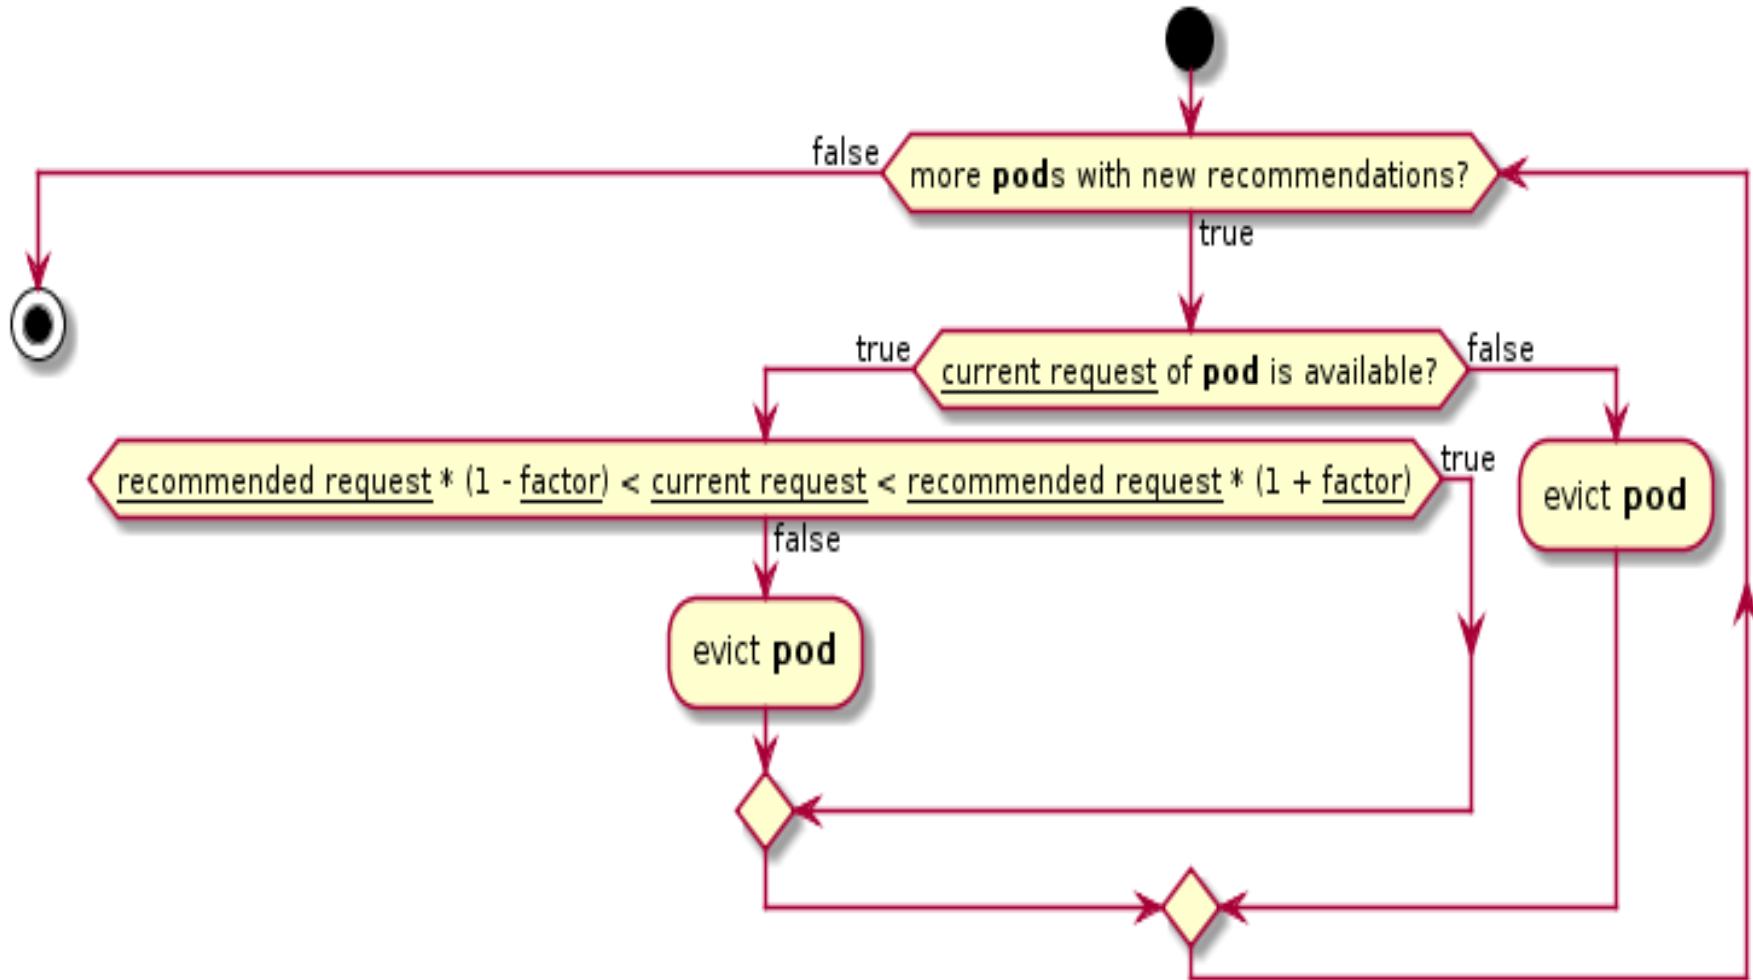




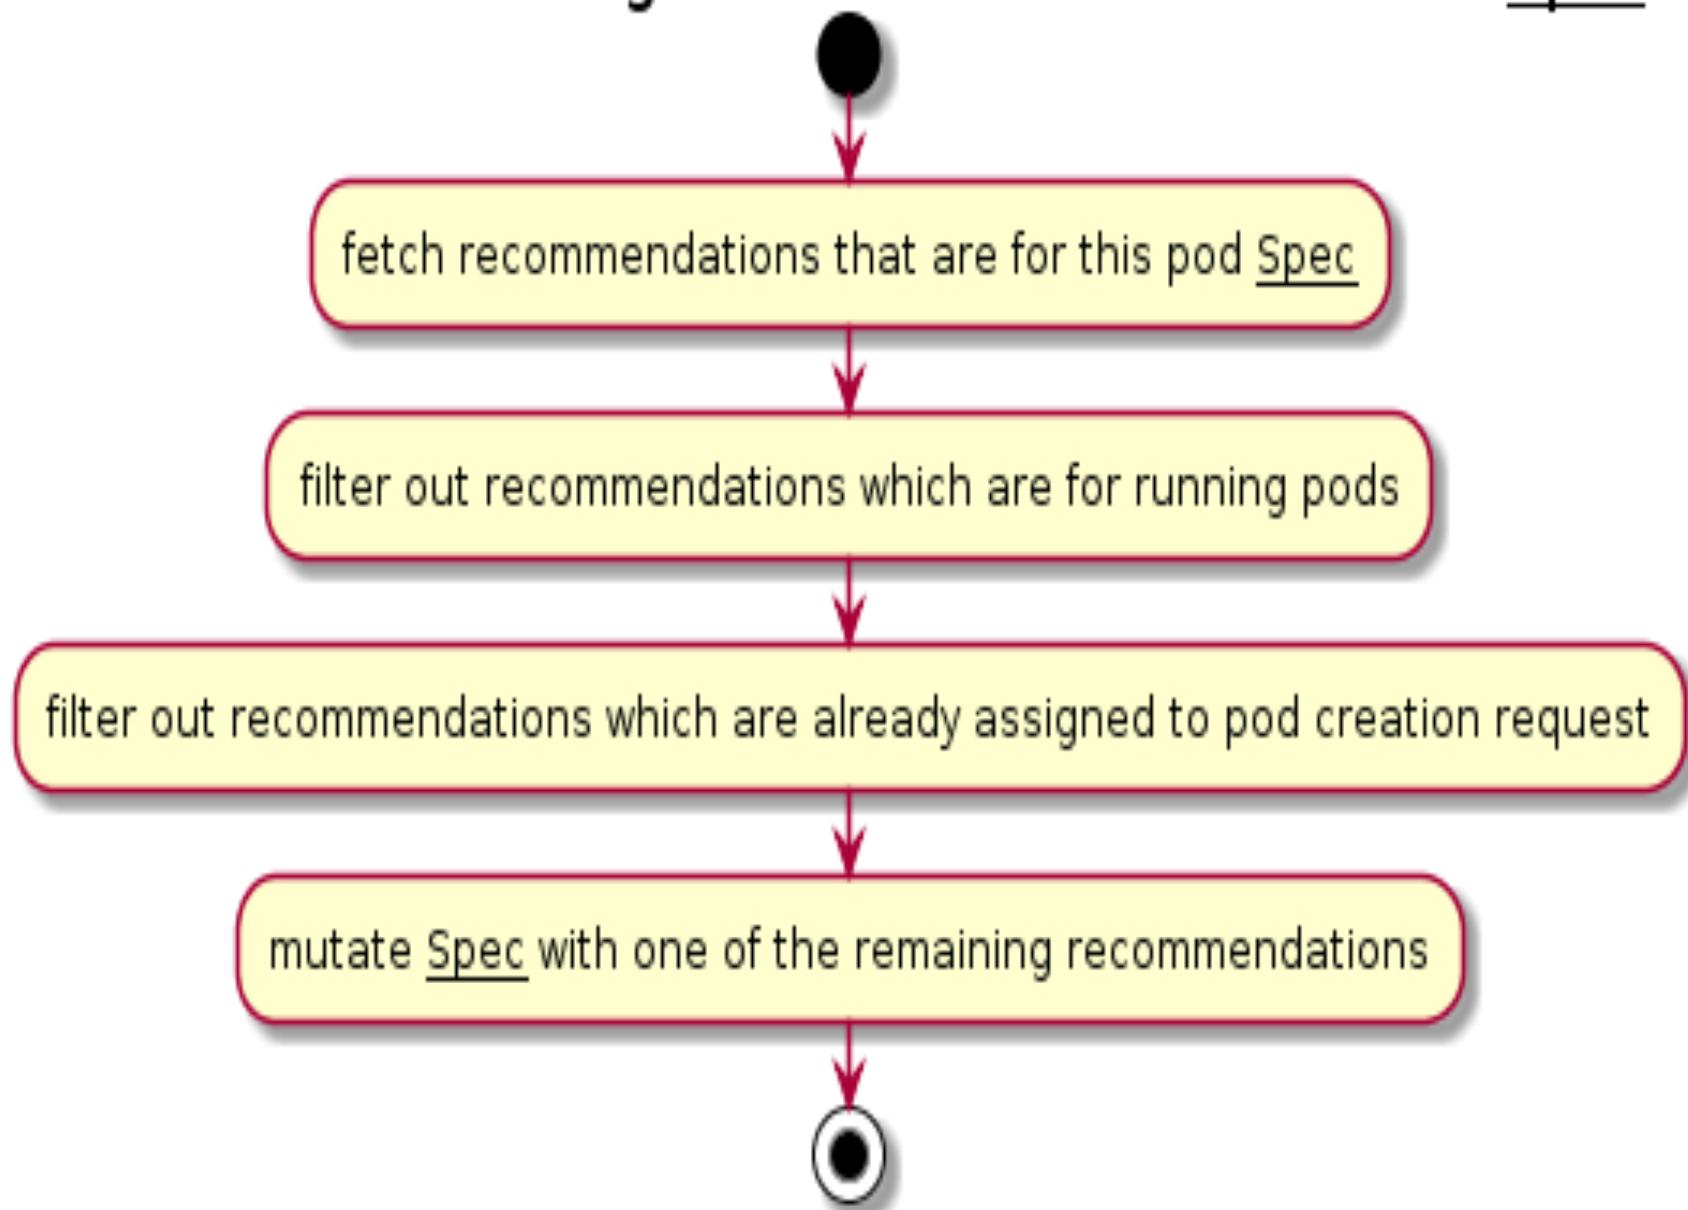
Recommendation Execution Workflow



Control Flow to Choose Pods for Eviction



Control Flow to Assign Recommendation for a Pod Spec



参考

- <https://pkg.go.dev/github.com/containers-ai/alameda>

参考

- <https://github.com/umbertogriffo/Predictive-Maintenance-using-LSTM>
- <https://github.com/AICoE/prometheus-anomaly-detector>
- <https://github.com/nfrumkin/forecast-prometheus>
- <https://www.slideshare.net/GeorgOettl/analyze-prometheus-metrics-like-a-data-scientist>
- <https://www.ericsson.com/en/blog/2019/6/automated-fault-management-machine-learning>
- <https://github.com/VictoriaMetrics/VictoriaMetrics>
- <https://github.com/nfrumkin/forecast-prometheus>
 - 《[*Distributed Systems Observability*](#)》
- <https://dbaplus.cn/news-134-3008-1.html>
- <https://mc.ai/review-of-fault-management-with-machine-learning-techniques/>
- <https://medium.com/@prasad.kumkar/review-of-fault-management-with-machine-learning-techniques-5d603fbbef5>
- <https://next.redhat.com/2019/11/18/prometheus-anomaly-detection/>